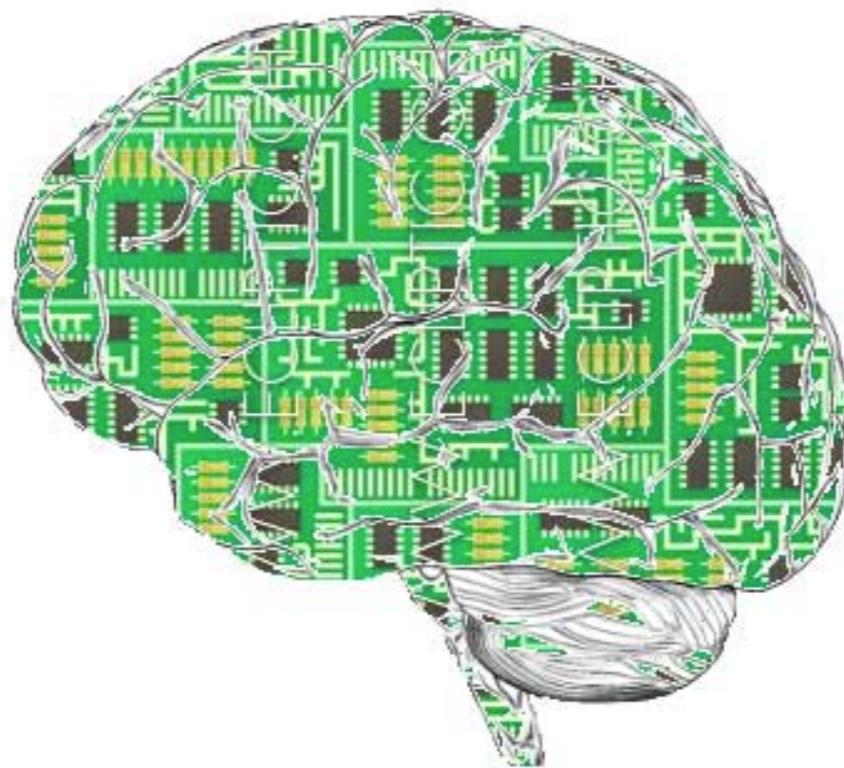




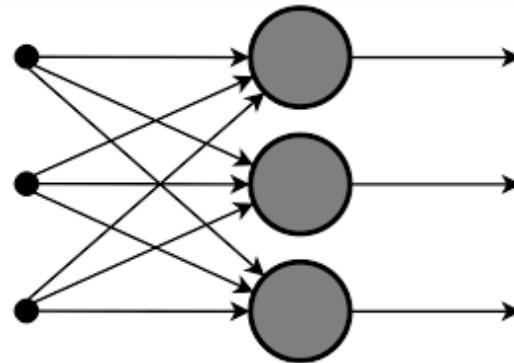
Uma introdução a Deep Learning

Dsc. Nauber Gois
Fametro
Serviço Federal de
Processamento de Dados

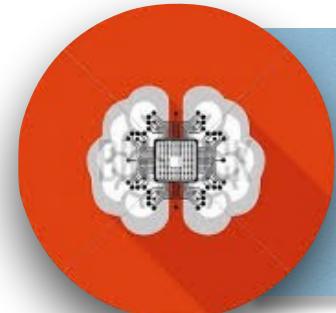


Cognitive Ai

Agenda



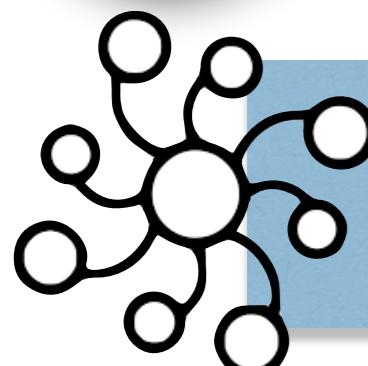
Introdução a Deep Learning



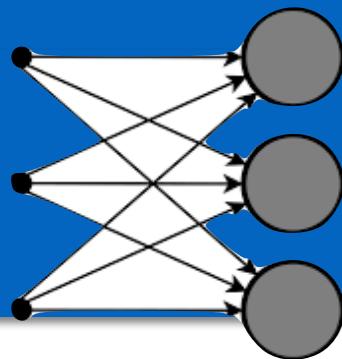
Redes Neurais Convolucionais



Recurrent Neural Networks



Generative Adversarial Networks

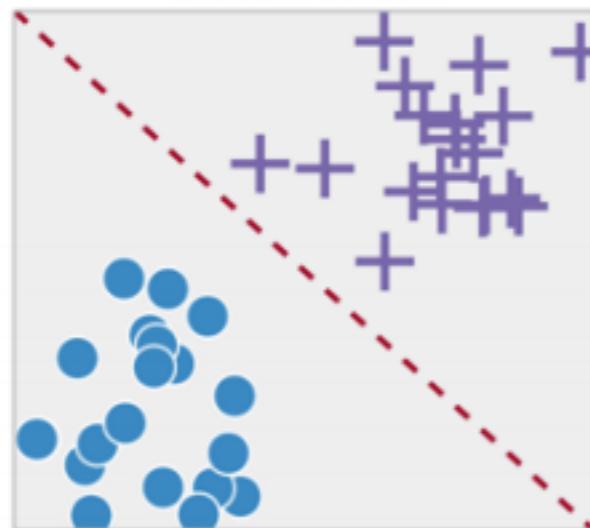


Classificação e Regressão

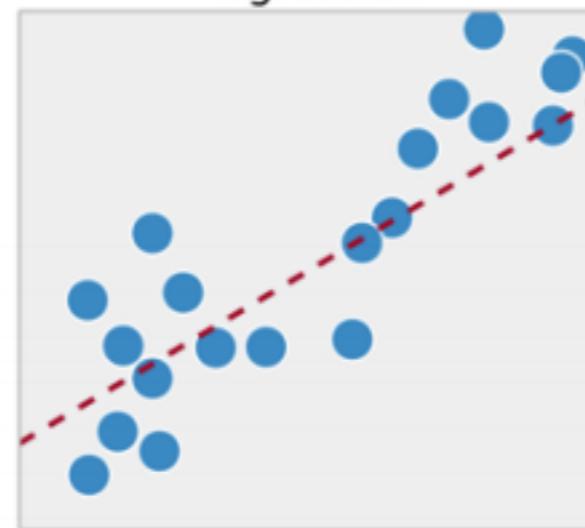
Identificação de uma área de interesse e o tipo de modelo



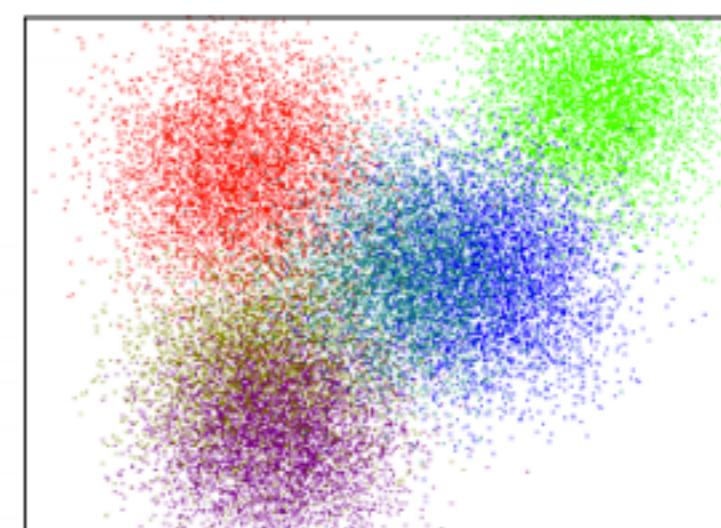
Classification

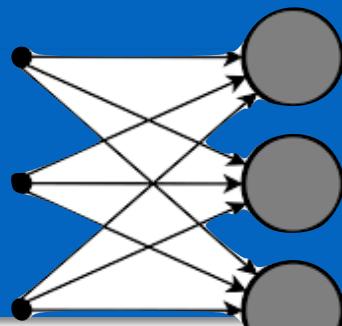


Regression

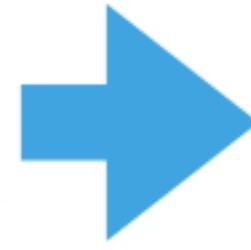


Clustering

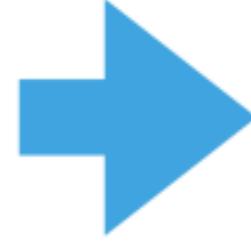




CLASSIFICAÇÃO



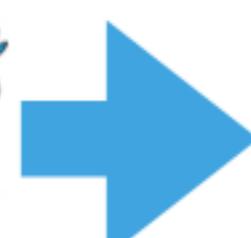
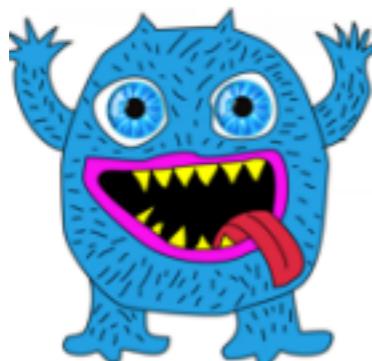
A



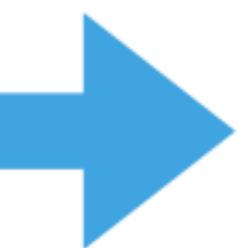
A



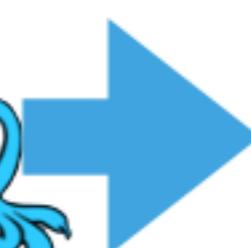
?



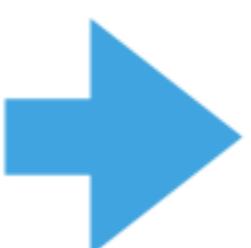
B



A

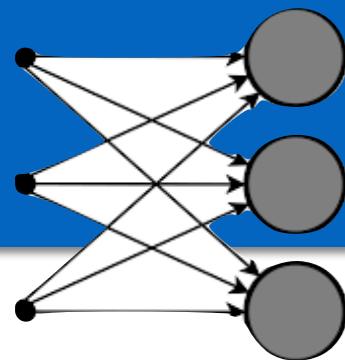


B



B

Classificação



Treinamento

DADOS DE TREINO VS DADOS DE TESTE

- ***Dados de Treino***

- Usados para treinar um modelo
- Exemplos



....

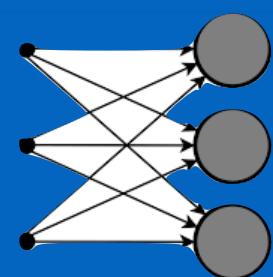


....

- ***Dados de Teste***

- Usados para testar a performance do modelo
- Dados de validação.

e.g. facial gender classification

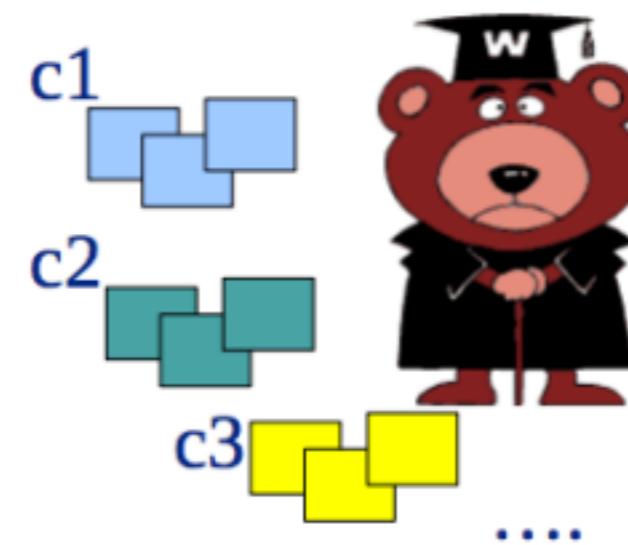


Aprendizado Supervisionado vs Não Supervisionado

APRENDIZADO SUPERVISIONADO VS NÃO SUPERVISIONADO

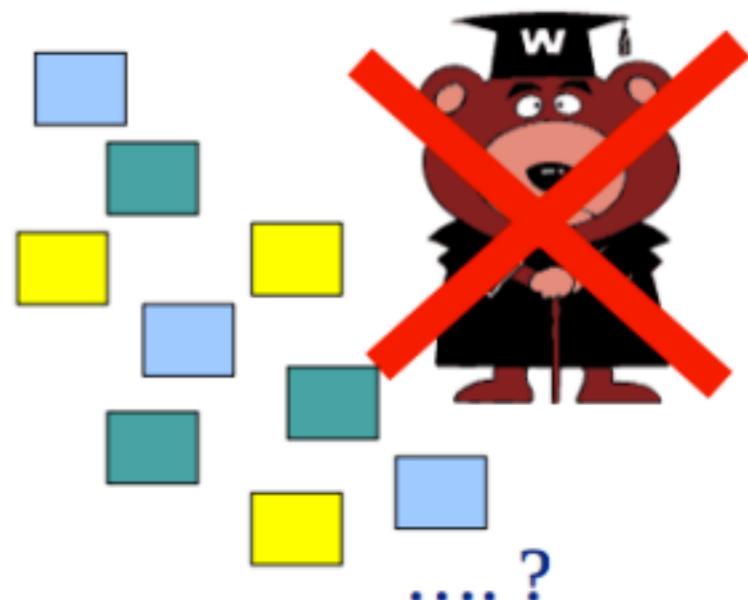
- ***Supervisionado***

- Conhecimento das entradas e saídas de dados
- Os dados possuem um label
- O objetivo é predizer a classe ou o label do dado

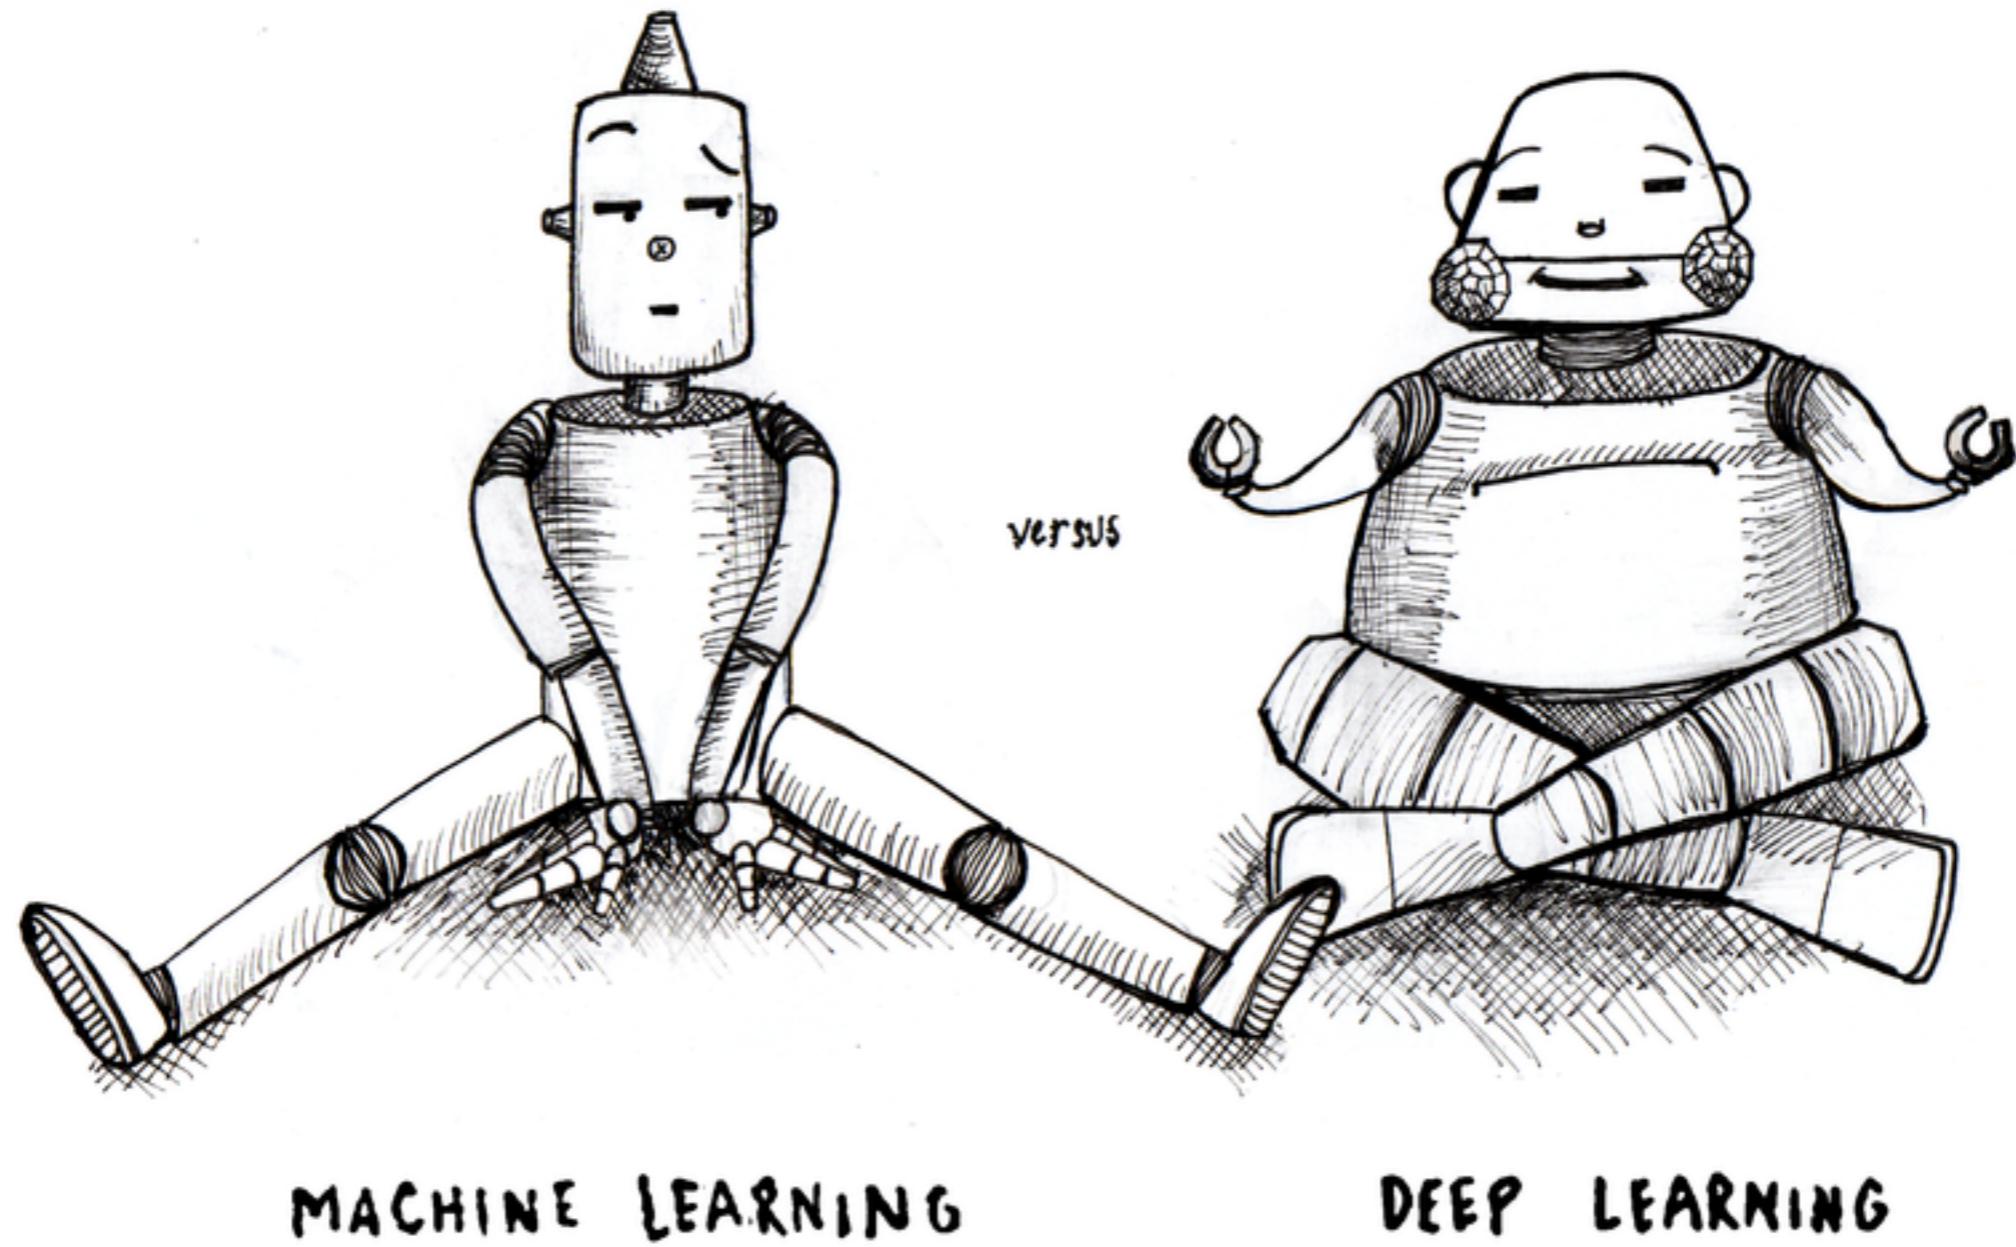


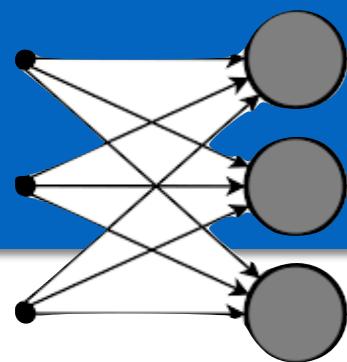
- ***Não Supervisionado***

- Sem conhecimento prévio dos dados
- O objetivo é determinar padrões



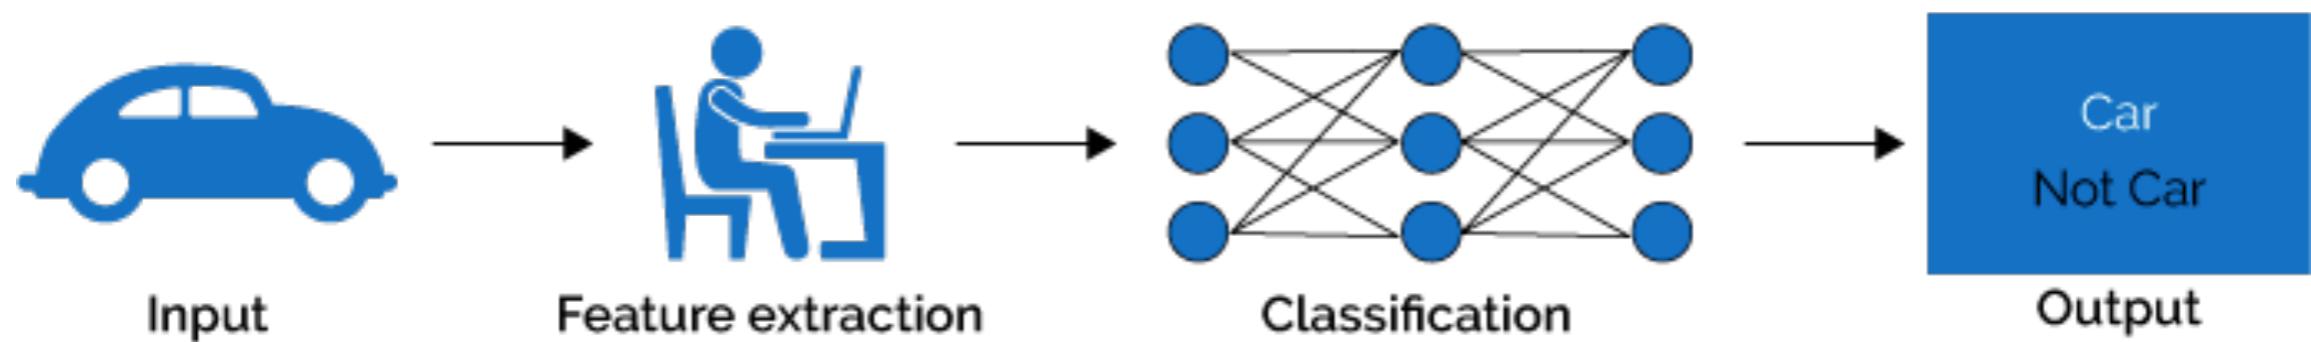
How can developments in deep learning make for a better approach to value investing?



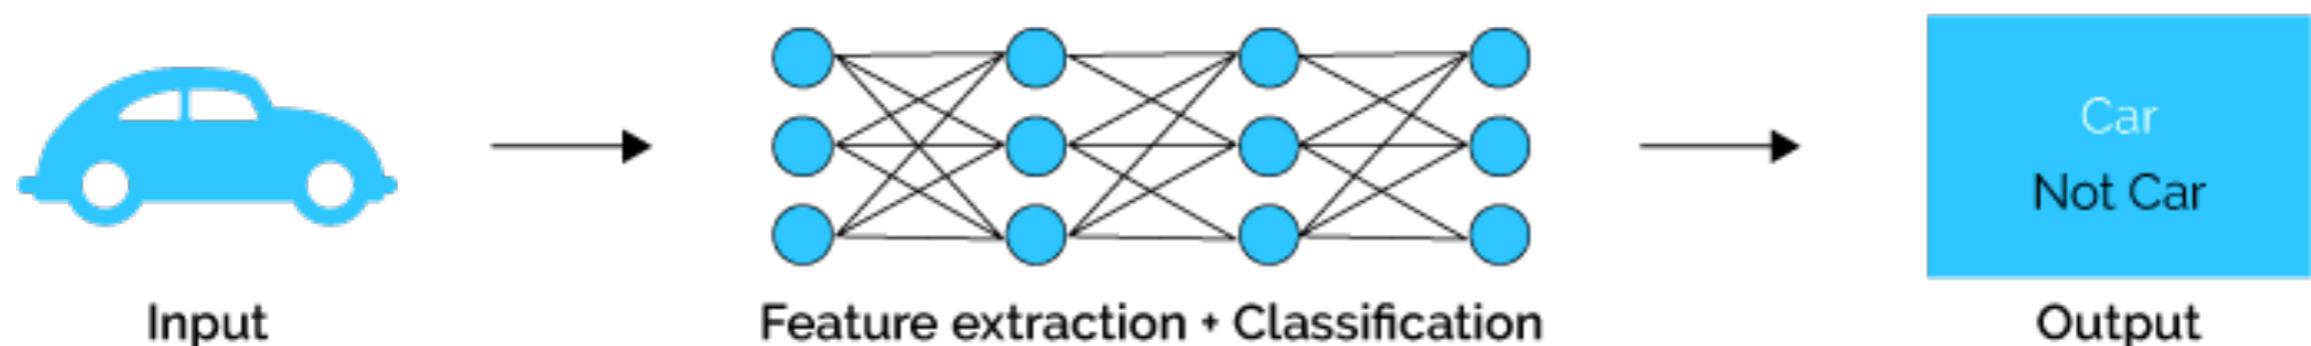


O que é Deep Learning?

Machine Learning



Deep Learning

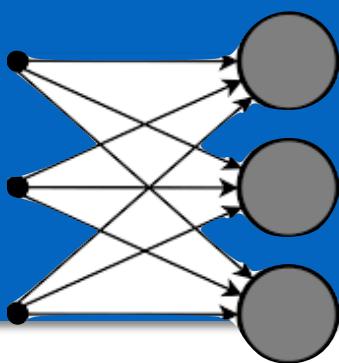


Introdução

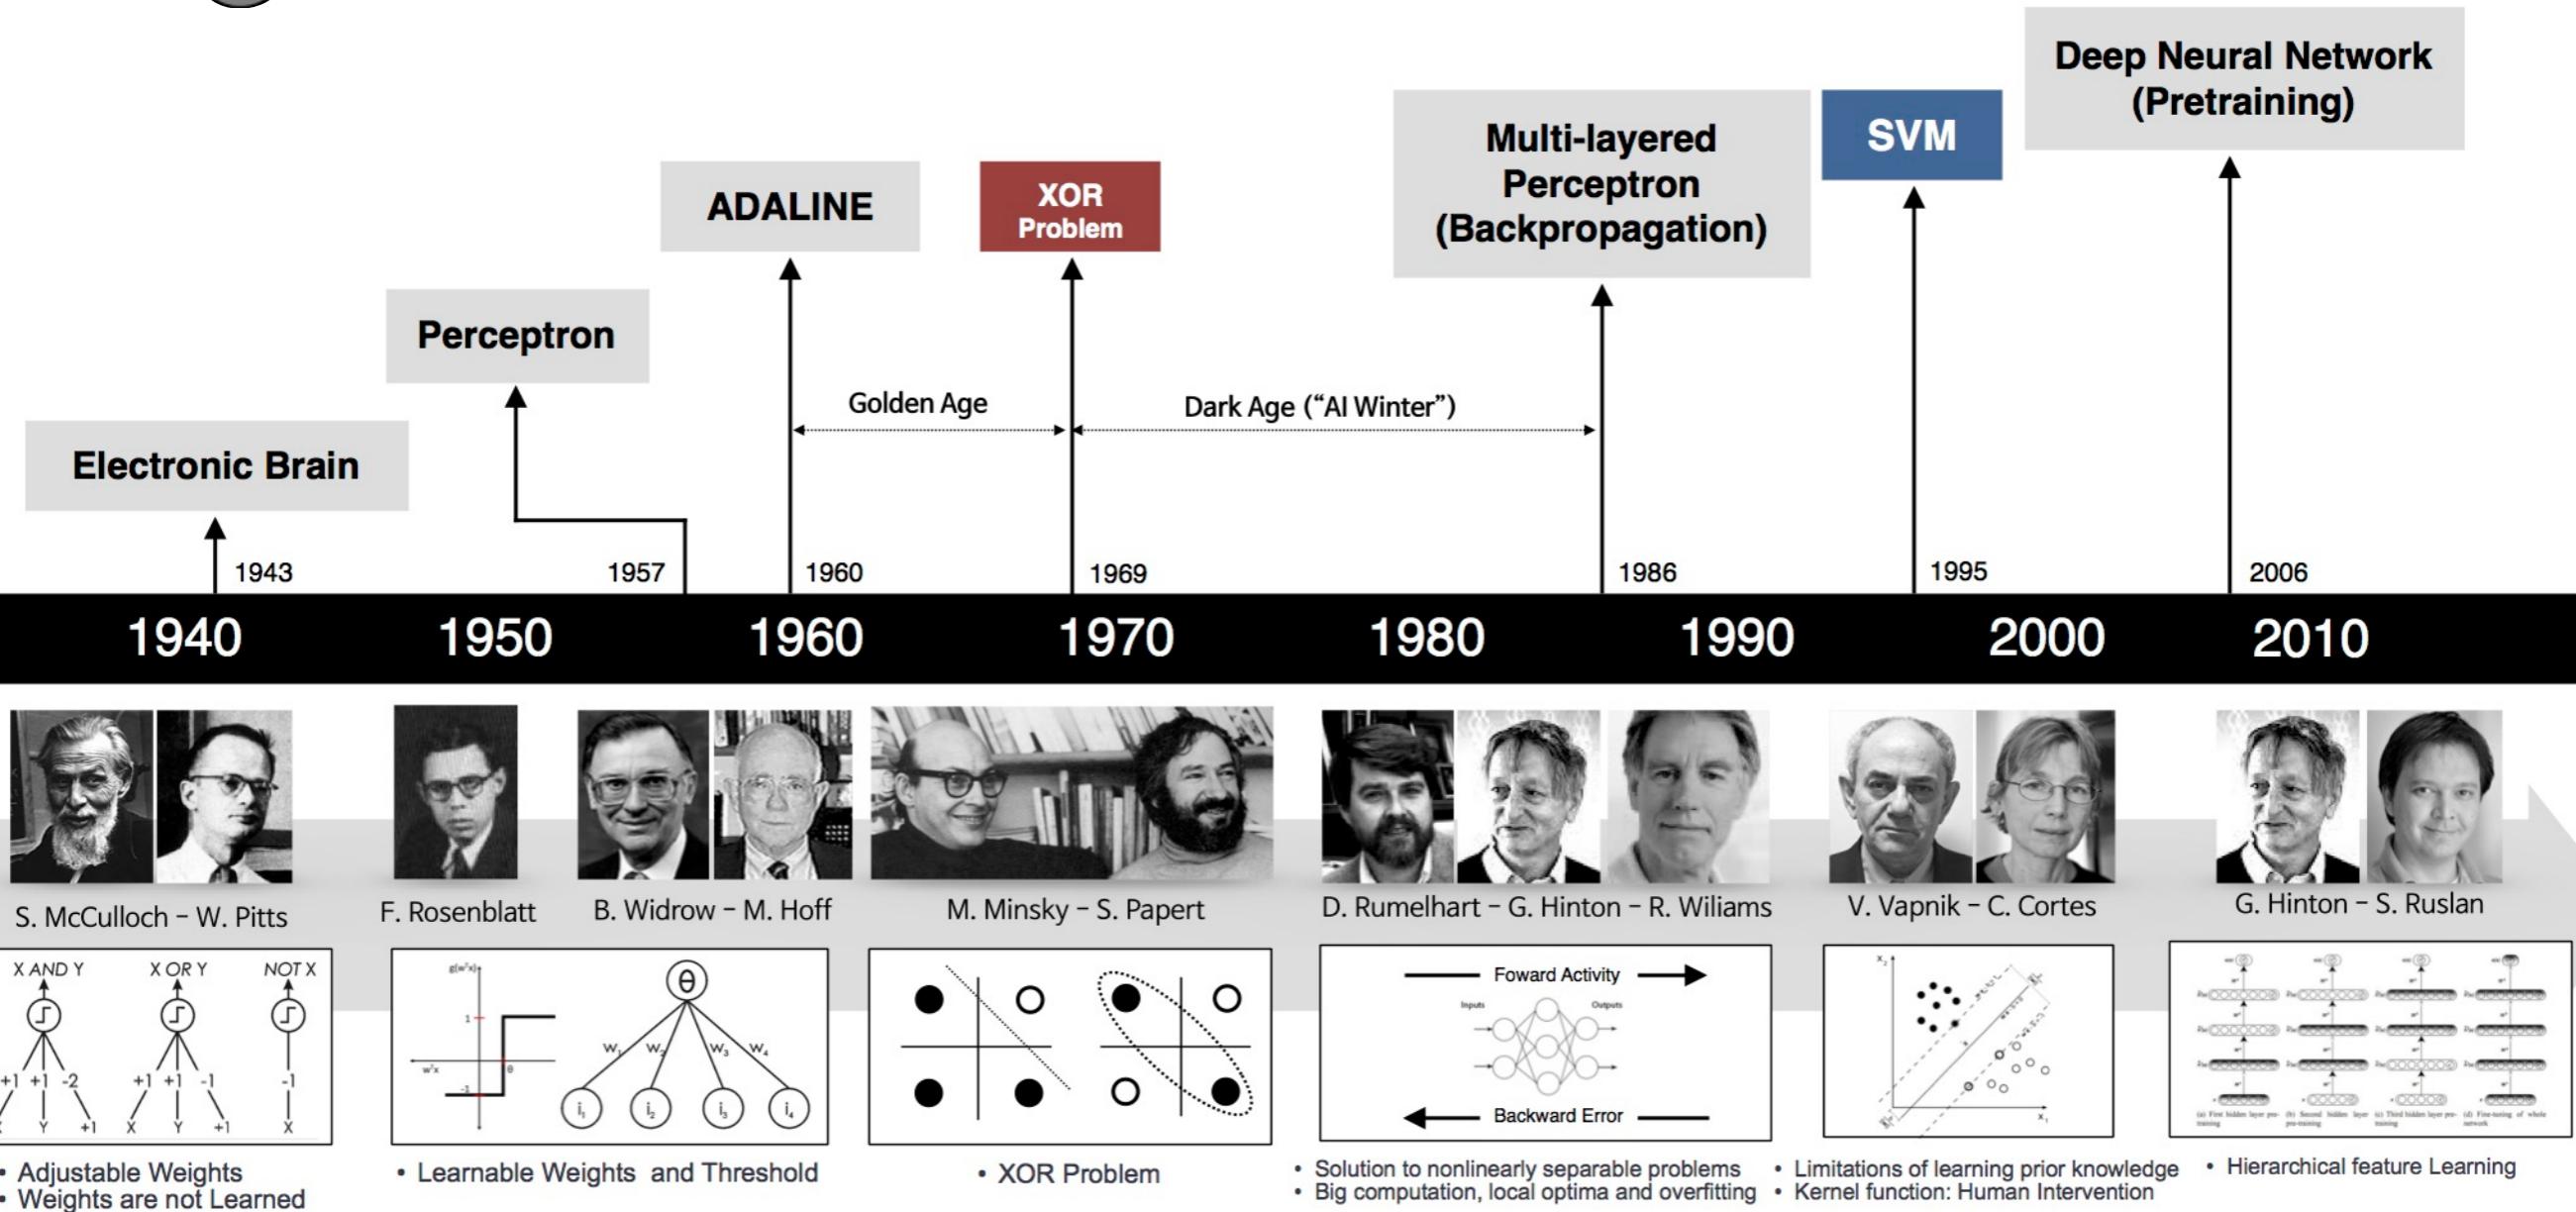
CNN

RNN

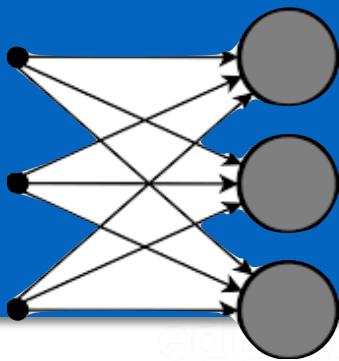
GAN



Histórico



Introdução



CNN

ARTIFICIAL INTELLIGENCE

Engineering of making Intelligent
Machines and Programs

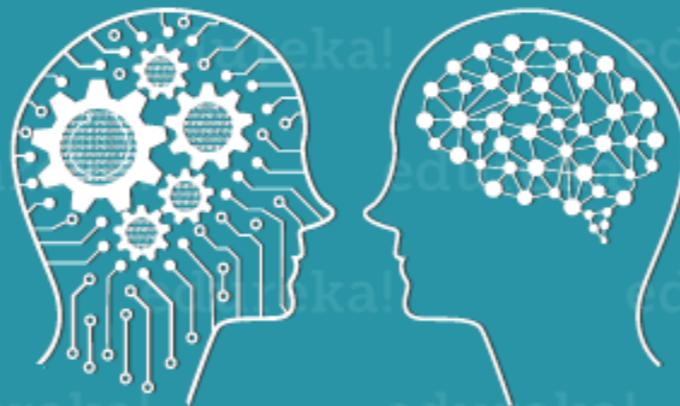


RNN

Histórico

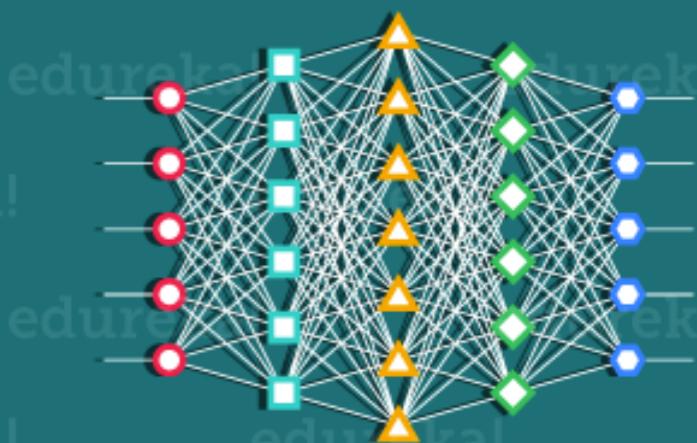
MACHINE LEARNING

Ability to learn without being
explicitly programmed



DEEP LEARNING

Learning based on Deep
Neural Network



1950's

1960's

1970's

1980's

1990's

2000's

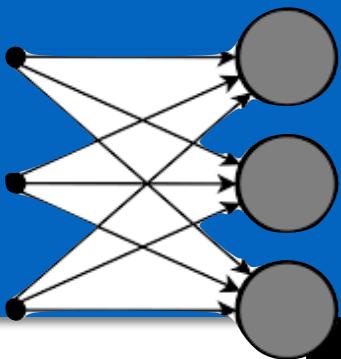
2006's

2010's

2012's

2017's

Introdução



CNN

RNN

GAN

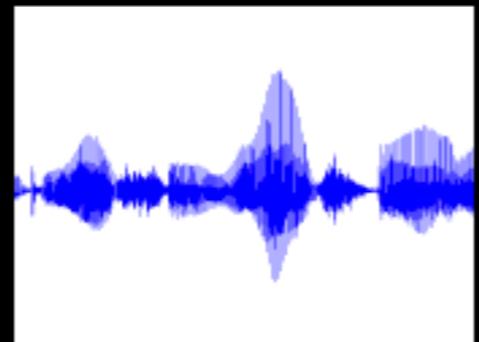
Onde utilizar Deep Learning

Images



Label image

Audio



Speech recognition

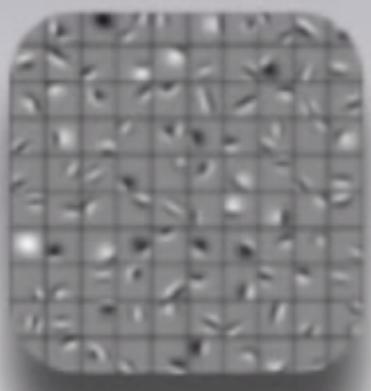
Text



Web search

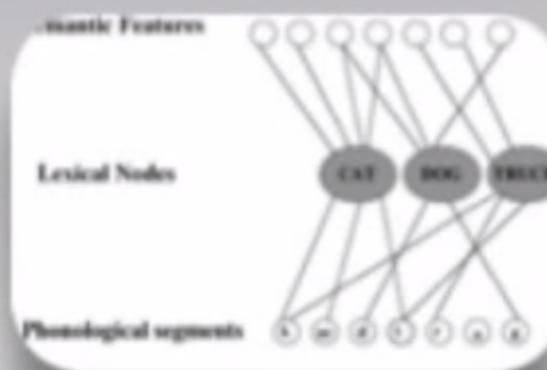
Application areas

Images



Jennifer Aniston

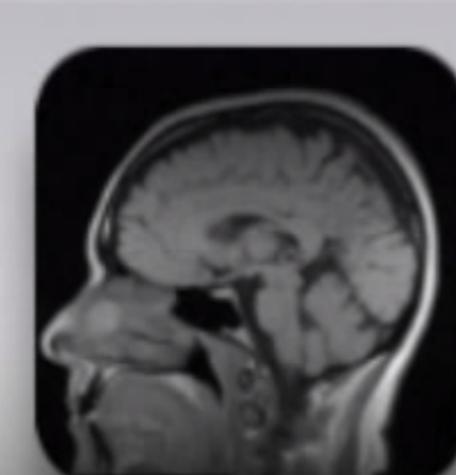
Speech



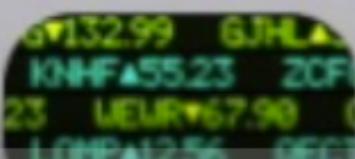
"This is what
I'm saying"



Network security
risks



Medical
diagnostics



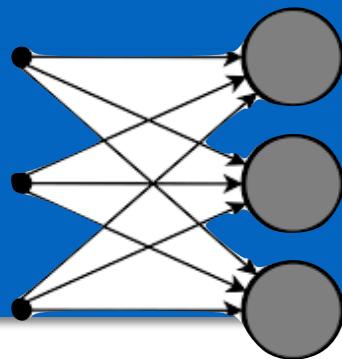
Trading signals

Introdução

CNN

RNN

GAN



Ferramentas



TensorFlow é uma biblioteca de software de código aberto para aprendizagem em máquina em uma variedade de tarefas. É um sistema para construir e treinar redes neurais para detectar e decifrar padrões e correlações, análogo ao (e não o mesmo) aprendizado e raciocínio humano



K Keras
theano

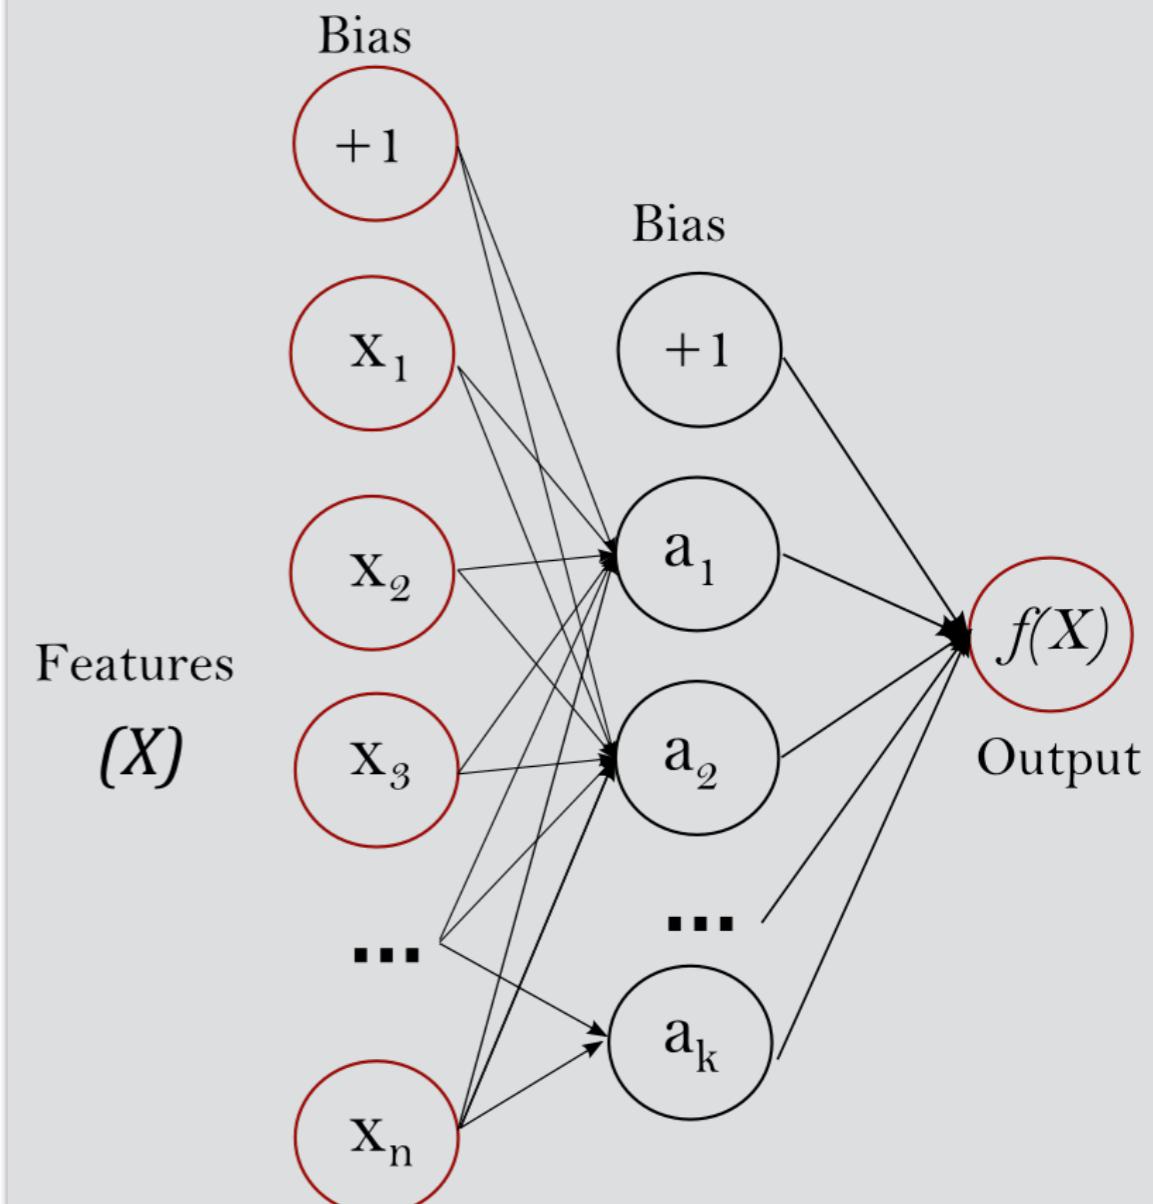
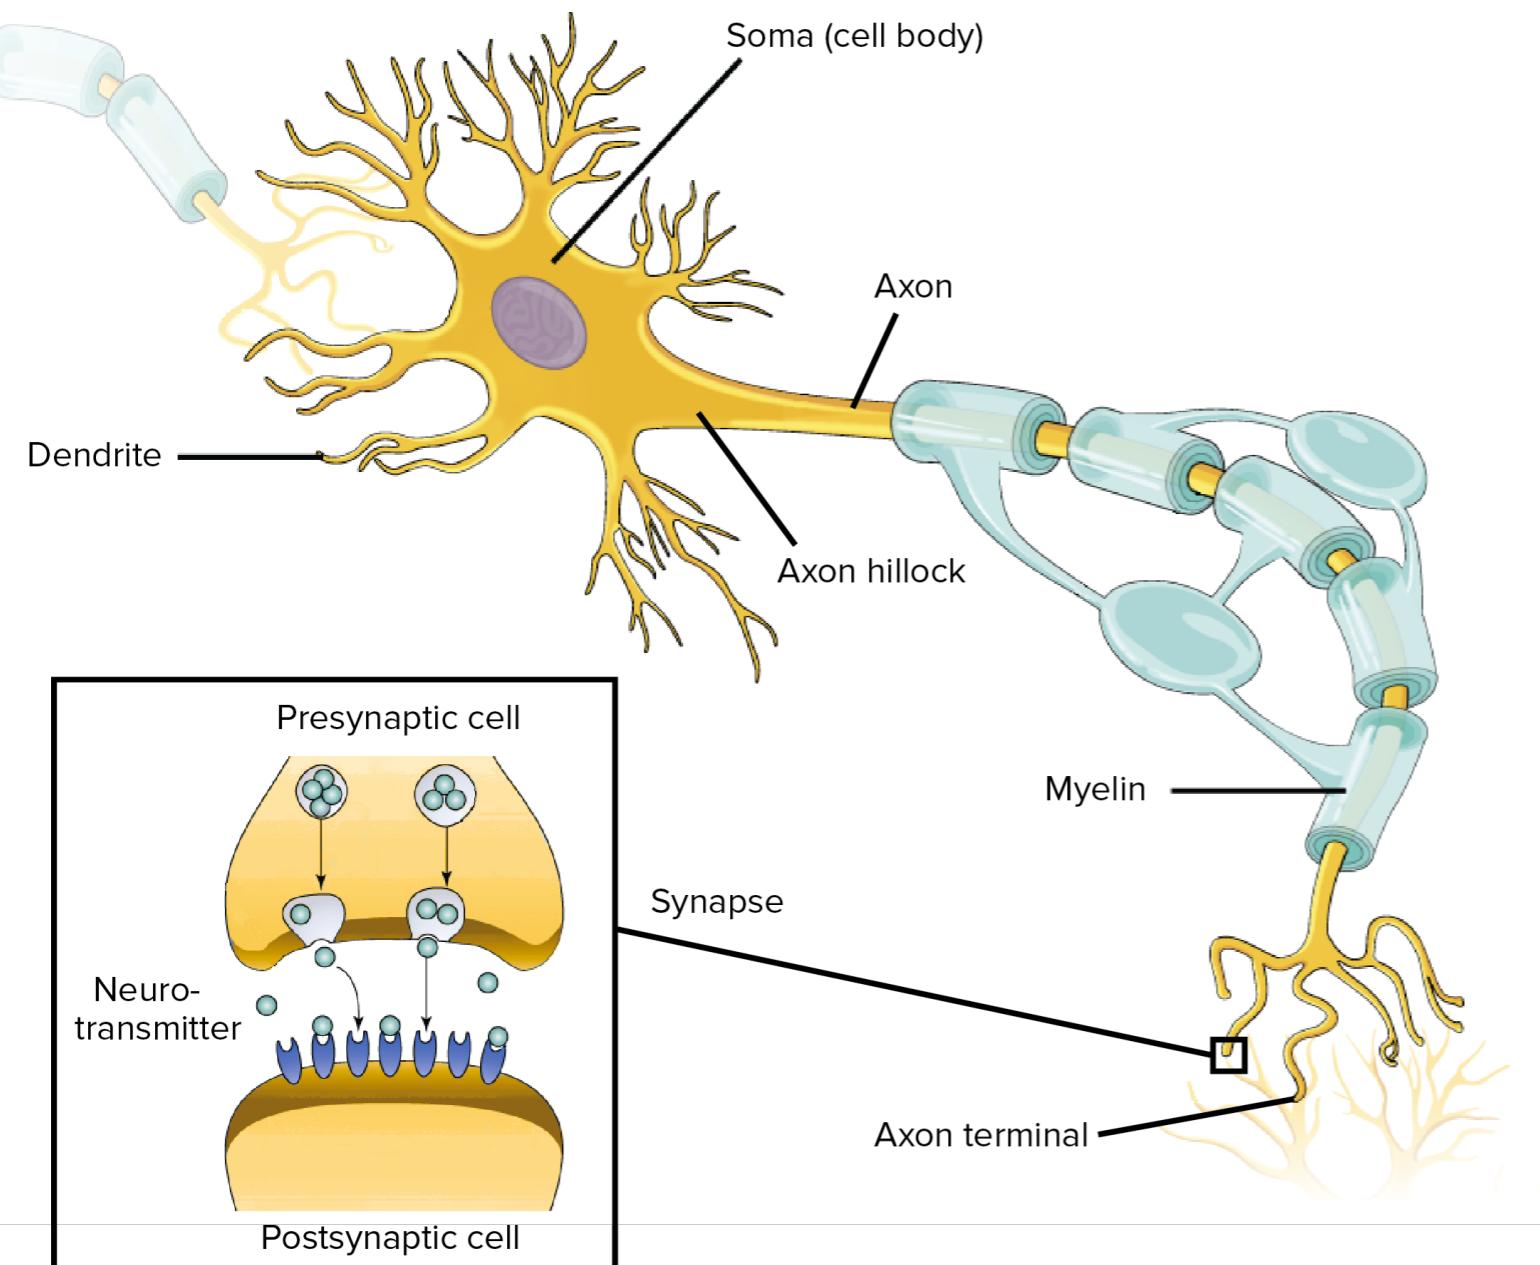
Introdução

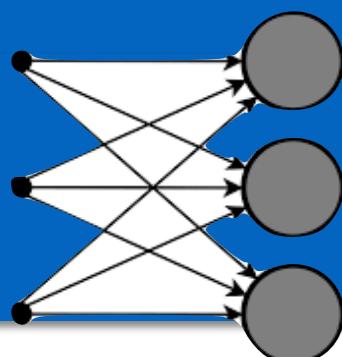
CNN

RNN

GAN

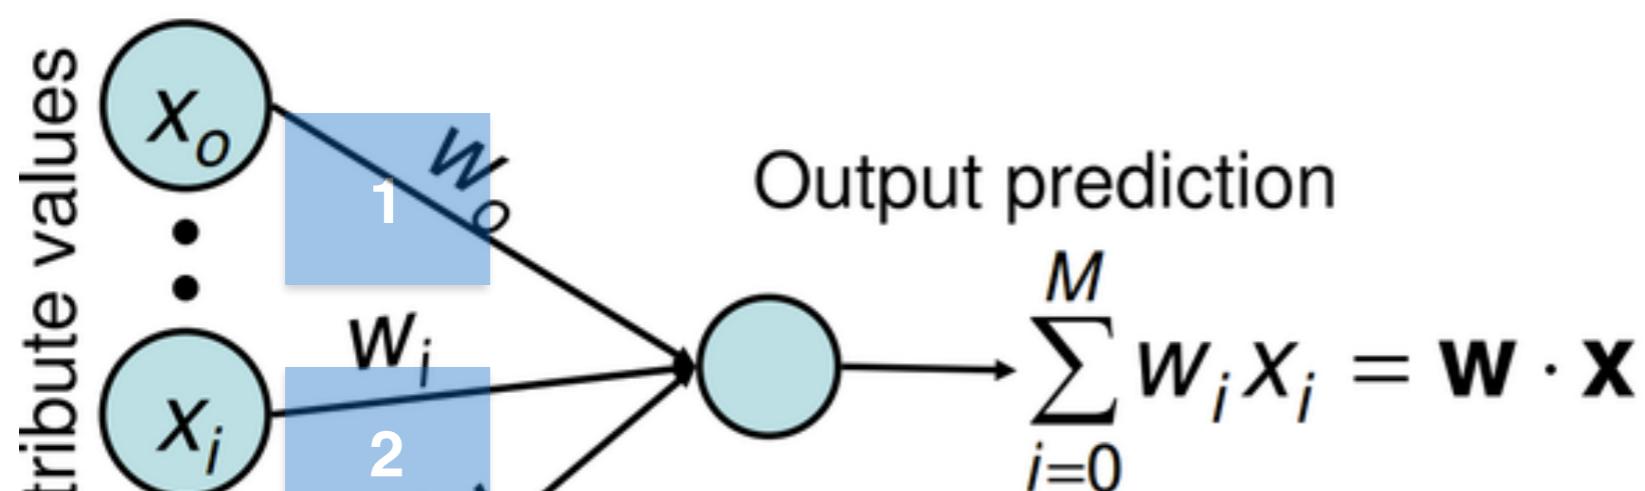
Abstração Biológica





Perceptrons

Neural Network: *Linear Perceptron*



Expected = 4
Predicted = 5

Erro = predicted value -
expected value

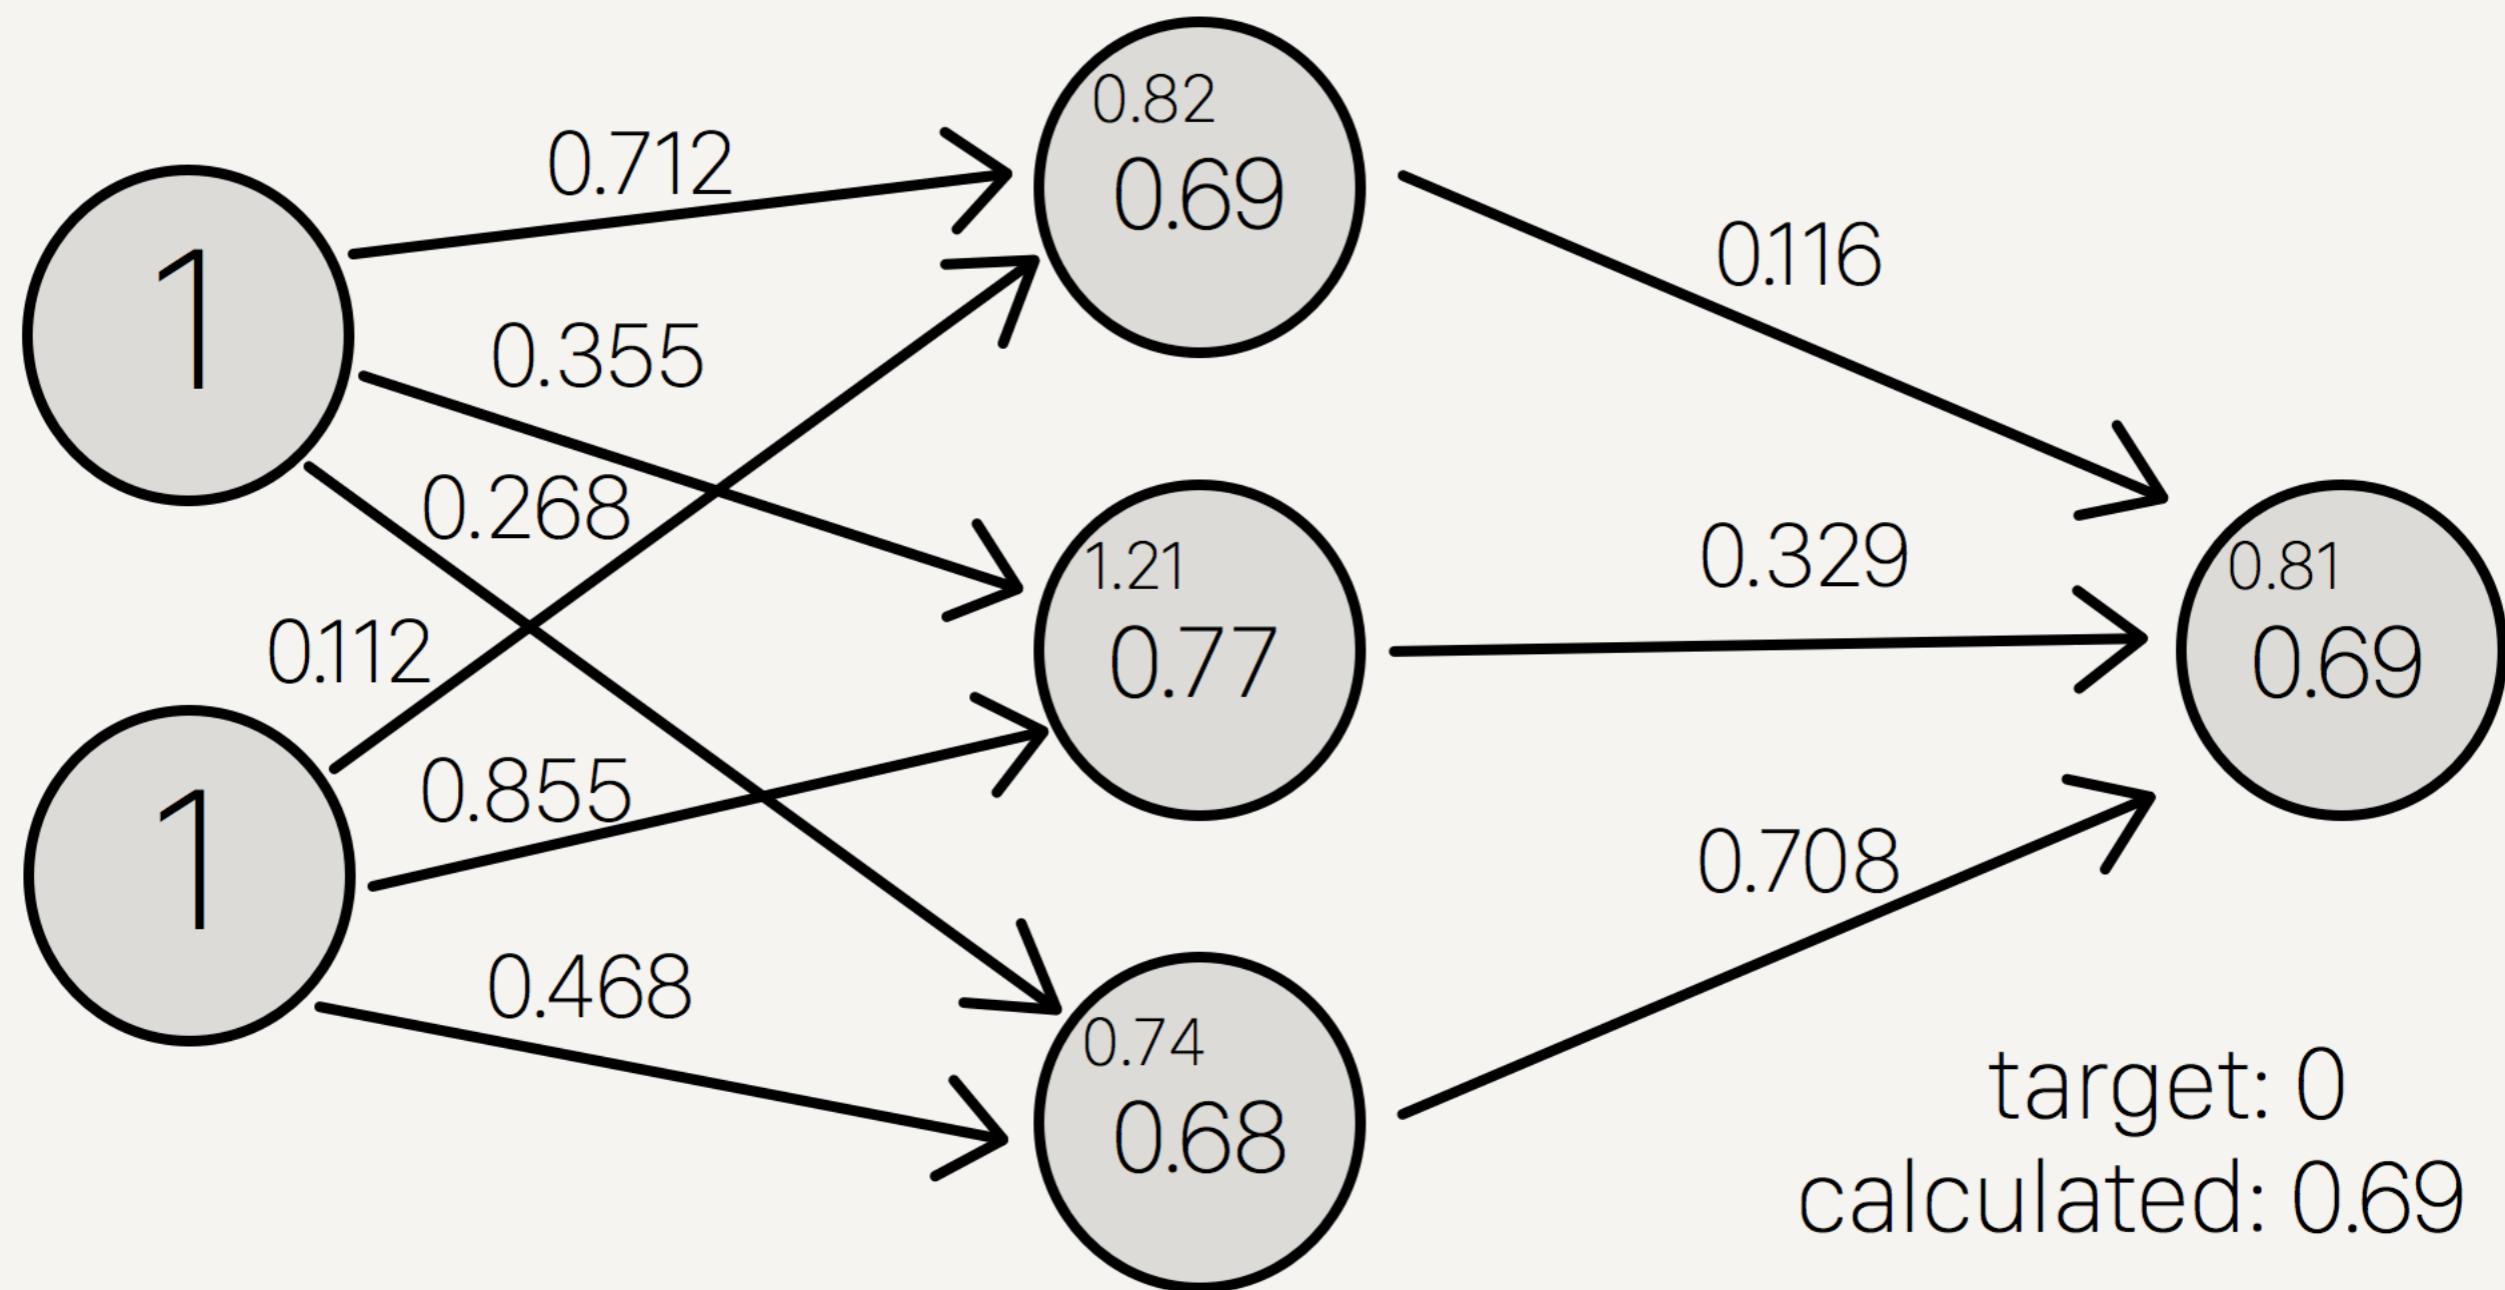
O perceptron é um tipo de rede neural artificial inventada em 1957 no Cornell Aeronautical Laboratory por Frank Rosenblatt. Ele pode ser visto como o tipo mais simples de rede neural feedforward: um classificador linear.



INPUT

HIDDEN

OUTPUT

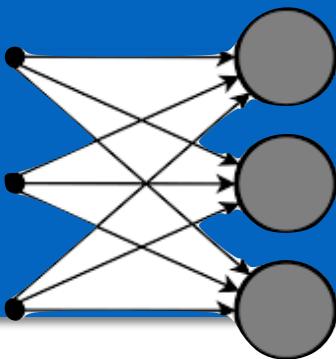


Introdução

CNN

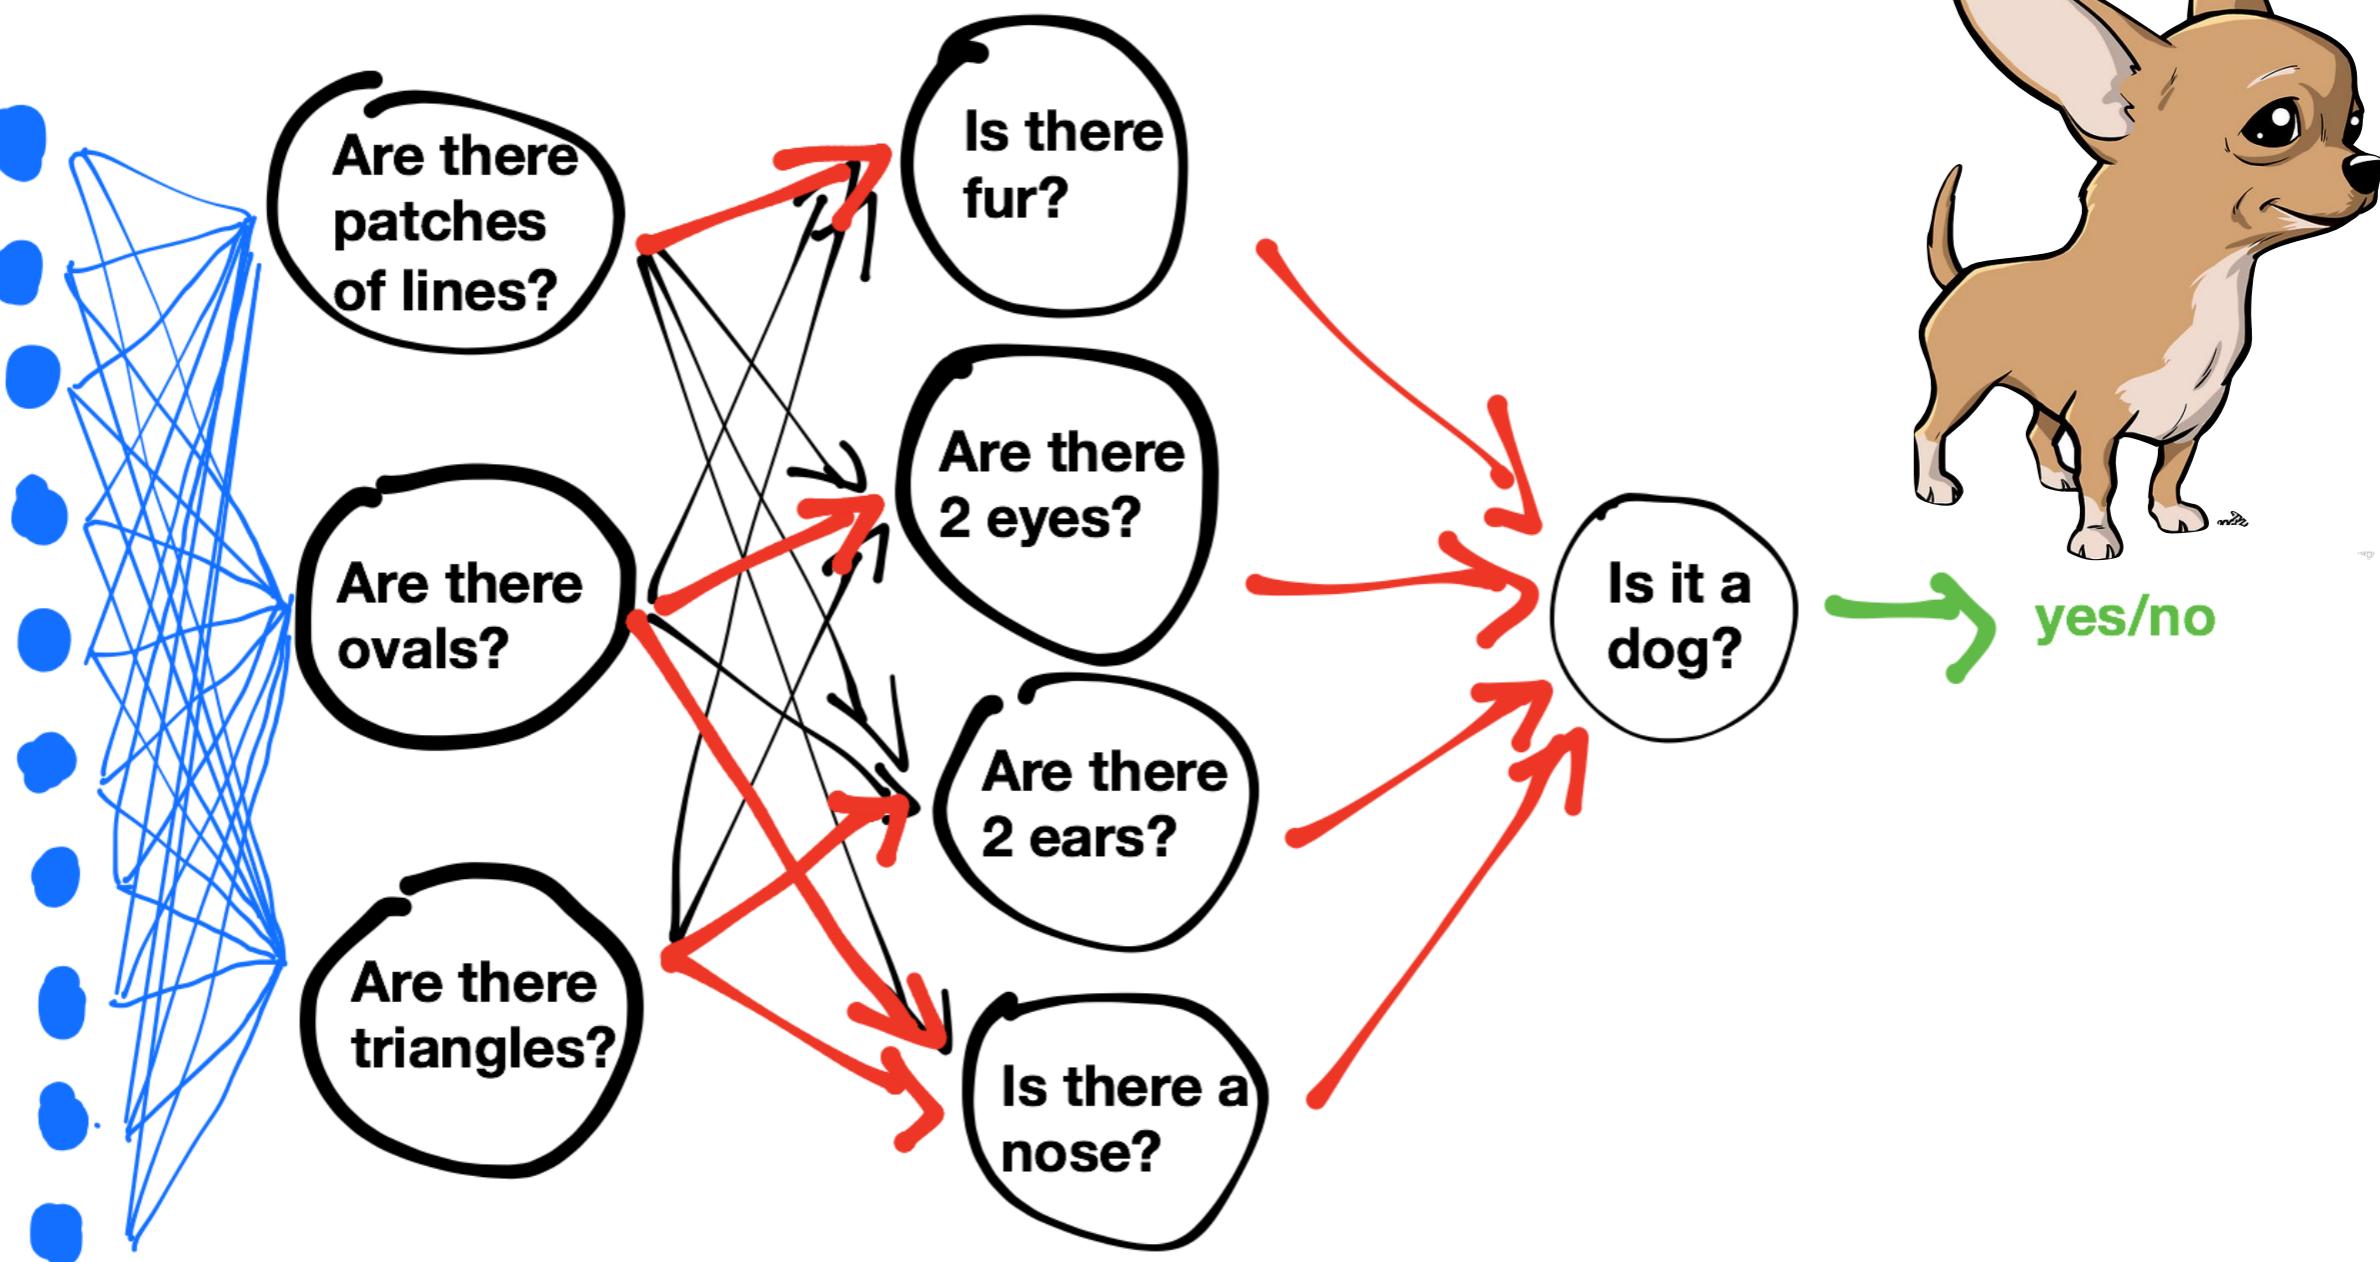
RNN

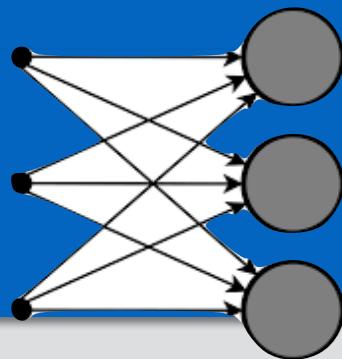
GAN



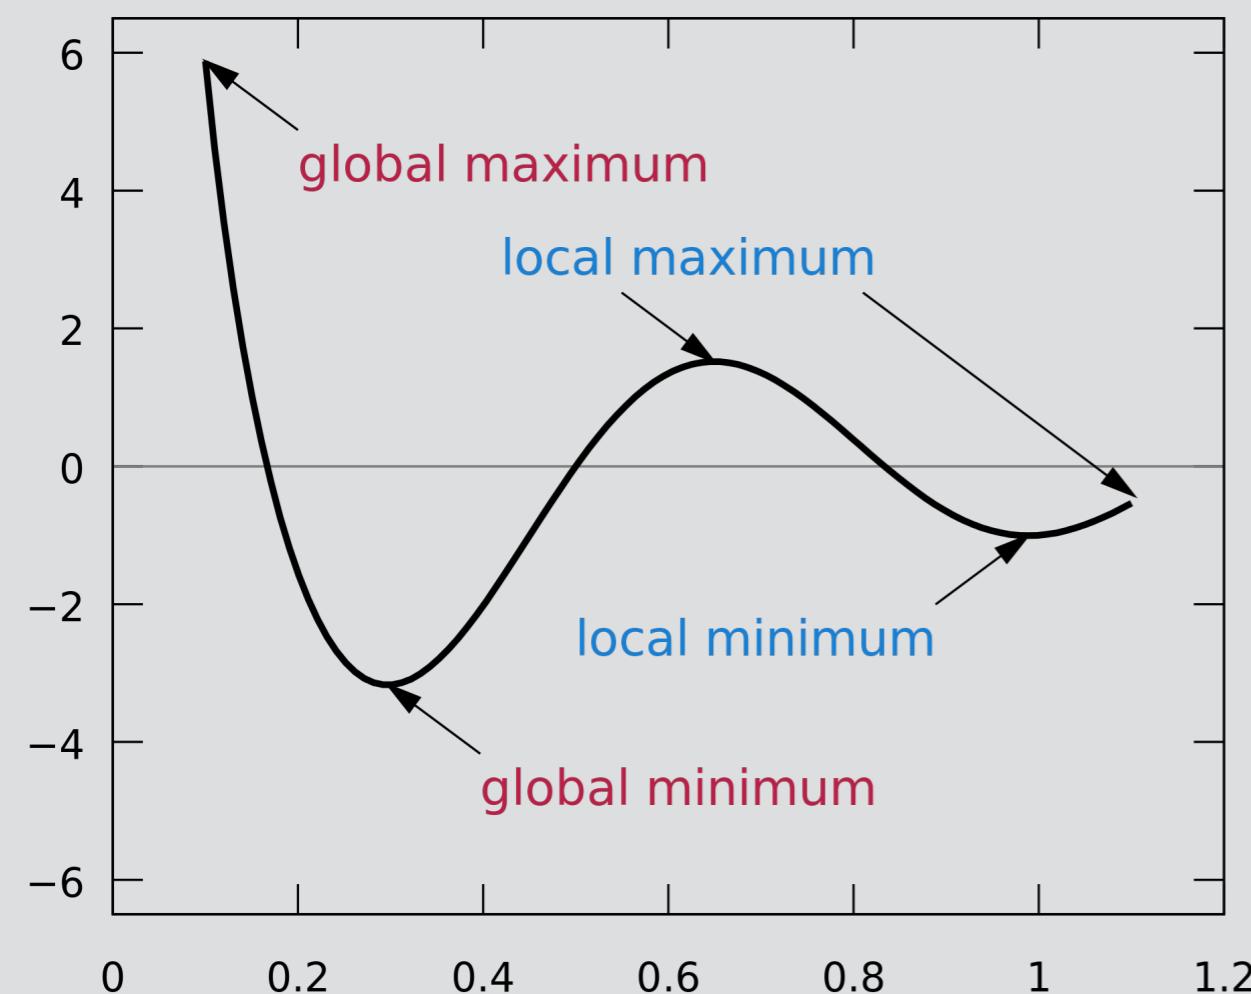
input: every single pixel of the image

Backpropagation

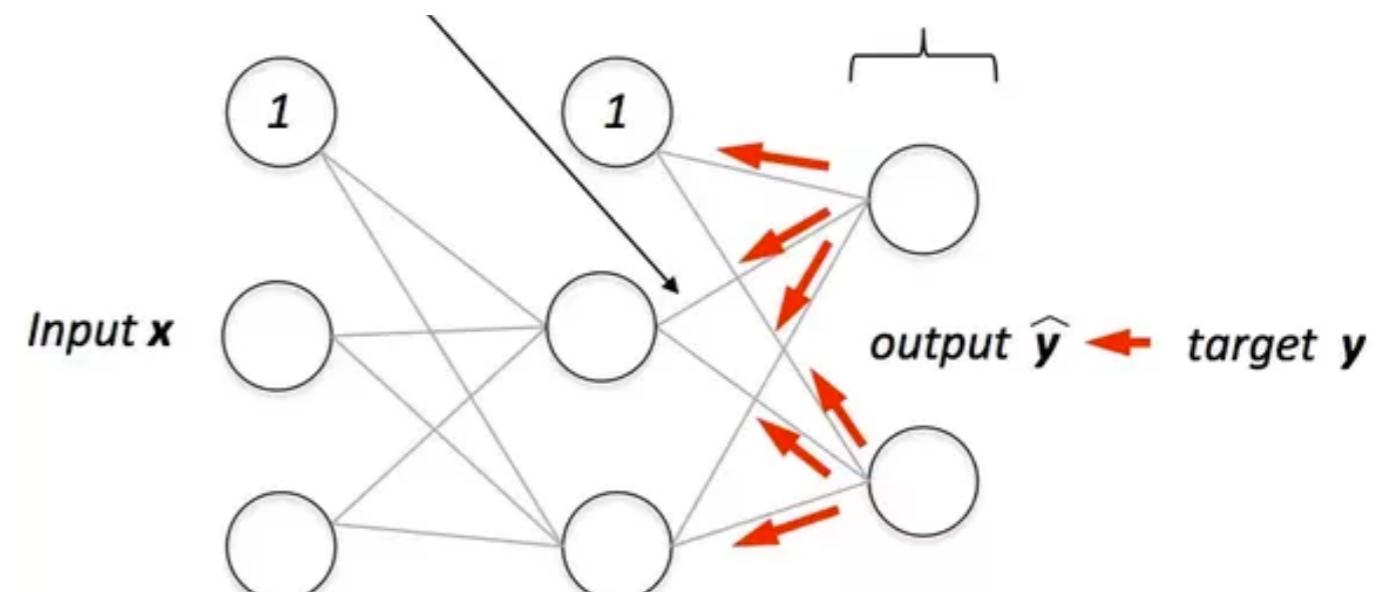


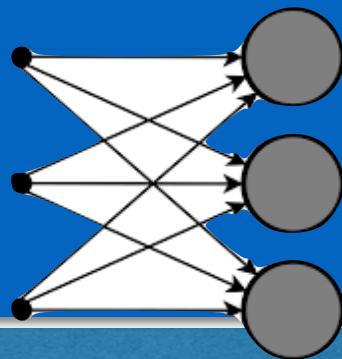


Backpropagation e Gradiente Descendente



**peso (w) = peso (w) +
taxa_aprendizado *
(esperado - predito) * x**





Função de Ativação

Nas redes neurais biologicamente inspirados, a função de ativação é geralmente uma abstração que representa a taxa de potencial de ação na célula.

$$x_1w_1 + x_2w_2 = (0.6 \times 0.5) + (1 \times 0.8) = 1.1$$

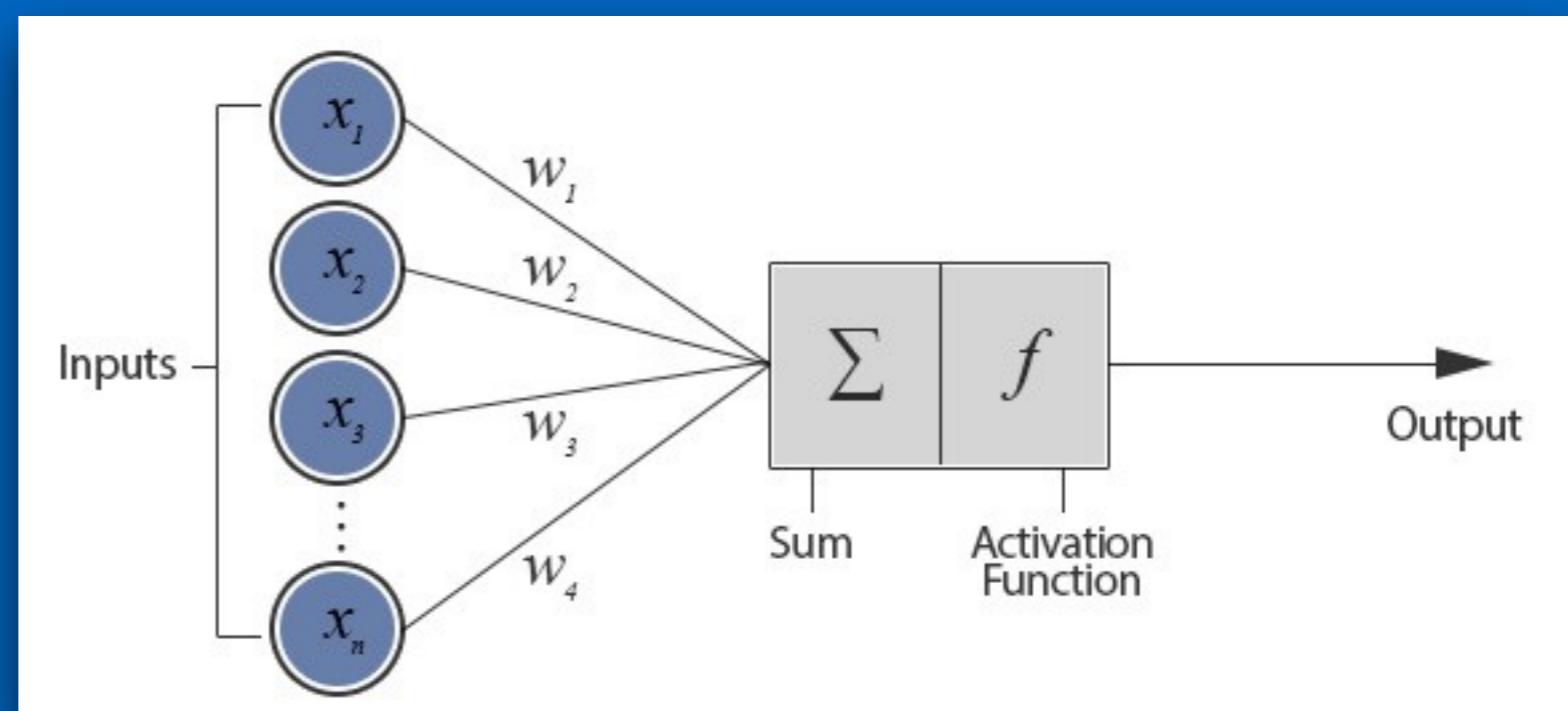
Input 1 (x_1) = 0.6

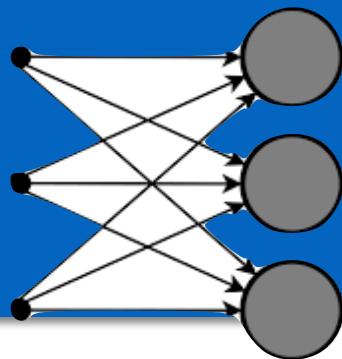
Input 2 (x_2) = 1.0

Weight 1 (w_1) = 0.5

Weight 2 (w_2) = 0.8

Threshold = 1.0



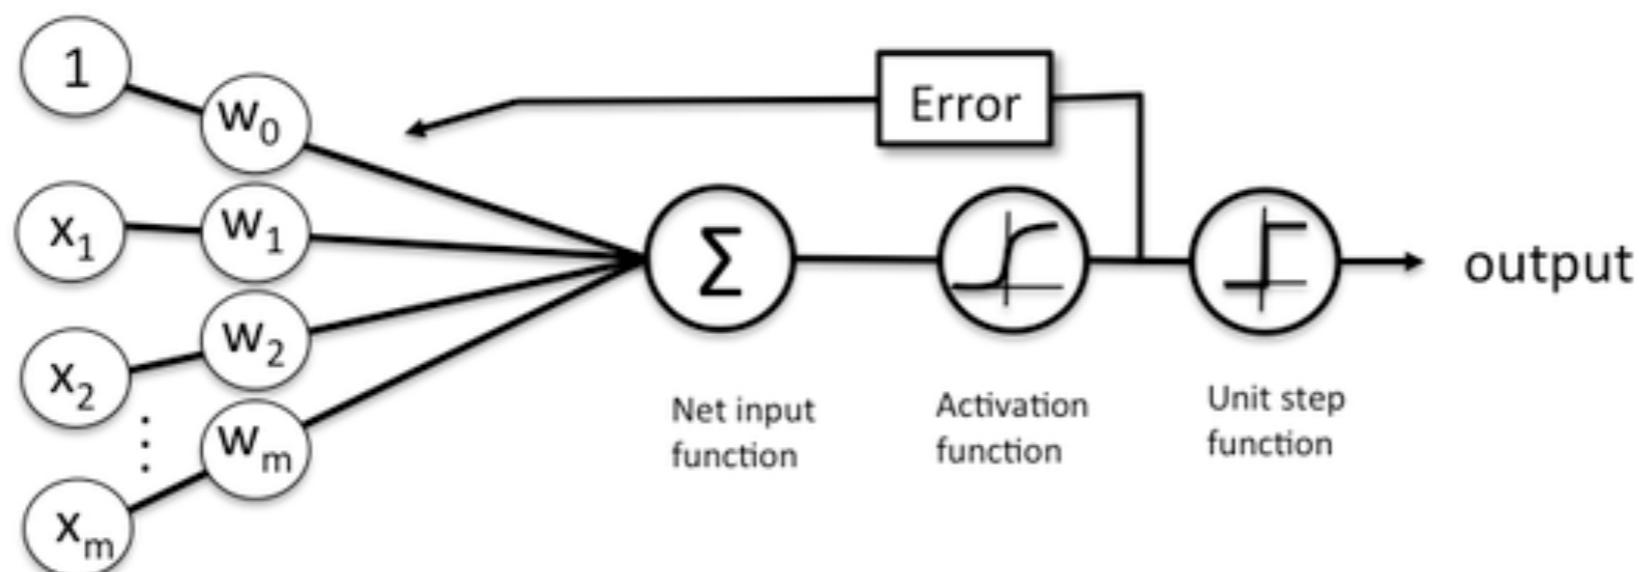


Função de Ativação

Uma função de activação serve como um limite,

$$\text{ativacao} = \text{sum}(\text{weight}_i * \text{x}_i) + \text{bias}$$

$$\text{predicao} = 1.0 \text{ if } \text{ativacao} \geq 0.0 \text{ else } 0.0$$



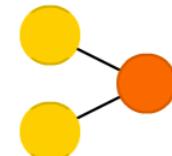
Schematic of a logistic regression classifier.

A mostly complete chart of
Neural Networks

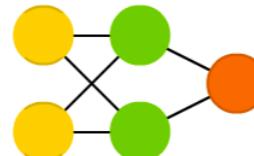
©2016 Fjodor van Veen - asimovinstitute.org

- (○) Backfed Input Cell
- (○) Input Cell
- (△) Noisy Input Cell
- (●) Hidden Cell
- (○) Probabilistic Hidden Cell
- (△) Spiking Hidden Cell
- (●) Output Cell
- (○) Match Input Output Cell
- (●) Recurrent Cell
- (○) Memory Cell
- (△) Different Memory Cell
- (●) Kernel
- (○) Convolution or Pool

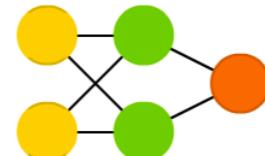
Perceptron (P)



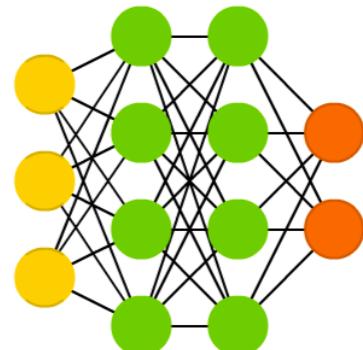
Feed Forward (FF)



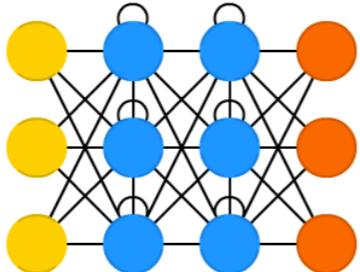
Radial Basis Network (RBF)



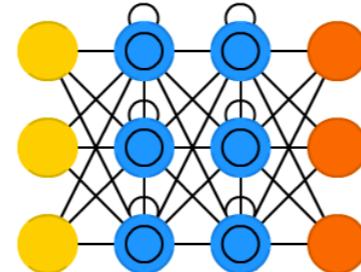
Deep Feed Forward (DFF)



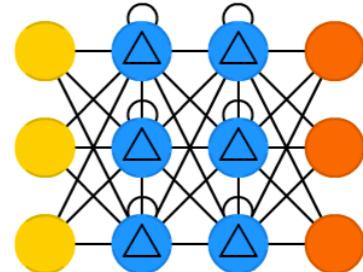
Recurrent Neural Network (RNN)



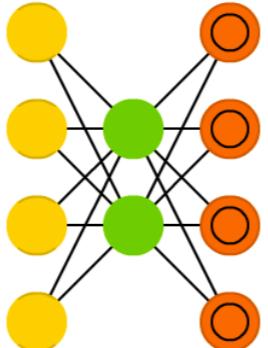
Long / Short Term Memory (LSTM)



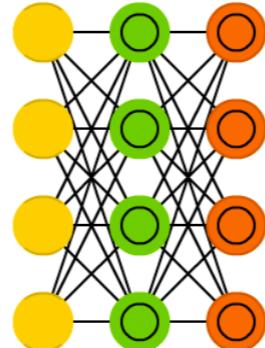
Gated Recurrent Unit (GRU)



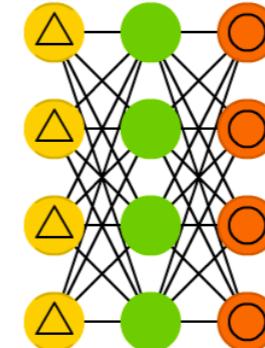
Auto Encoder (AE)



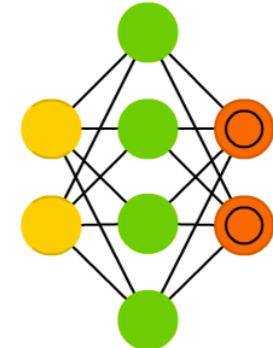
Variational AE (VAE)



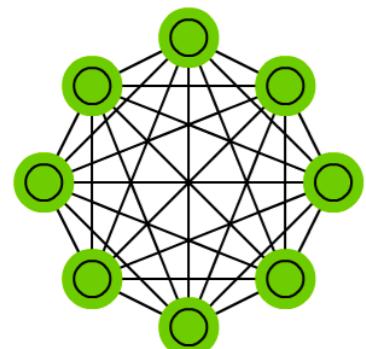
Denoising AE (DAE)



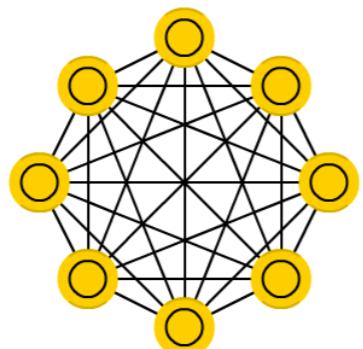
Sparse AE (SAE)



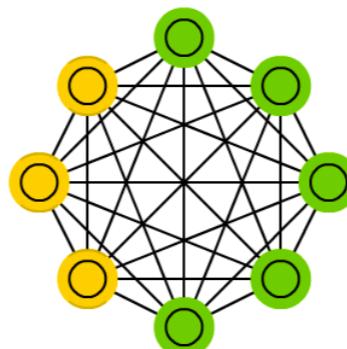
Markov Chain (MC)



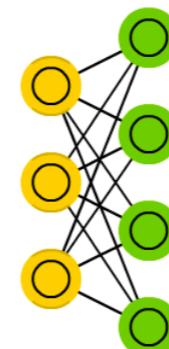
Hopfield Network (HN)



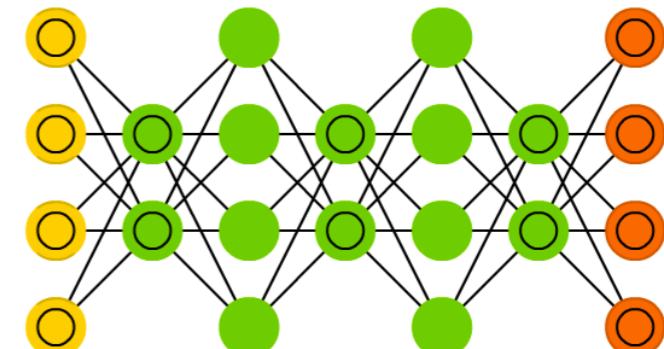
Boltzmann Machine (BM)



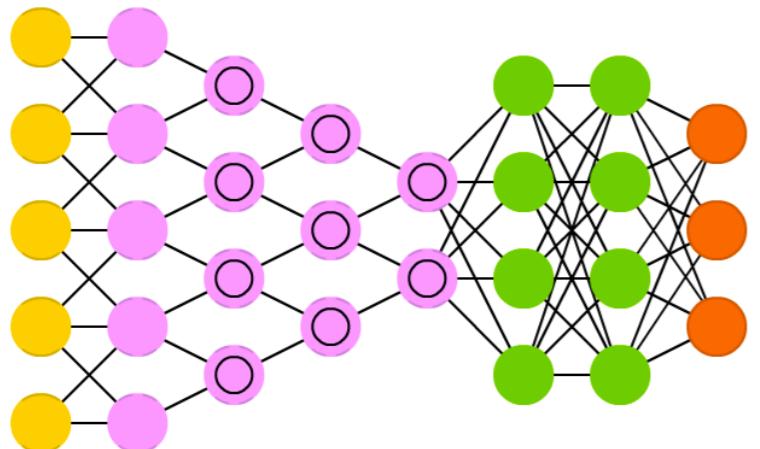
Restricted BM (RBM)



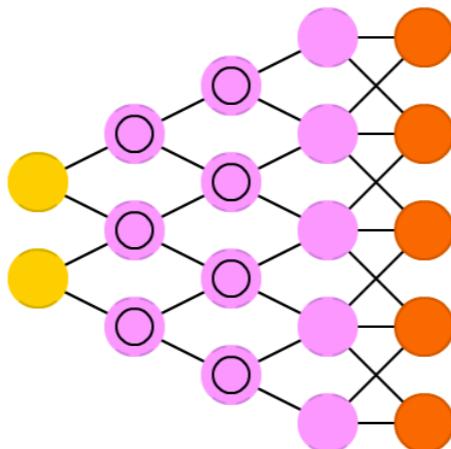
Deep Belief Network (DBN)



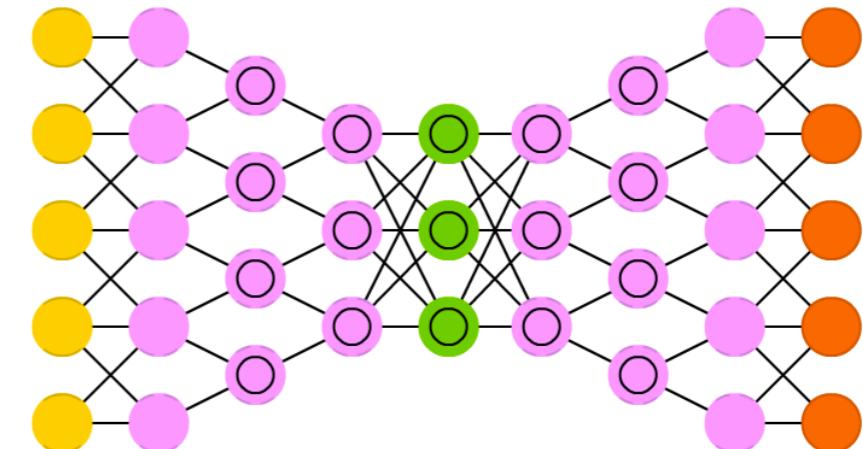
Deep Convolutional Network (DCN)



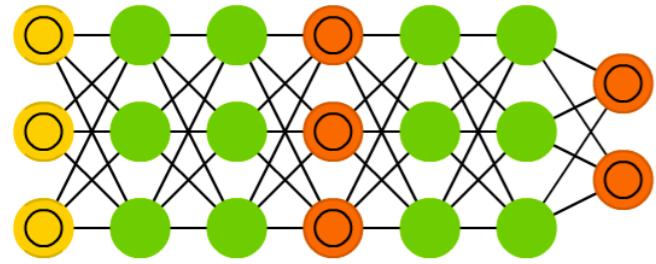
Deconvolutional Network (DN)



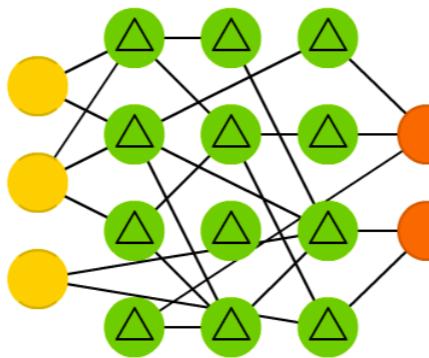
Deep Convolutional Inverse Graphics Network (DCIGN)



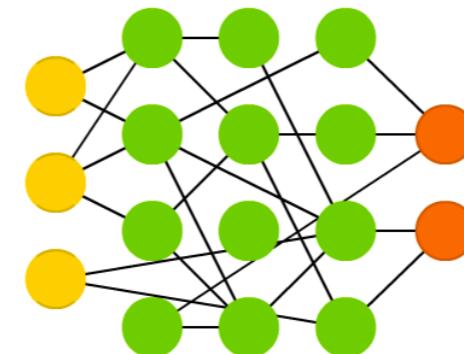
Generative Adversarial Network (GAN)



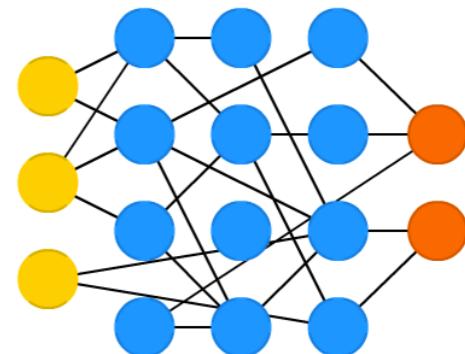
Liquid State Machine (LSM)



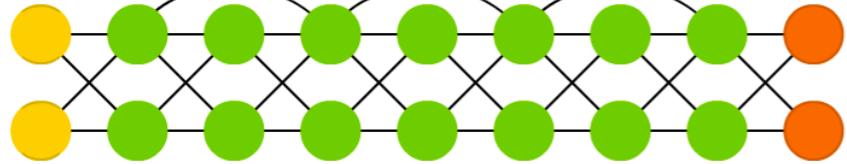
Extreme Learning Machine (ELM)



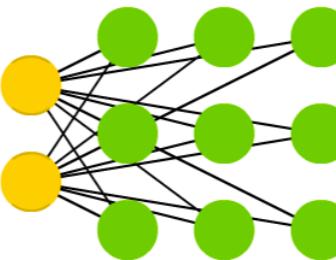
Echo State Network (ESN)



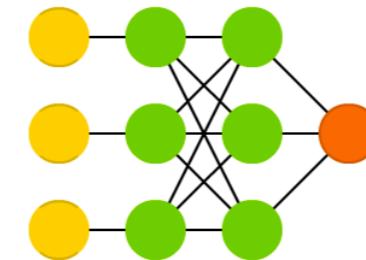
Deep Residual Network (DRN)



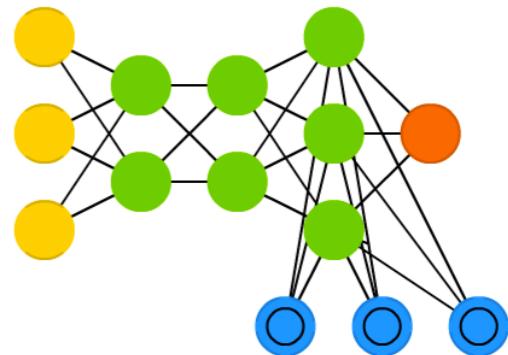
Kohonen Network (KN)



Support Vector Machine (SVM)



Neural Turing Machine (NTM)

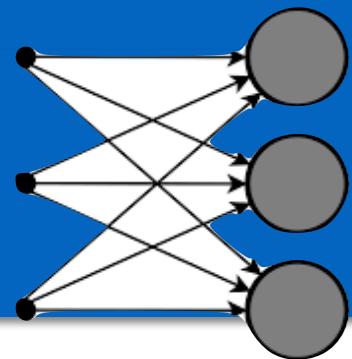


Introdução

CNN

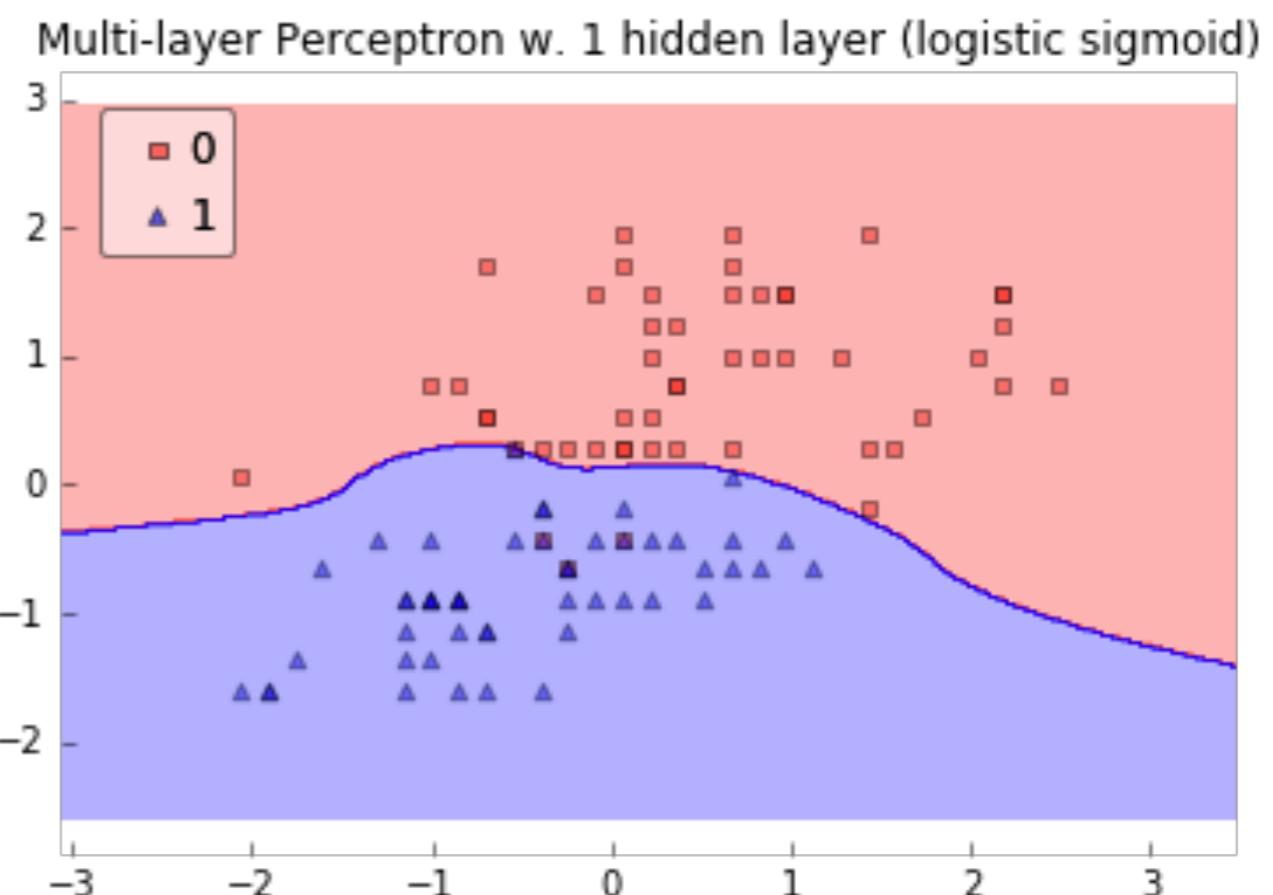
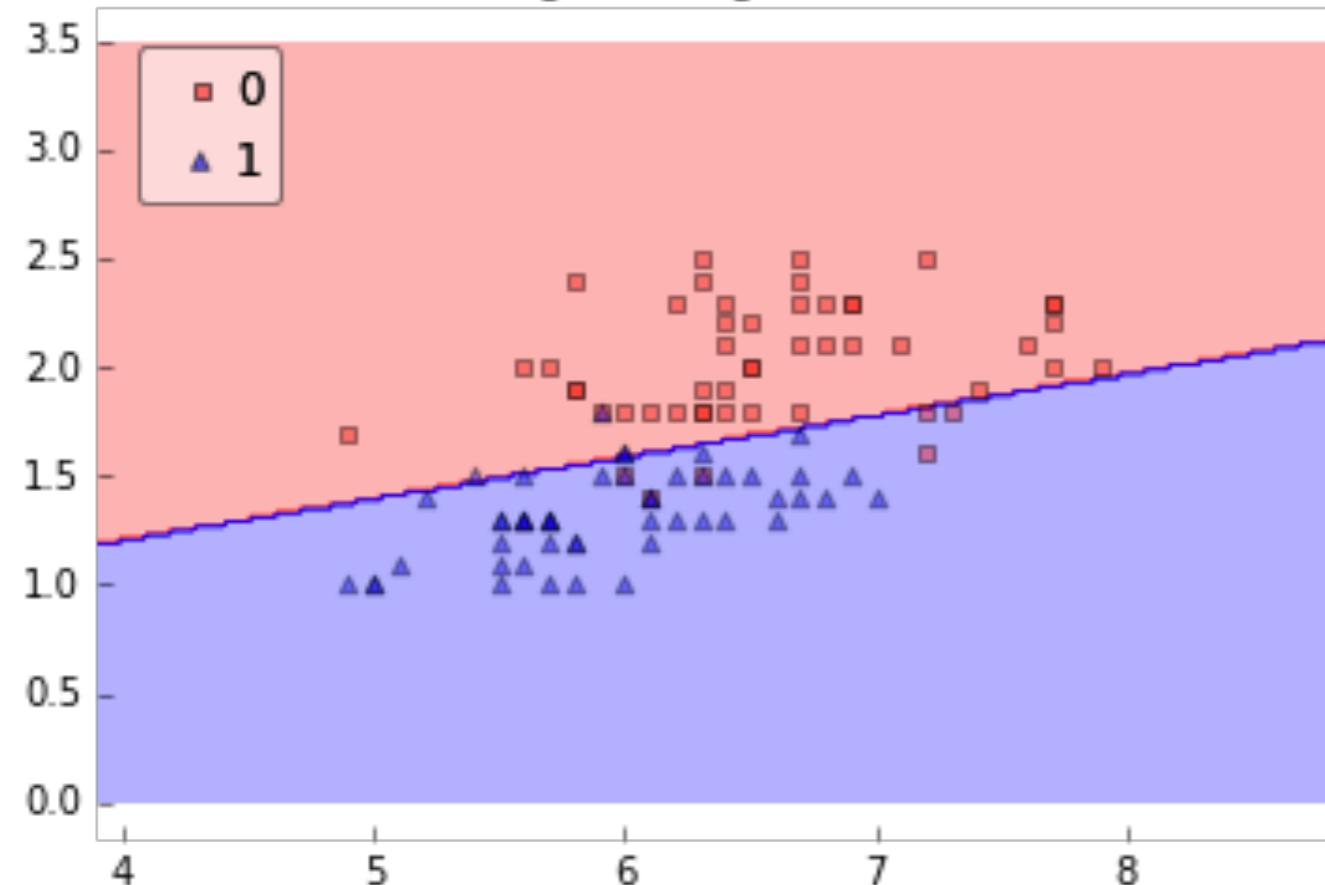
RNN

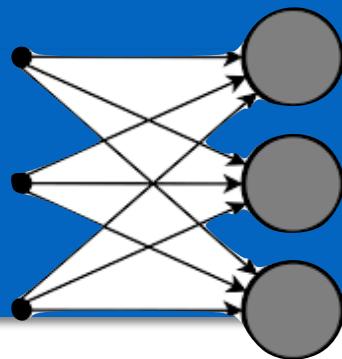
GAN



Função de Ativação

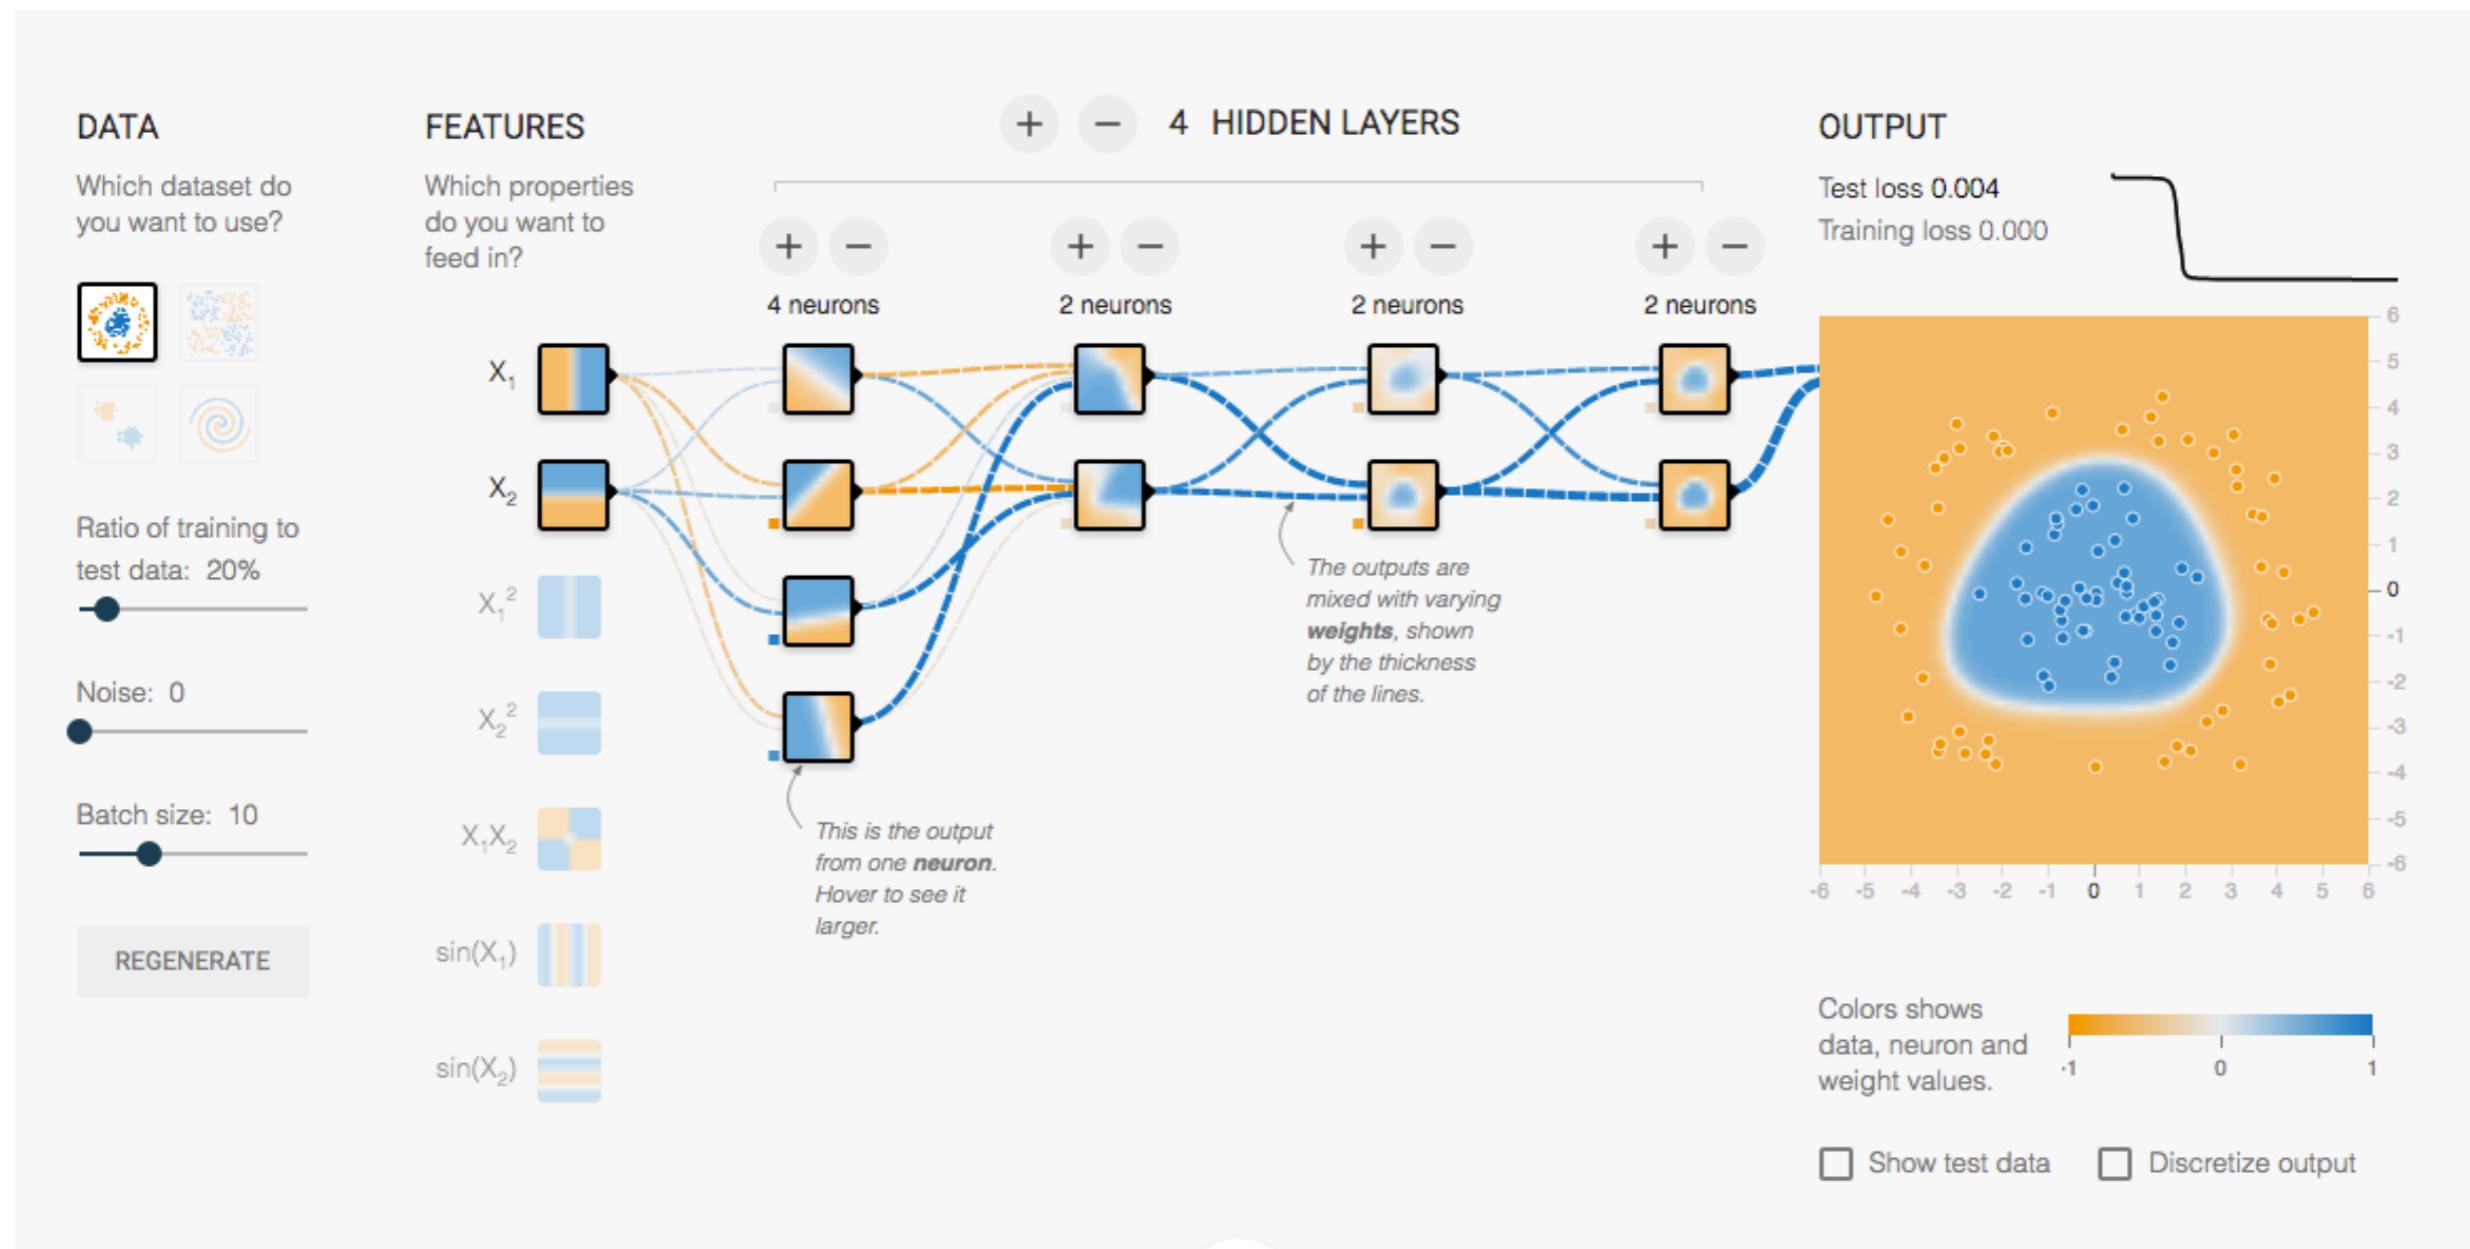
Logistic Regression 2

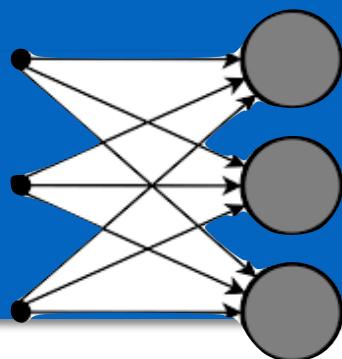




Experimento com Redes Neurais

<http://playground.tensorflow.org/>

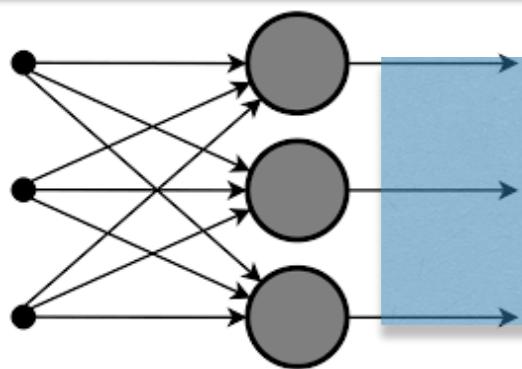




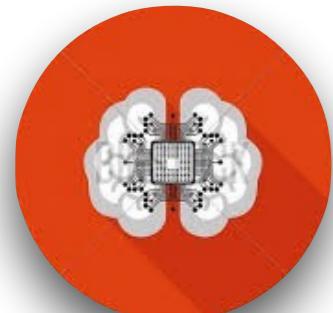
Função de Ativação

Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Sign (Signum)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Linear	$\phi(z) = z$	Adaline, linear regression	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	
Hyperbolic tangent	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multi-layer NN	

Agenda



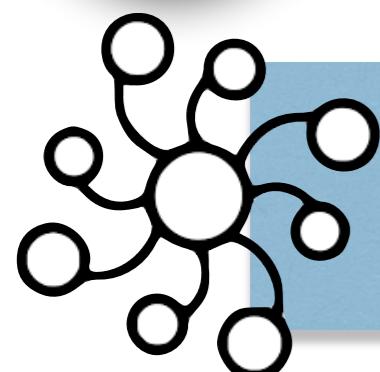
Introdução a Deep Learning



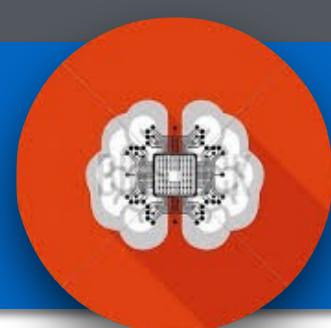
Redes Neurais Convolucionais



Recurrent Neural Networks

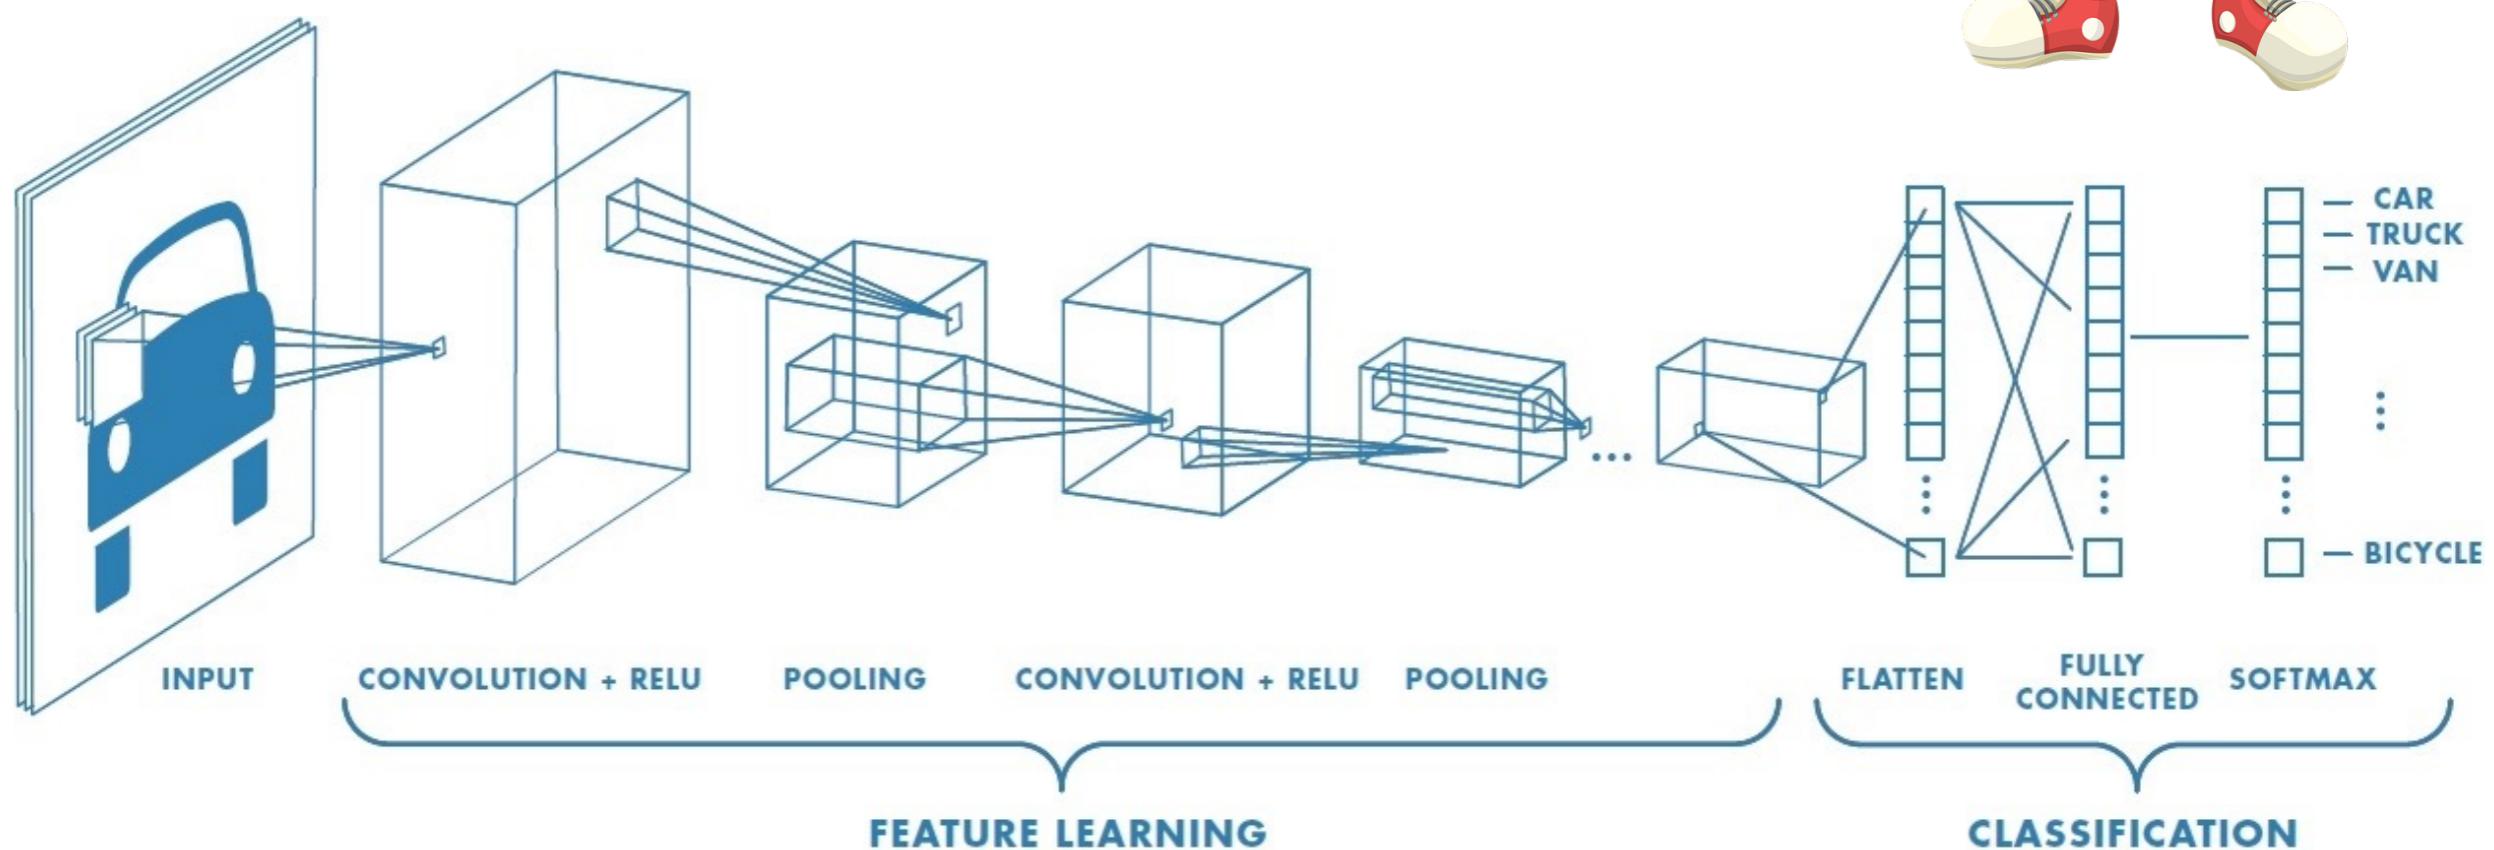


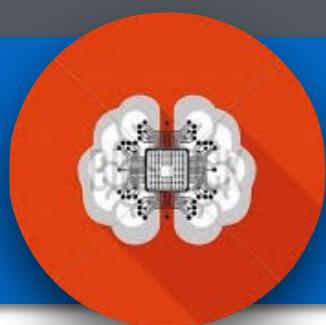
Generative Adversarial Networks



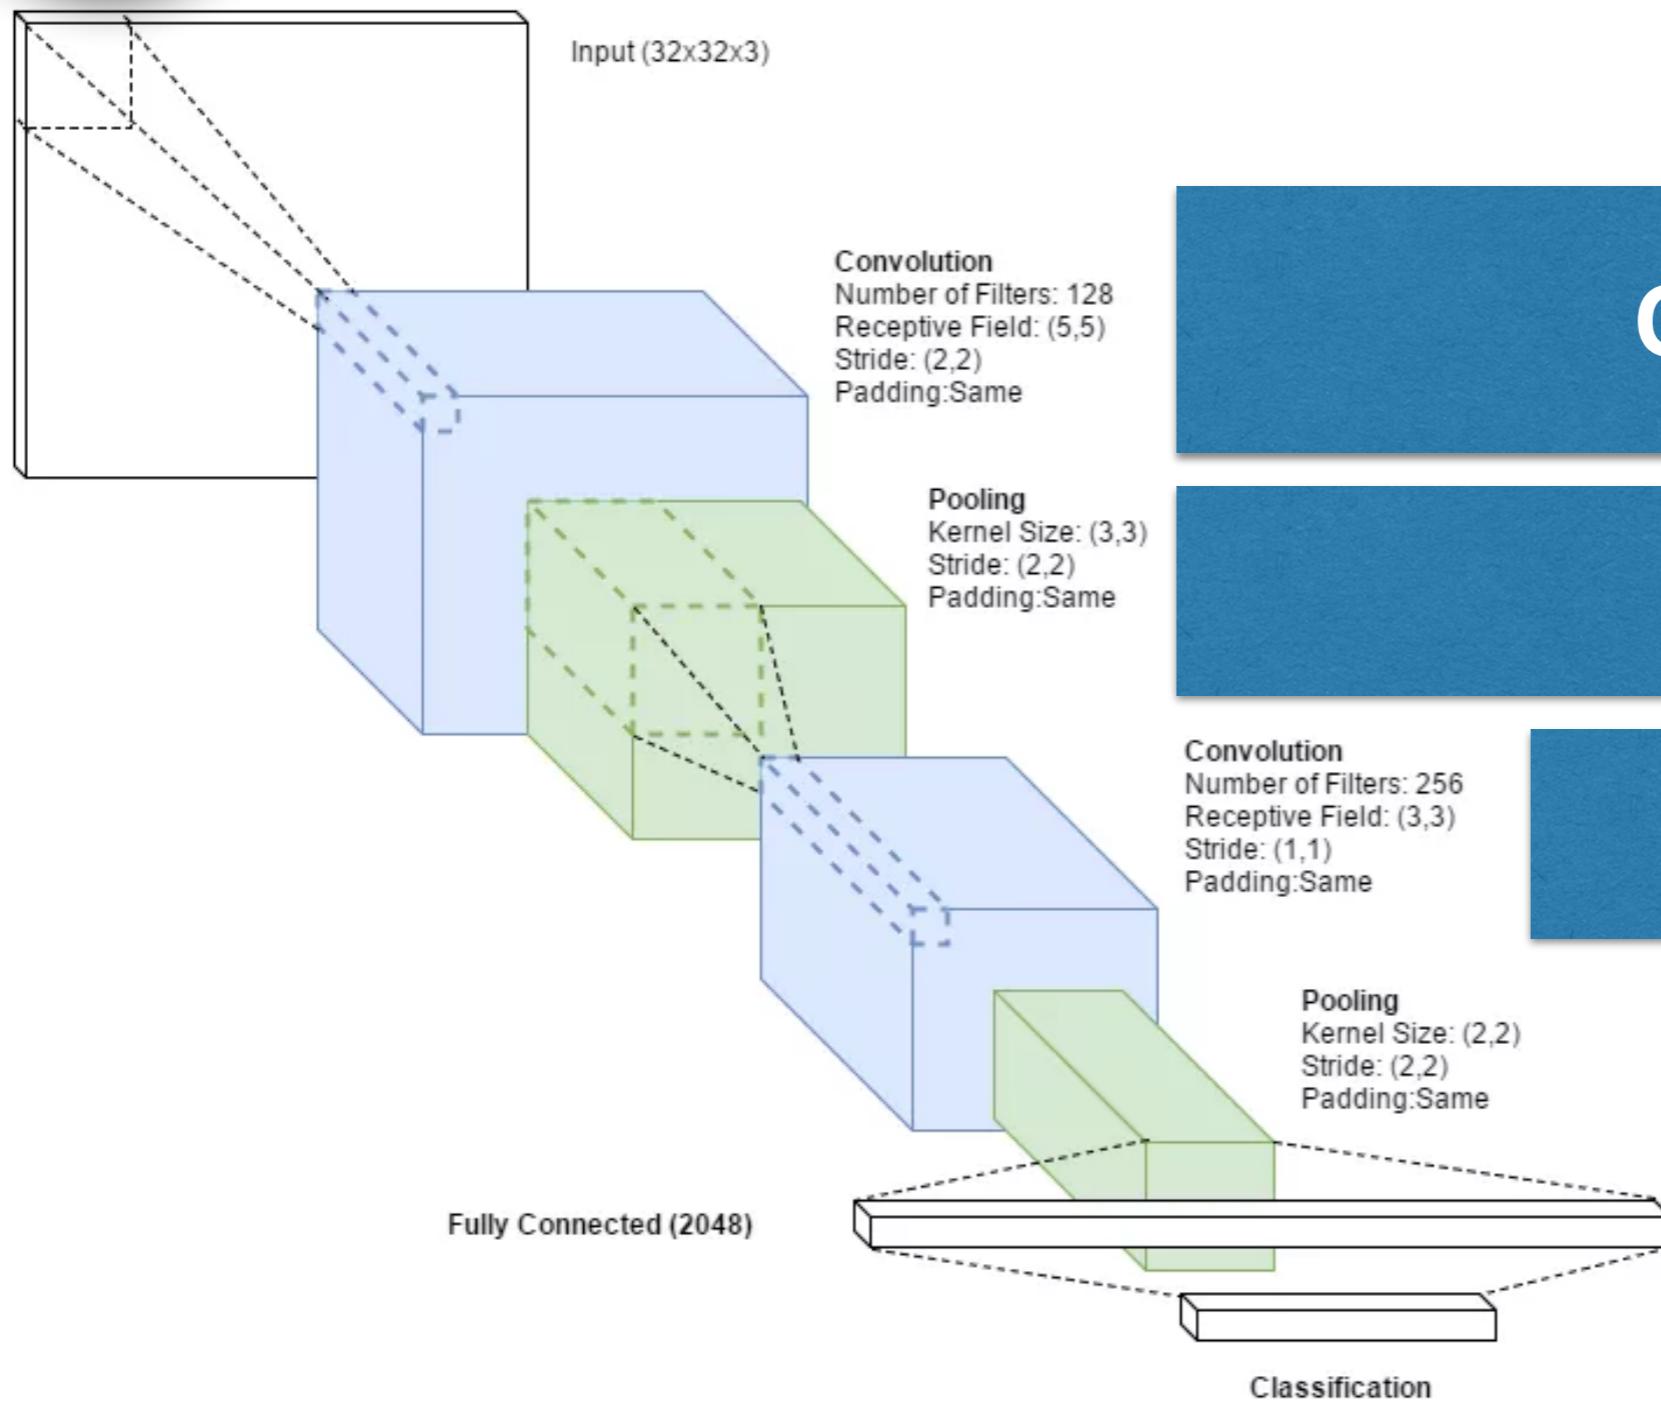
Convolutional Neural Network

Tipo de rede foi inspirada em processos biológicos e são variações de redes multilayer perceptron que são desenhadas o mínimo de pré-processamento possível. Tipo de rede é geralmente usada em reconhecimento de imagens e vídeo.





Convolutional Neural Network

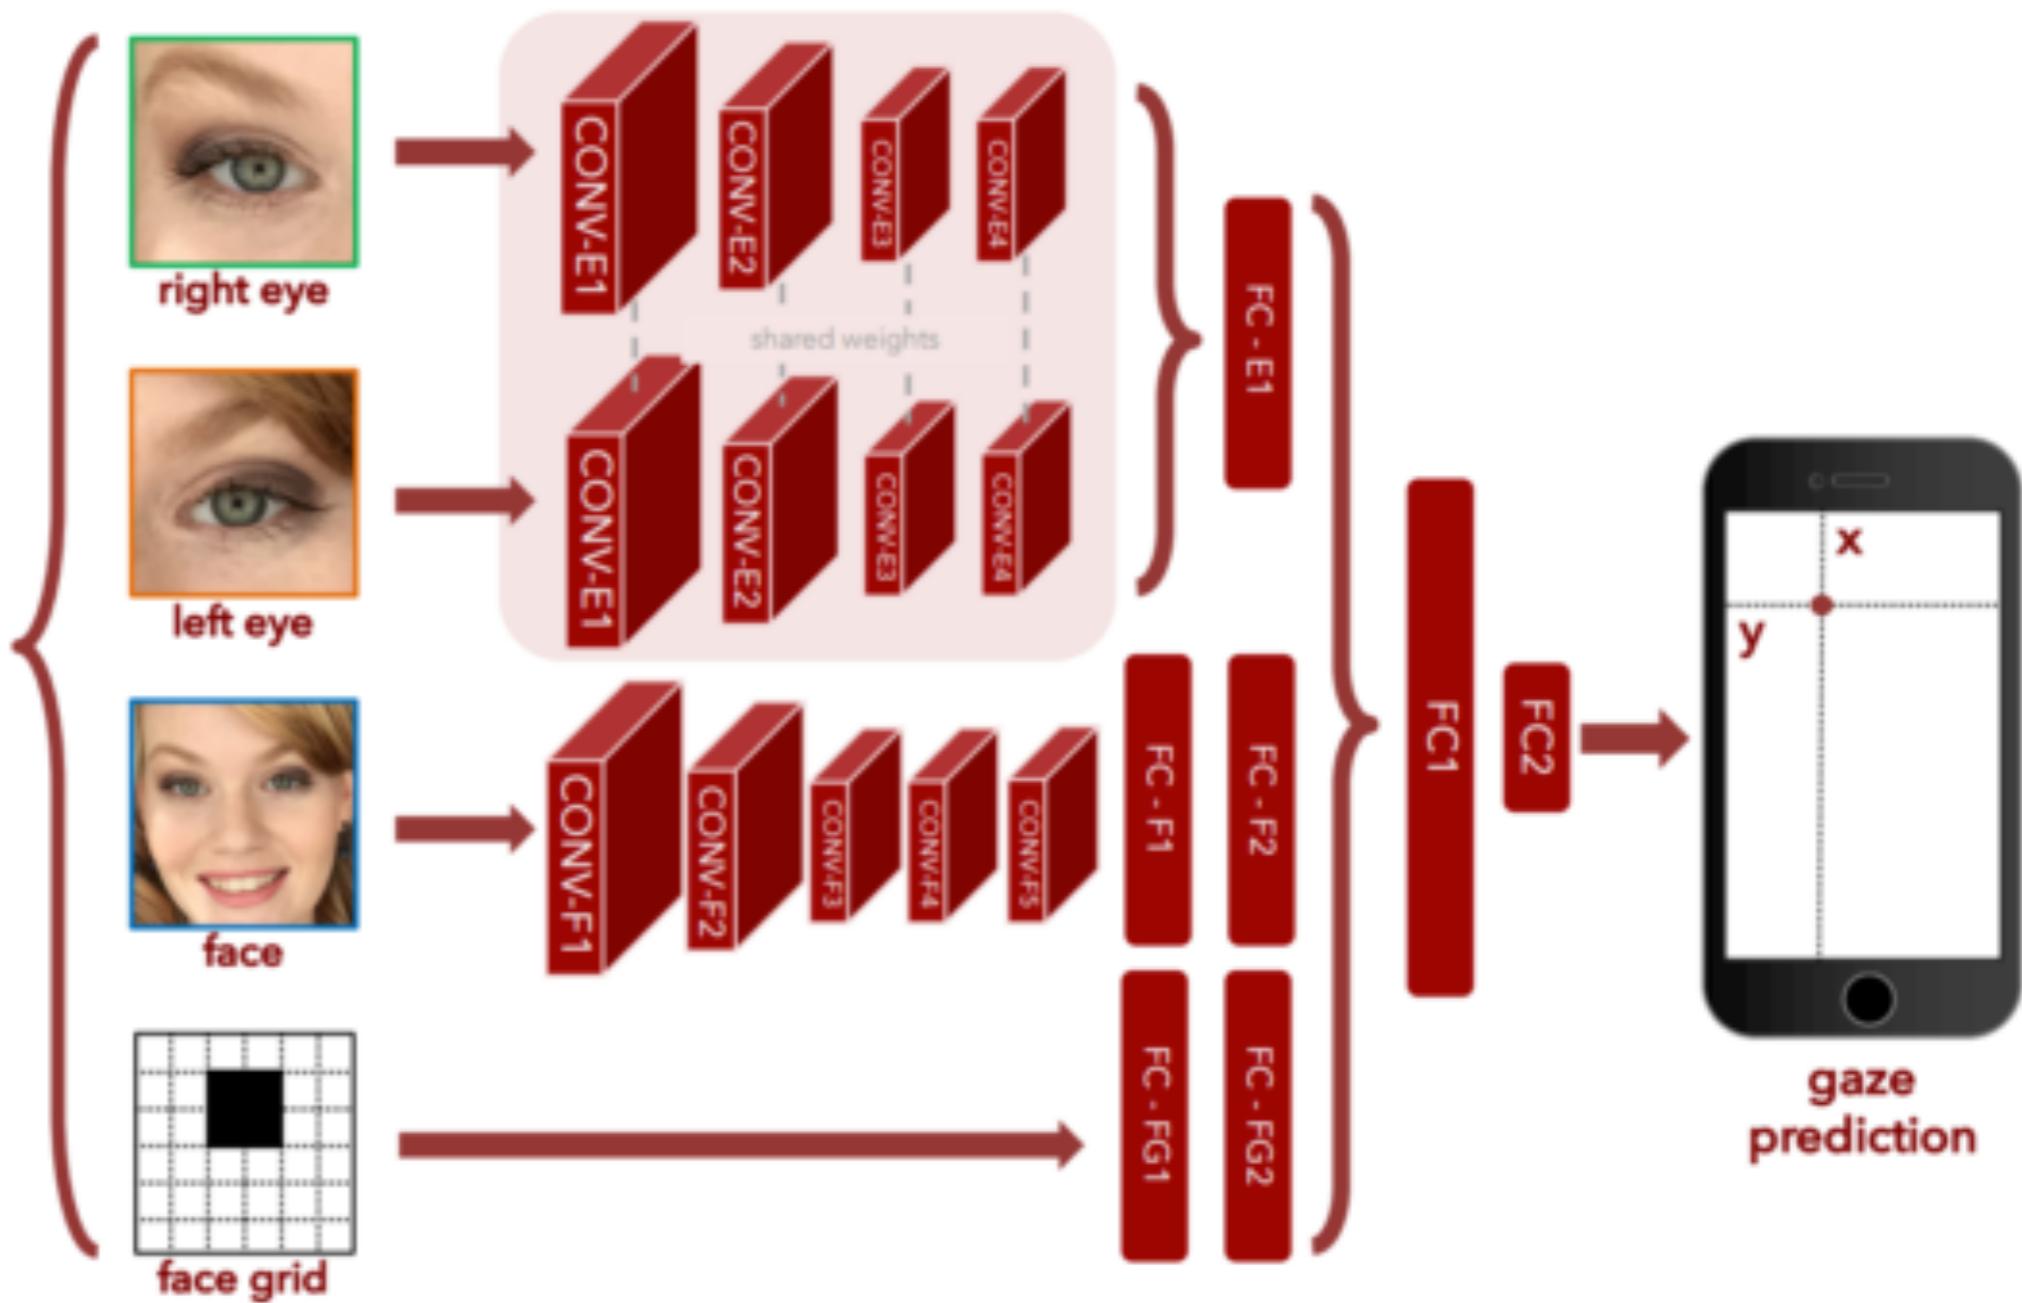
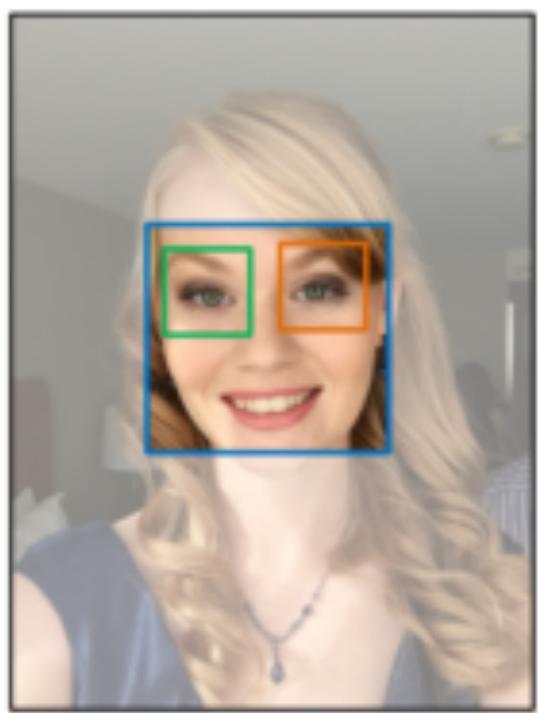


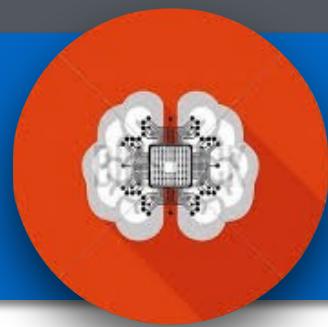
Convolução

Pooling

Convolução

Pooling





$$\begin{pmatrix} 0 & 1 & 2 \\ 2 & 2 & 0 \\ 0 & 1 & 2 \end{pmatrix} \text{ Kernel}$$



Input Vector



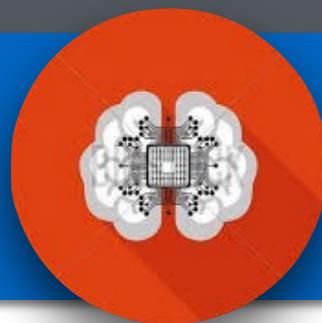
Padding

3 ₀	3 ₁	2 ₂	1	0
0 ₂	0 ₂	1 ₀	3	1
3 ₀	1 ₁	2 ₂	2	3
2	0	0	2	2
2	0	0	0	1

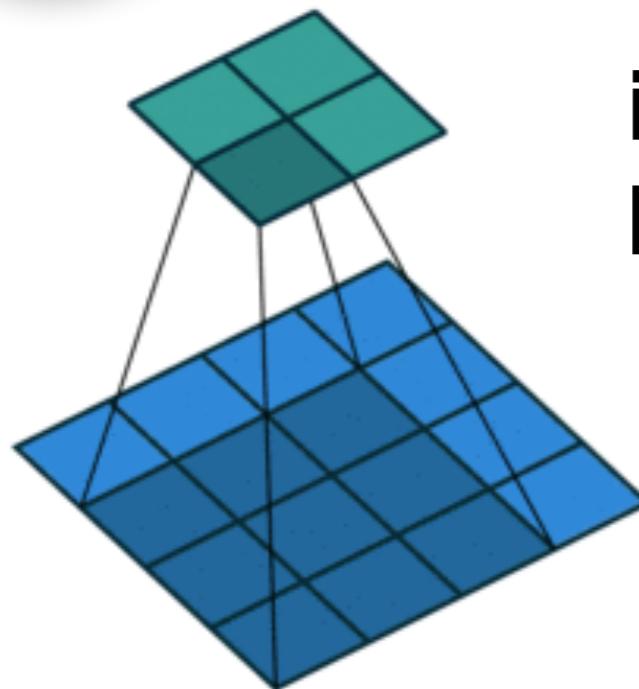
12	12	17
10	17	19
9	6	14

3	3 ₀	2 ₁	1 ₂	0
0	0 ₂	1 ₂	3 ₀	1
3	1 ₀	2 ₁	2 ₂	3
2	0	0	2	2
2	0	0	0	1

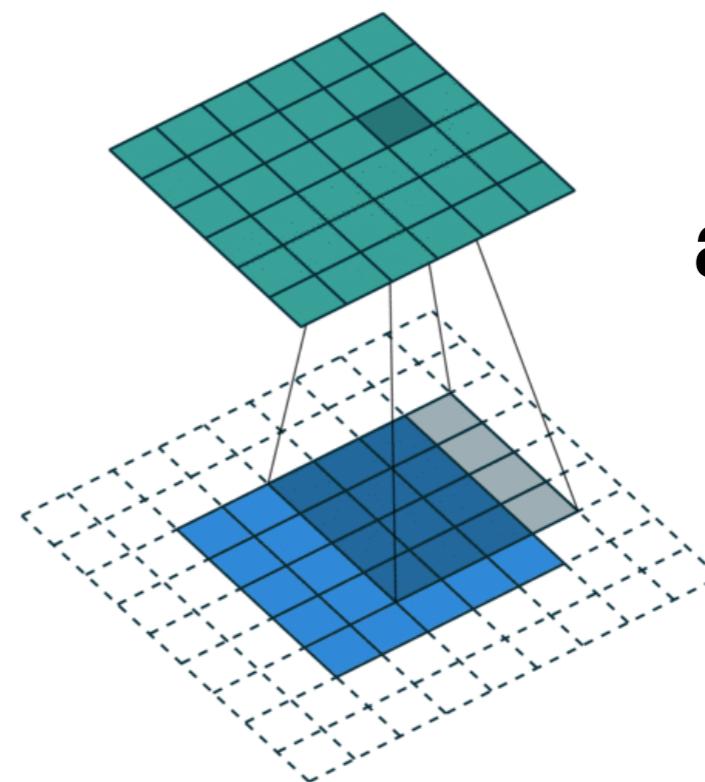
12	12	17
10	17	19
9	6	14



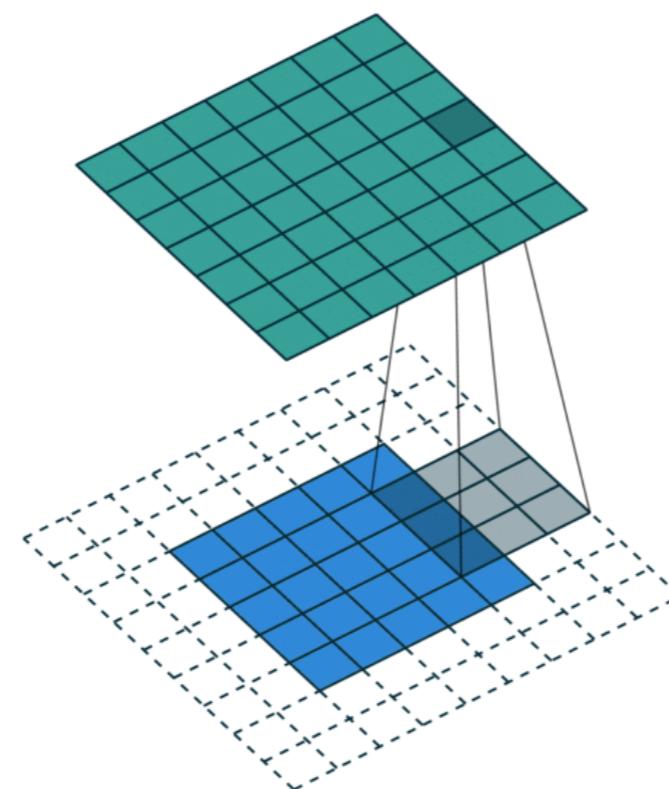
Convolução



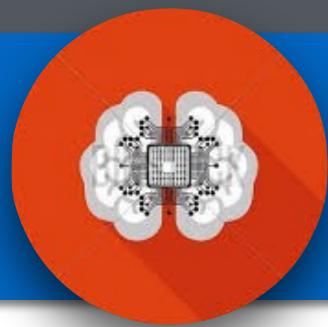
$i = 4$ and
 $k = 3$



$i = 5$, $k = 4$
and $p = 2$:

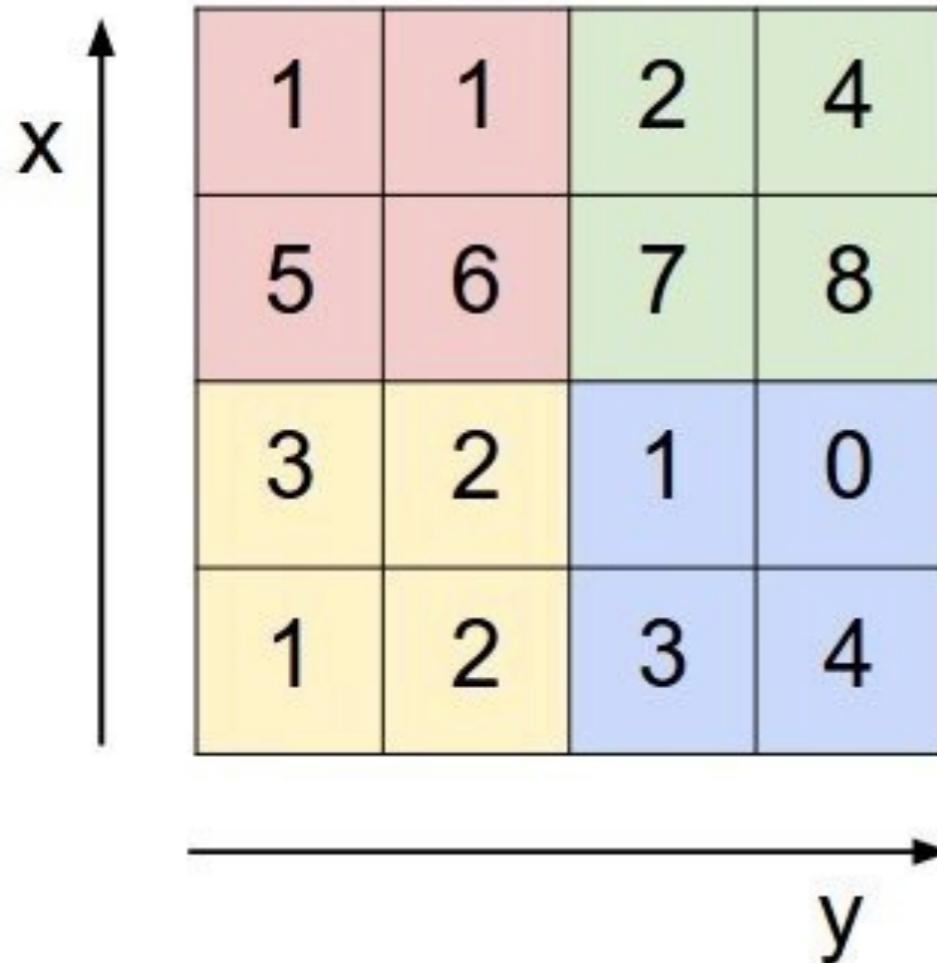


$i = 5$, $k = 3$ and
 $p = 2$:



Pooling

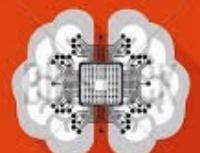
Single depth slice



max pool with 2x2 filters
and stride 2

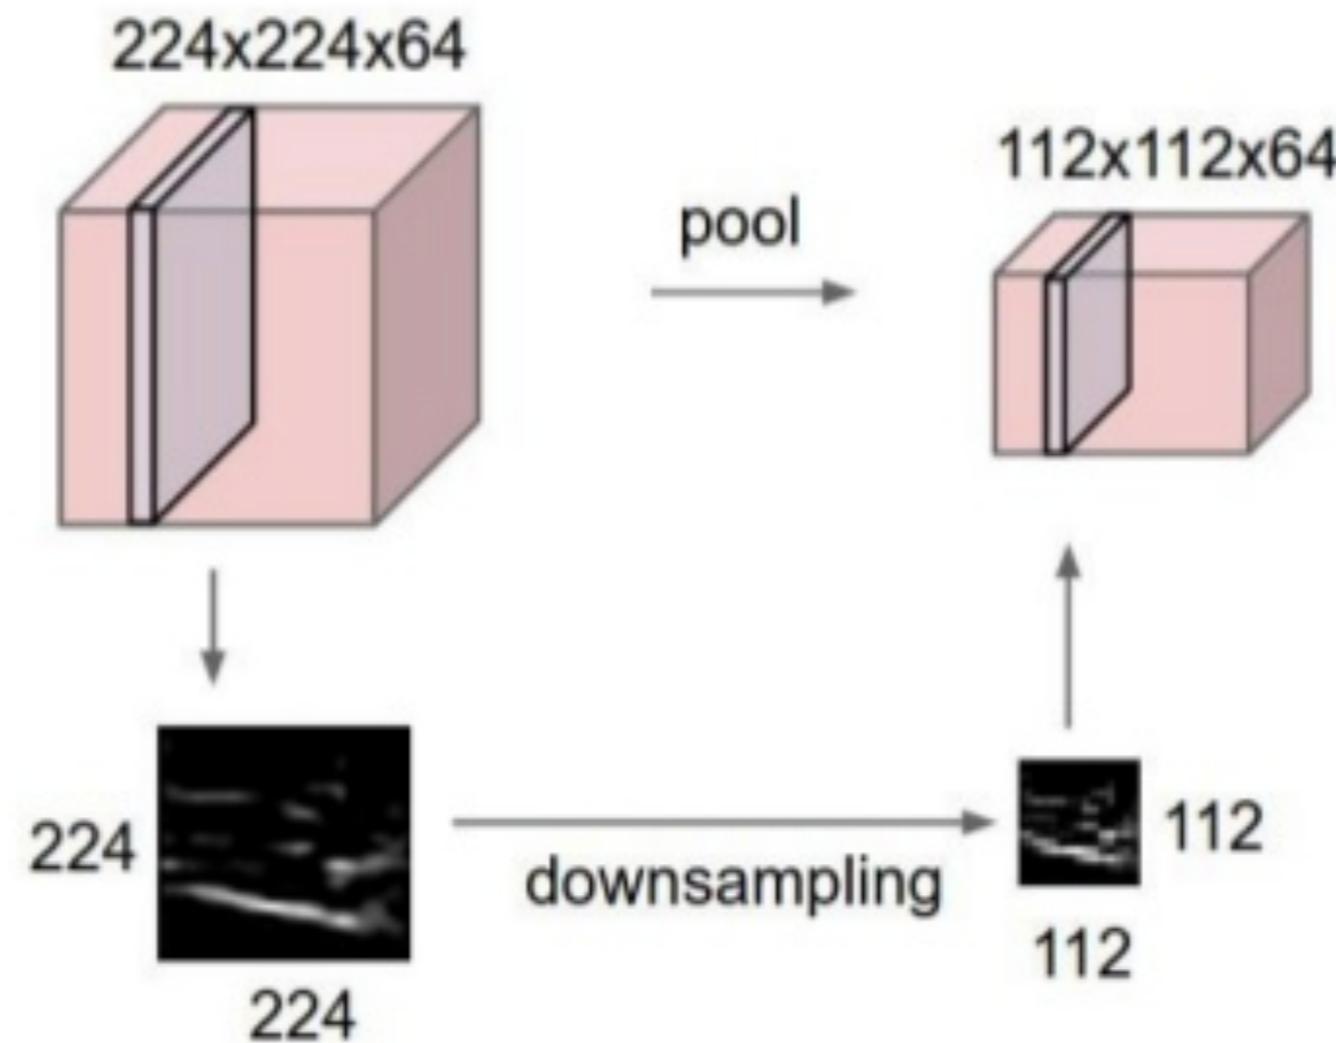
The output slice is a 2x2 matrix resulting from the max pooling operation. The values are:

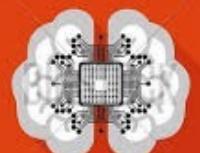
6	8
3	4



Pooling

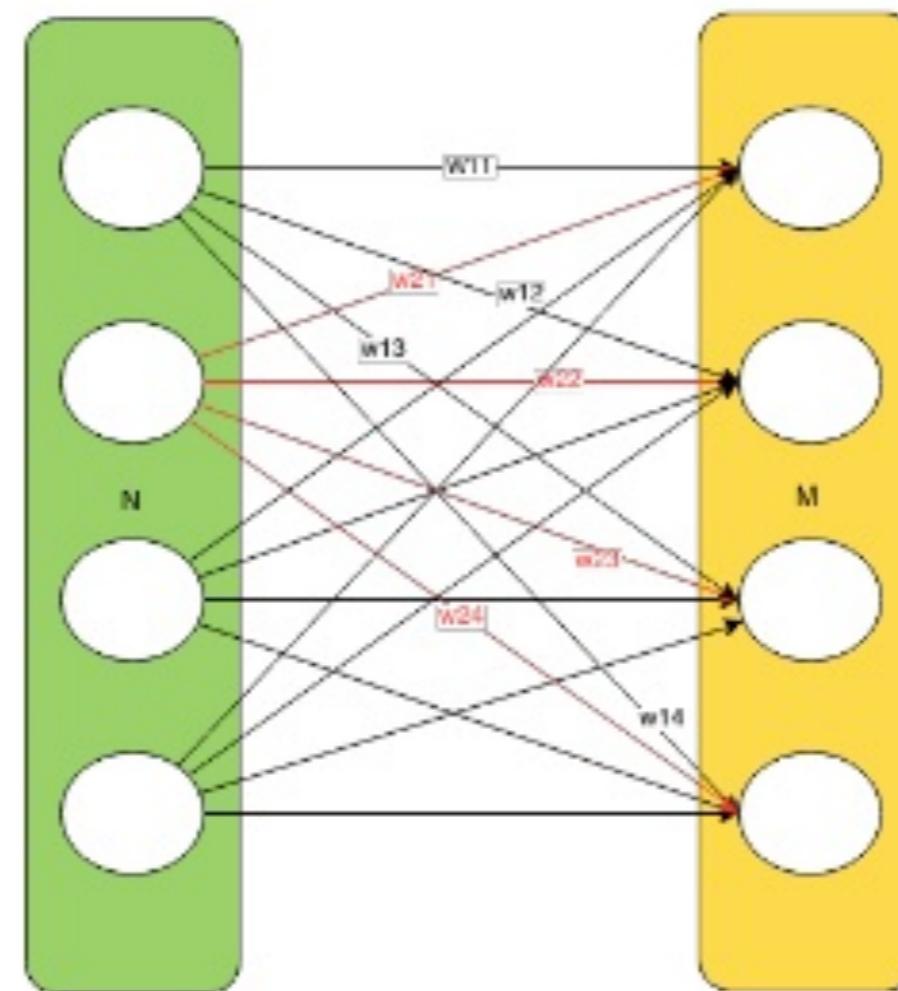
Pooling Layer

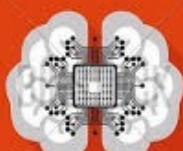




Rede Neural

Fully-connected layer





Casos de Sucesso

Case studies

LeNet: The first successful applications of CNN

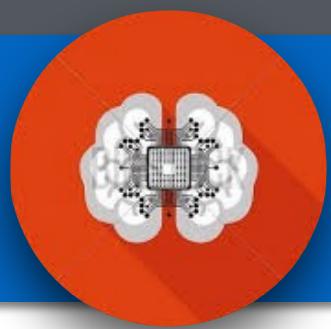
AlexNet: The first work that popularized CNN in Computer Vision

ZF Net: The ILSVRC 2013 winner

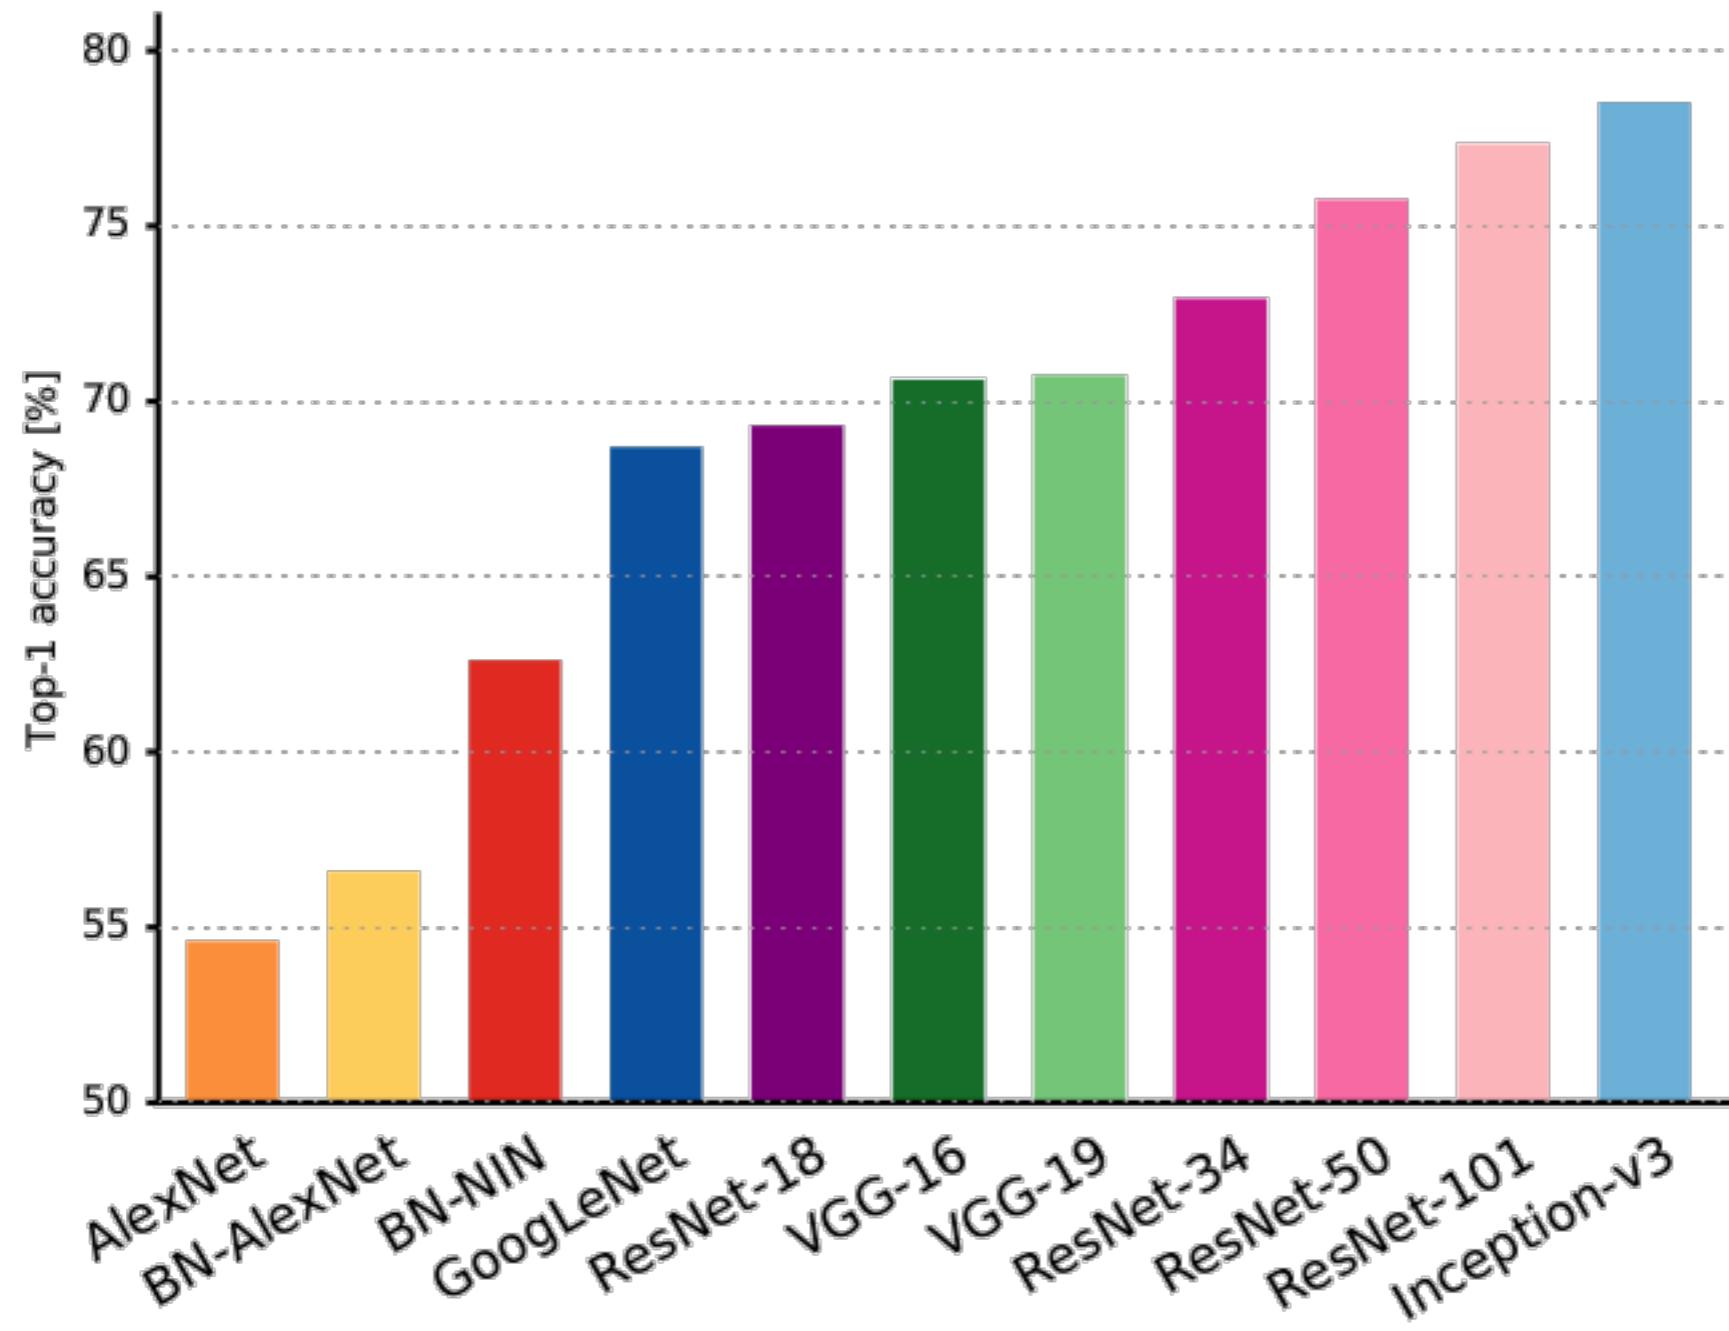
GoogLeNet: The ILSVRC 2014 winner

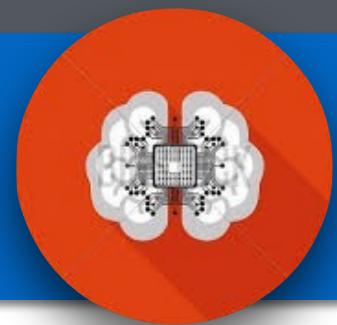
VGGNet: The runner-up in ILSVRC 2014

ResNet: The winner of ILSVRC 2015



Redes Treinadas

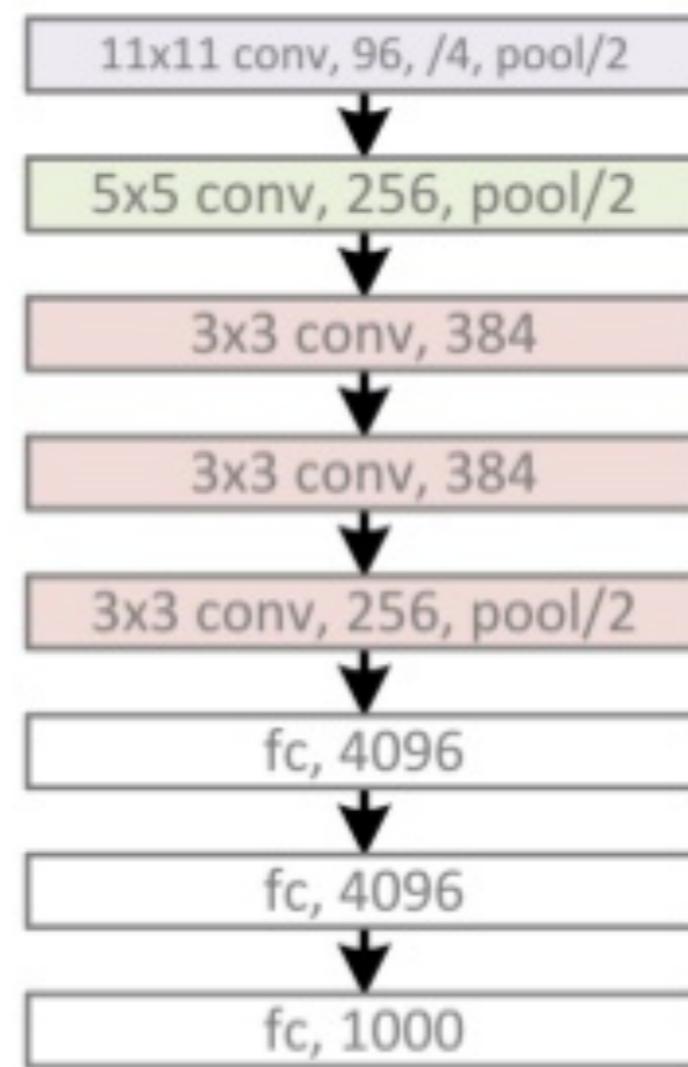


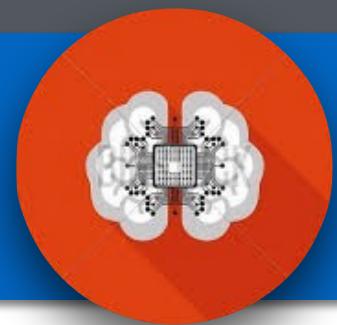


Redes Treinadas

AlexNet

AlexNet, 8 layers
(ILSVRC 2012)





Redes Treinadas

Revolution of Depth

AlexNet, 8 layers
(ILSVRC 2012)



VGG, 19 layers
(ILSVRC 2014)



ResNet, 152 layers
(ILSVRC 2015)



LeNet

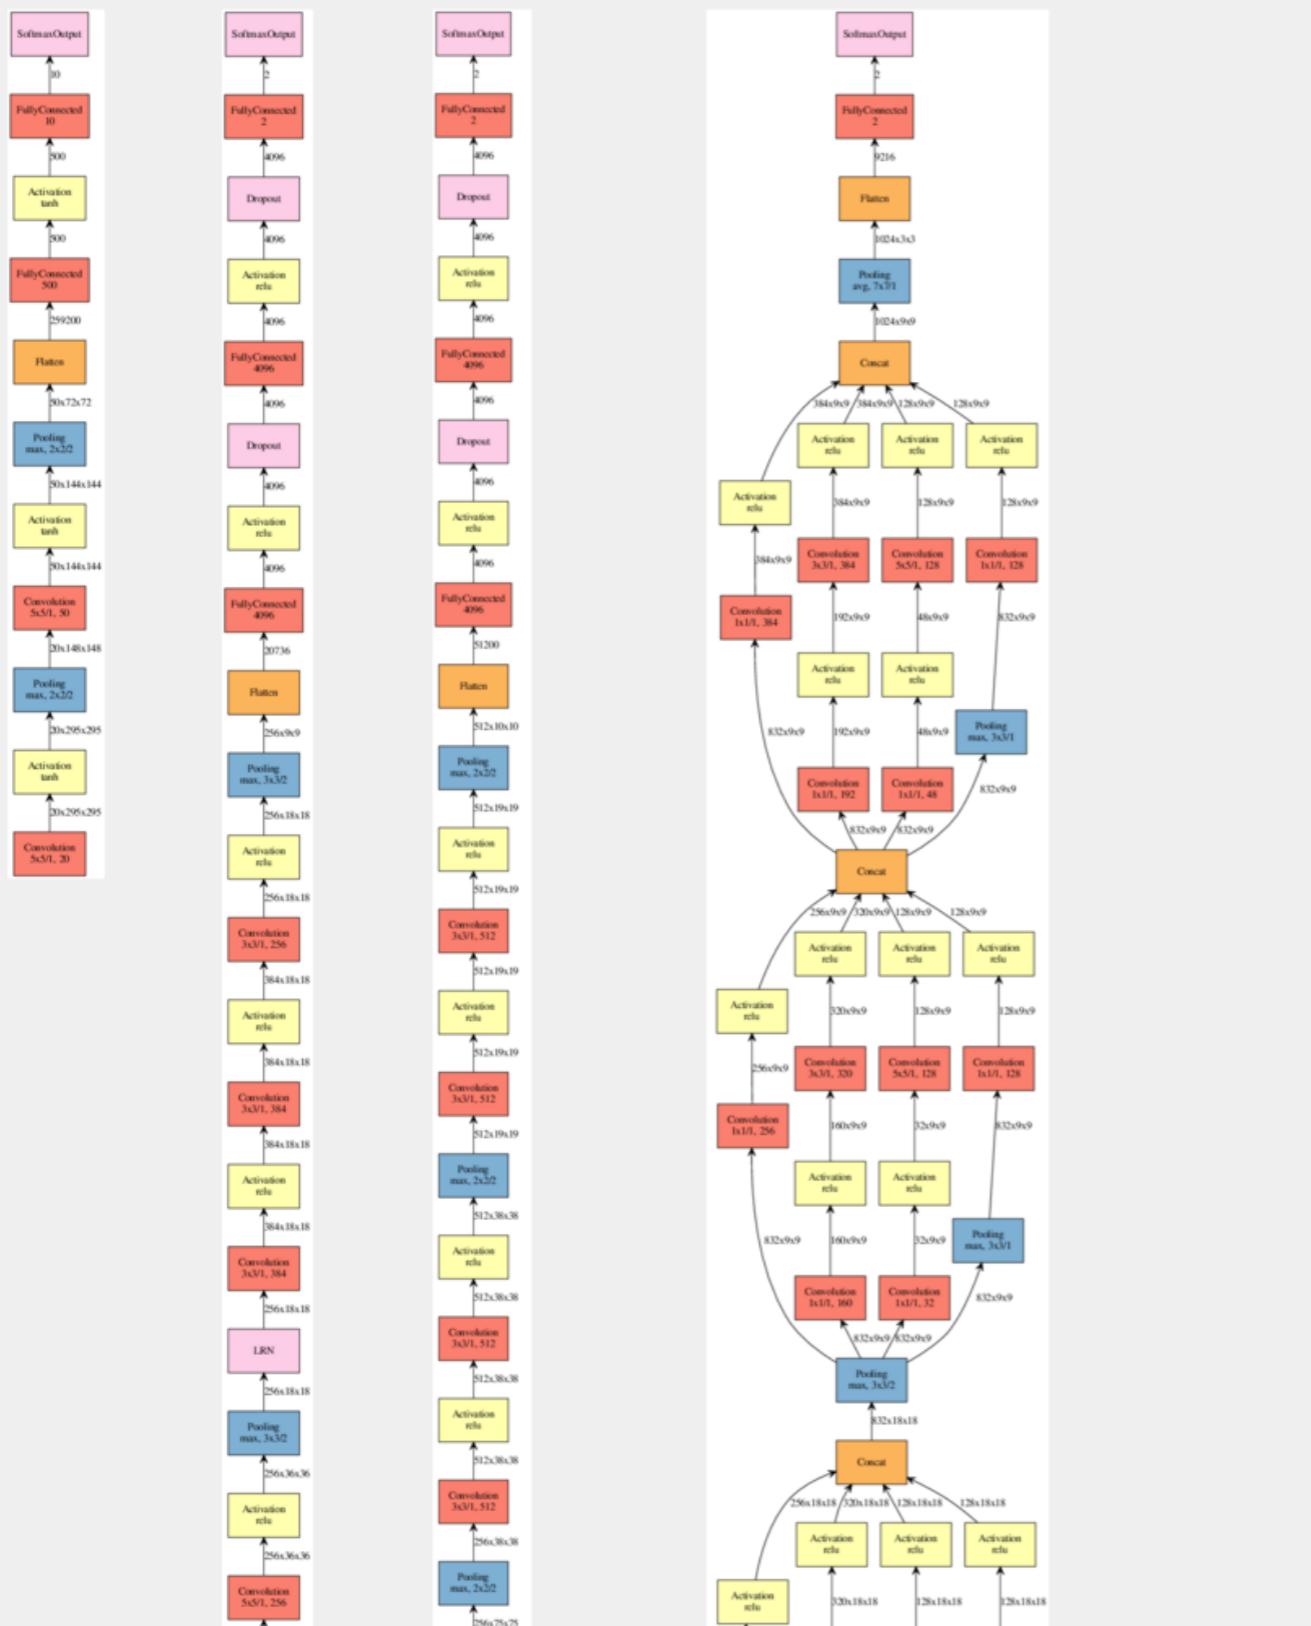
AlexNet

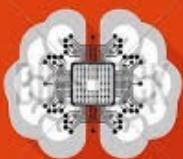
VGG

GoogLeNet

Inception V3

Inception BN





Datasets

Datasets Benchmark

MNIST Handwritten digits – 60000 Training + 10000 Test Data

Google House Numbers from street view - 600,000 digit images

CIFAR-10 60000 32x32 colour images in 10 classes

IMAGENET >150 GB

Tiny Images 80 Million tiny images

Flickr Data 100 Million Yahoo dataset

MNIST Dataset CIFAR-10 Dataset



airplane



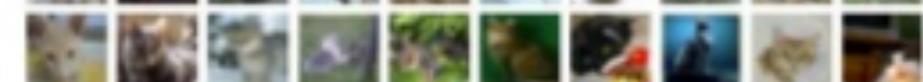
automobile



bird



cat



deer



dog



frog



horse



ship

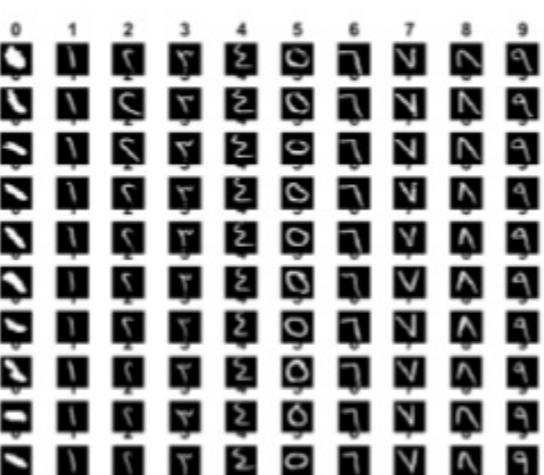


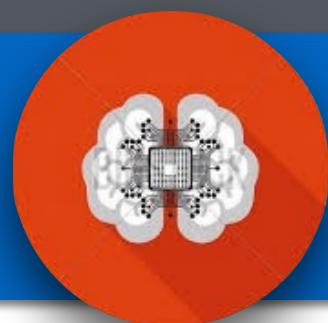
truck



Arabic Handwritten Digits Database

- By Sheriff Abdelazeem, Ezzat El-Sherif
- 70,000 digits





Implementação no Keras



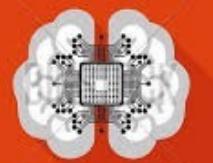
```
model = Sequential()
```

```
model.add(Convolution2D(12, 5, 5, activation = 'relu', input_shape=in_shape,  
init='he_normal'))
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model.add(Convolution2D(25, 5, 5, activation = 'relu', init='he_normal'))
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
```



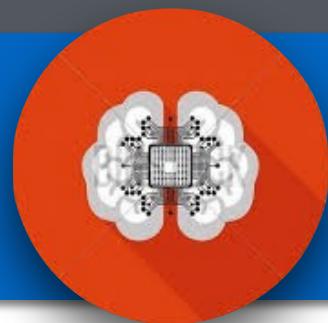
Implementação no TensorFlow

```
# Convolution Layer with 32 filters and a kernel size of 5
conv1 = tf.layers.conv2d(x, 32, 5, activation=tf.nn.relu)
# Max Pooling (down-sampling) with strides of 2 and kernel size of 2
conv1 = tf.layers.max_pooling2d(conv1, 2, 2)

# Flatten the data to a 1-D vector for the fully connected layer
fc1 = tf.contrib.layers.flatten(conv2)

# Fully connected layer (in tf contrib folder for now)
fc1 = tf.layers.dense(fc1, 1024)
# Apply Dropout (if is_training is False, dropout is not applied)
fc1 = tf.layers.dropout(fc1, rate=dropout, training=is_training)

# Output layer, class prediction
out = tf.layers.dense(fc1, n_classes)
```



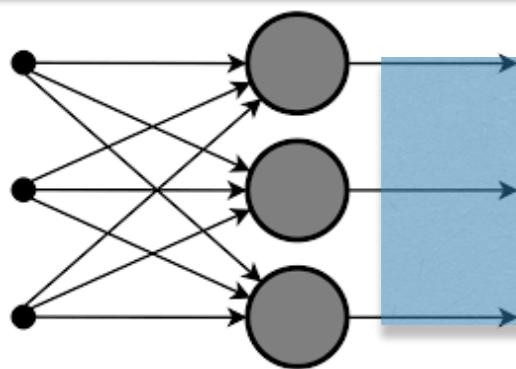
Exemplos

<https://github.com/hiteshvaidya/Star-wars-classifier>

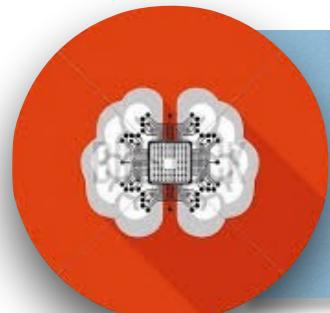
**[https://github.com/rdcolema/keras-image-classification/
blob/master/cats_n_dogs_BN.ipynb](https://github.com/rdcolema/keras-image-classification/blob/master/cats_n_dogs_BN.ipynb)**

**[https://github.com/rajshah4/image_keras/blob/master/
notebook.ipynb](https://github.com/rajshah4/image_keras/blob/master/notebook.ipynb)**

Agenda



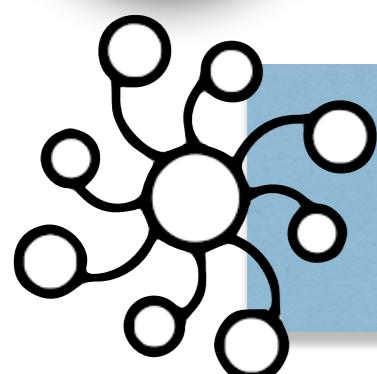
Introdução a Deep Learning



Redes Neurais Convolucionais



Recurrent Neural Networks



Generative Adversarial Networks



Recurrent Neural Networks

ANALISANDO TEXTO

Passados quase três meses do final dos Jogos, o Comitê

Rio-2016 ainda deve reembolso a

8.000 torcedores que utilizaram sua plataforma online para revender ingressos.

A entidade reduziu o contingente de consumidores a quem devia pagamentos, que chegou a 140 mil em 19 de outubro, data até a qual prometeu quitar os débitos. Mas ainda não deu fim ao problema.

A entidade afirmou que tem dificuldades para ressarcir o restante. Alega problemas para encontrar os credores e inconsistência nos dados bancários fornecidos —muitos depósitos não foram completados.

De acordo com o **comitê**, 3.500 pessoas foram procuradas mas não responderam às mensagens eletrônicas, 2.500 até deram retorno, porém as informações repassadas continham algum erro e 2.000 devem receber o reembolso até esta segunda (12), após terem dados checados.

Uma mutação aparentemente insignificante no

DNA dos ancestrais da humanidade

pode ter contribuído para que nosso cérebro alcançasse o tamanho descomunal que tem hoje (três vezes maior que o dos grandes macacos).

Bastou inserir o gene que contém essa mutação em fetos de camundongo para que dobrasse o número de células que dão origem aos neurônios do córtex, a área cerebral mais "nobre".

A **pesquisa**, conduzida por

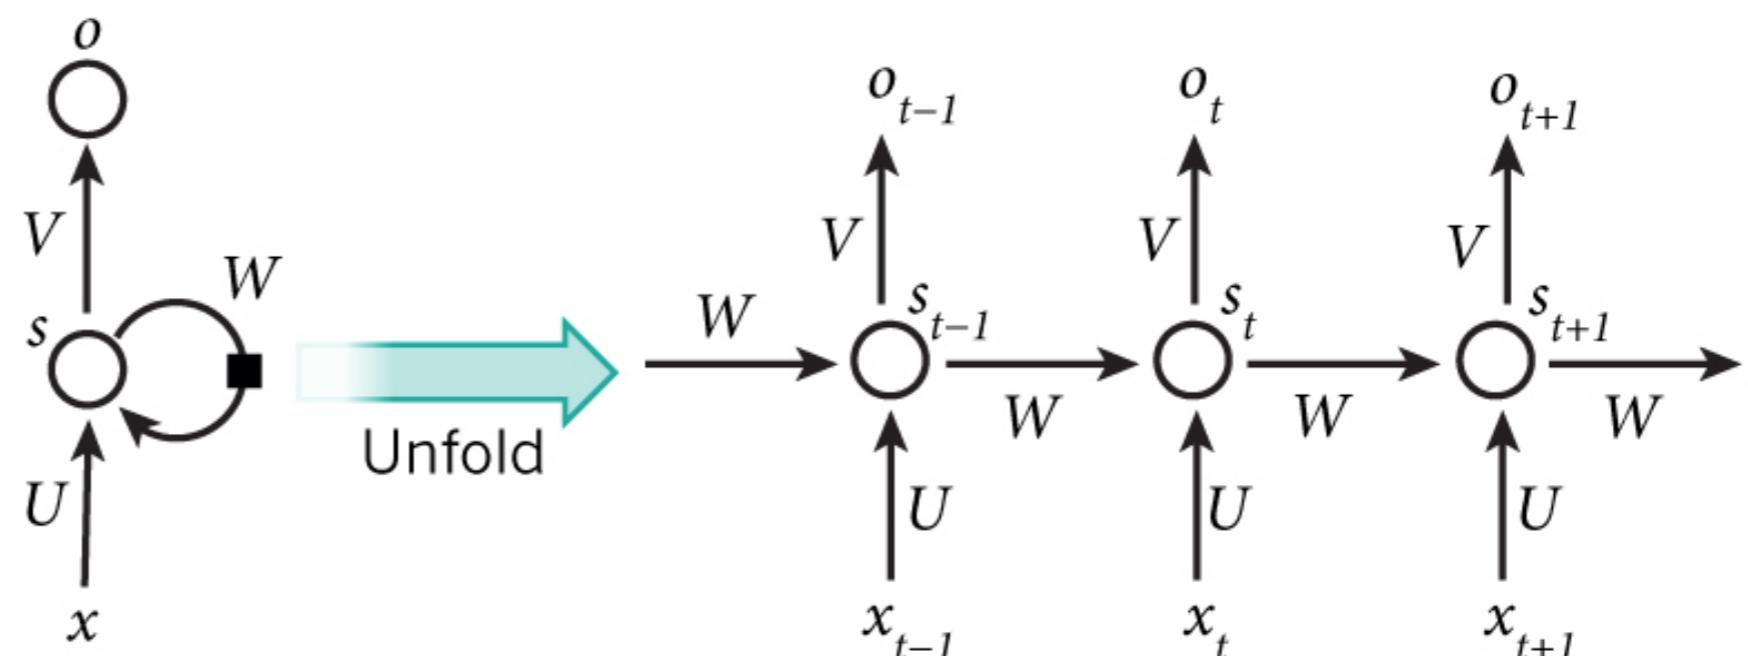
Instituto Max

Planck (Alemanha), é um dos primeiros frutos da tentativa de usar o genoma para entender como a evolução humana se desenrolou. Por enquanto, isso não tem sido fácil —tanto que o gene analisado pelos pesquisadores no novo estudo, designado pela indigesta sigla ARHGAP11B, é o único específico da linhagem humana a ser associado com a proliferação das tais células do córtex cerebral.



Recurrent Neural Networks

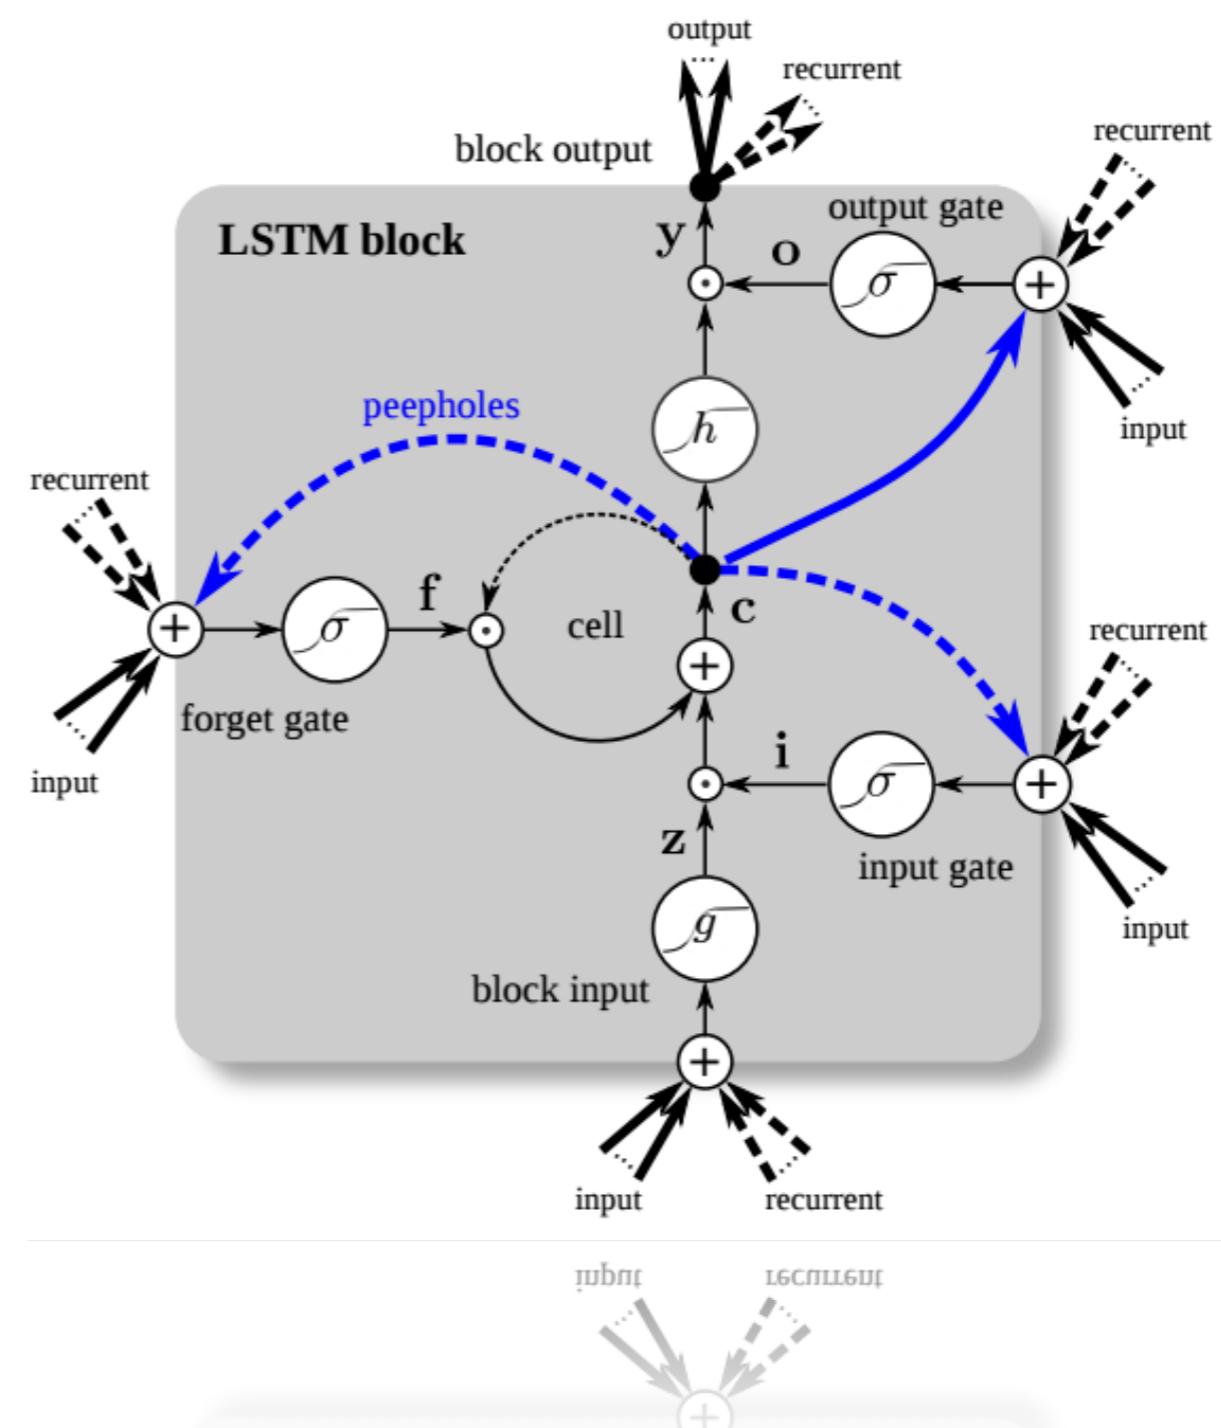
Rnns por trás da idéia é fazer uso da informação sequencial. Em uma rede neural tradicional assumimos que todas as entradas (e saídas) são independentes um do outro. Mas para muitas tarefas que é uma idéia muito ruim.





Longa memória de curto prazo (LSTM)

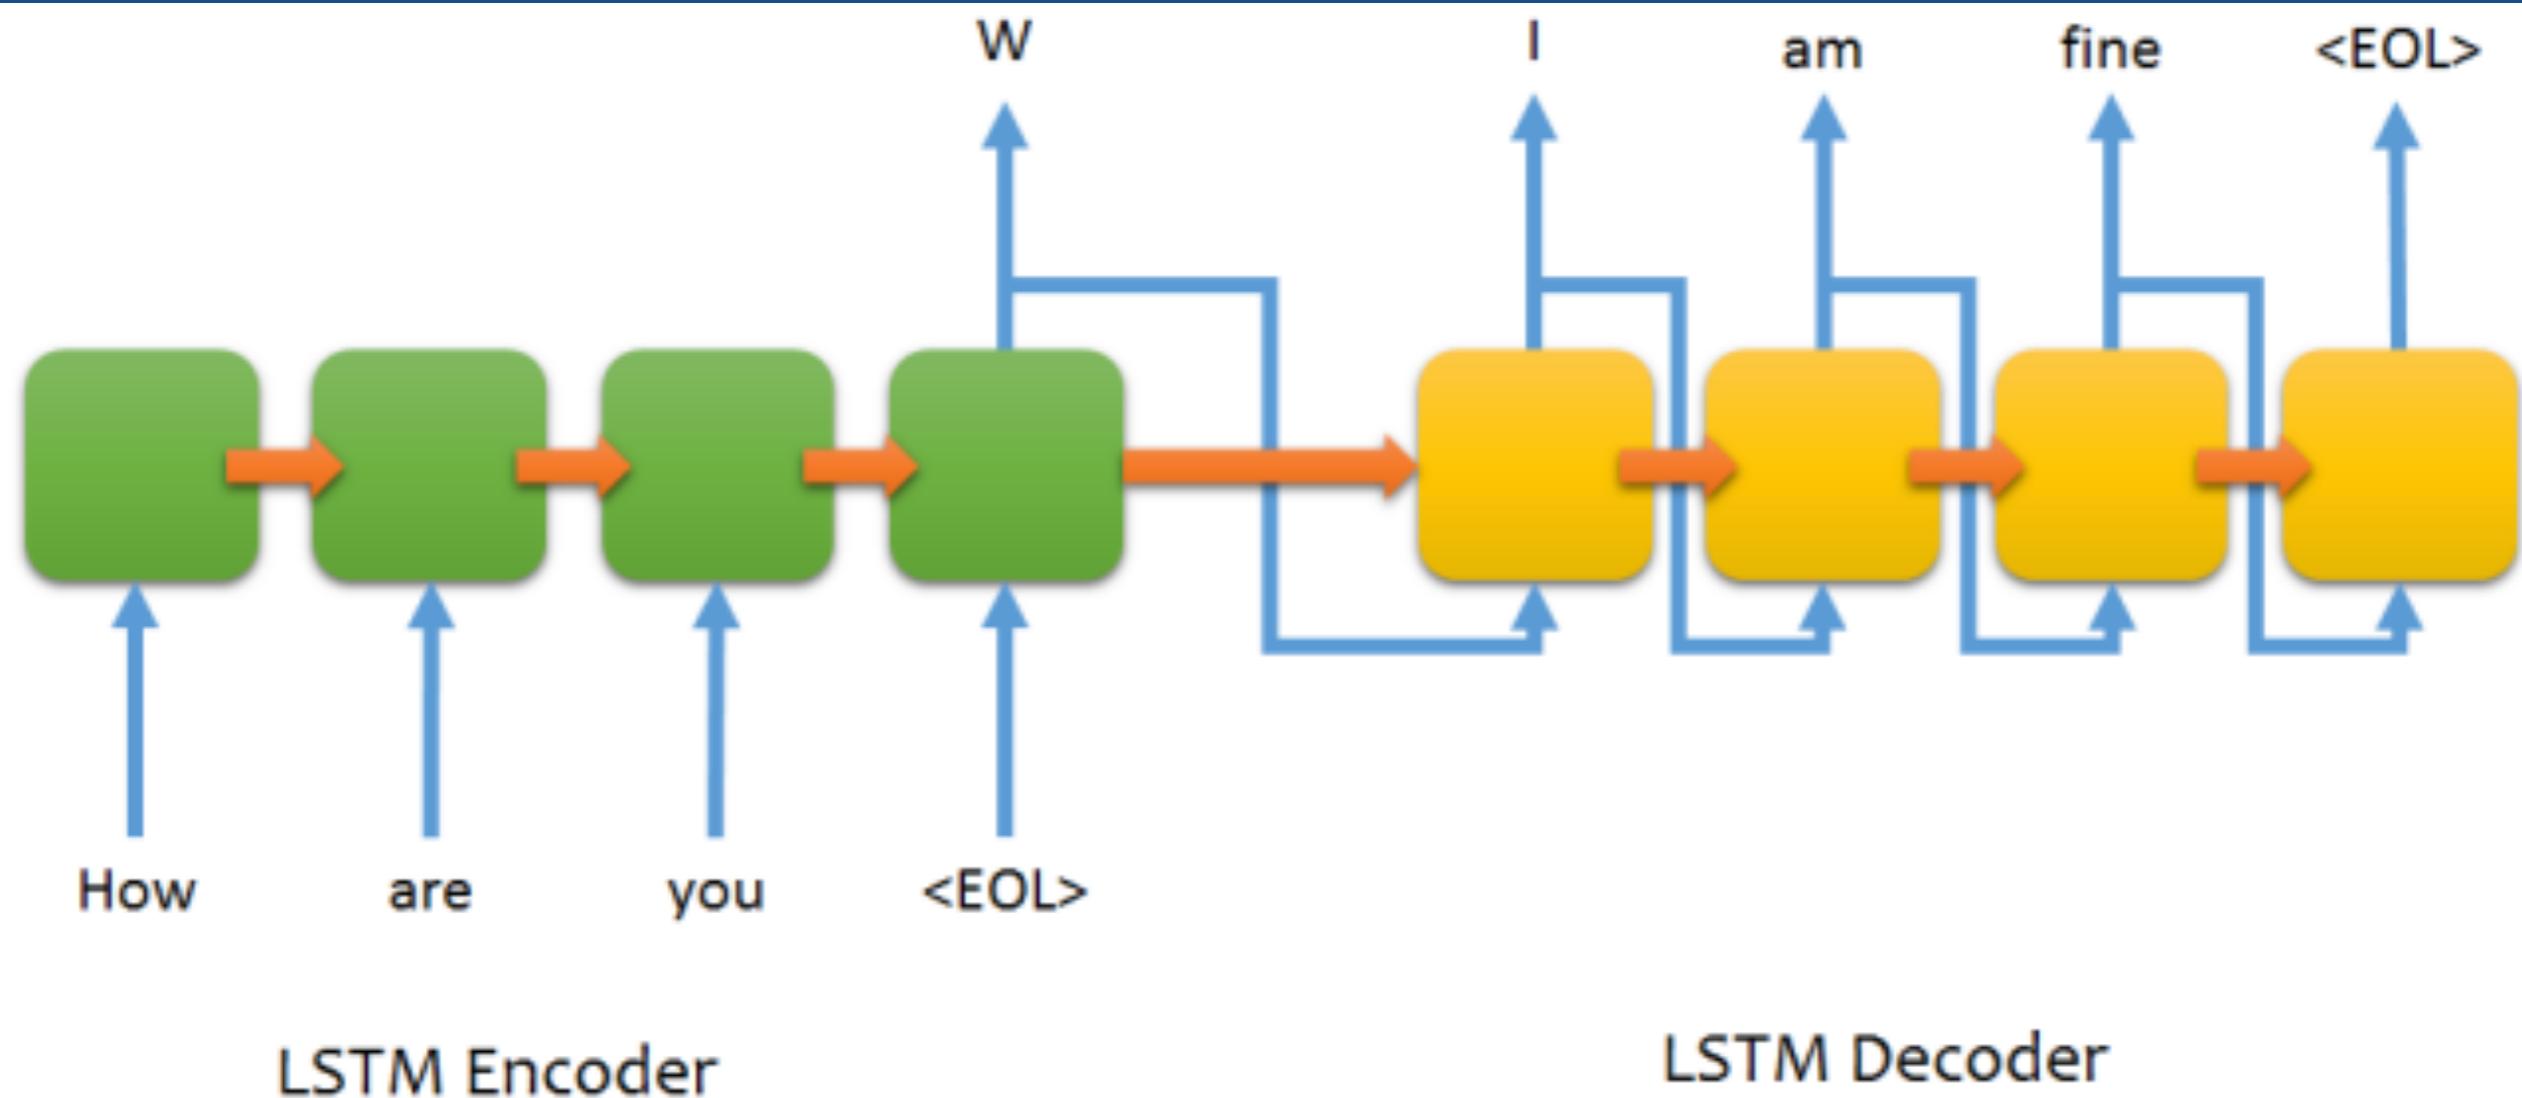
- Longa memória de curto prazo (LSTM) é uma arquitetura de rede neural recorrente (RNN) que lembra os valores em intervalos arbitrários.
- Valores armazenados não são modificados à medida que prossegue aprendizagem.





Longa memória de curto prazo (LSTM)

Um modelo básico sequência-a-sequência, tal como foi introduzido em Cho et al, 2014 (pdf), consiste de duas rede neural recorrente (rnns): um codificador que processa a entrada e um descodificador que gera a saída.

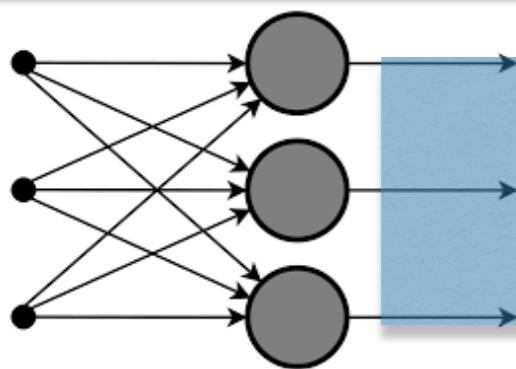




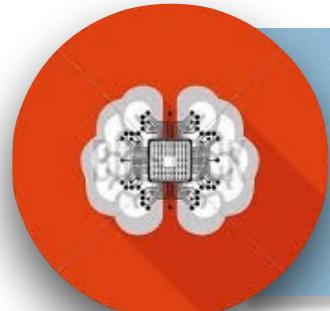
Exemplos

<https://github.com/RyanCCollins/deep-learning-nd>

Agenda



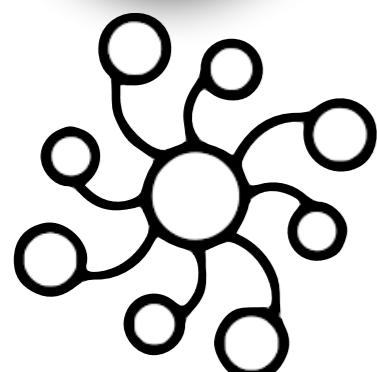
Introdução a Deep Learning



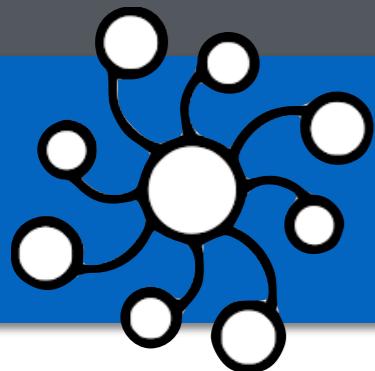
Redes Neurais Convolucionais



Recurrent Neural Networks



Generative Adversarial Networks



Generative Adversarial Network



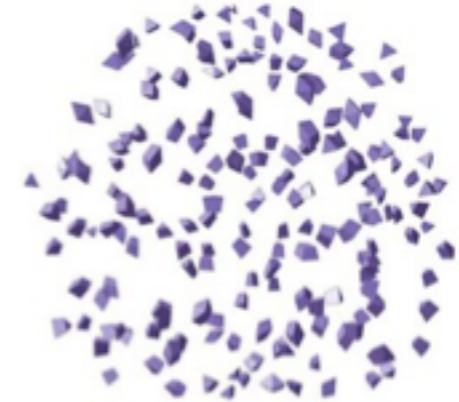
D: Detective



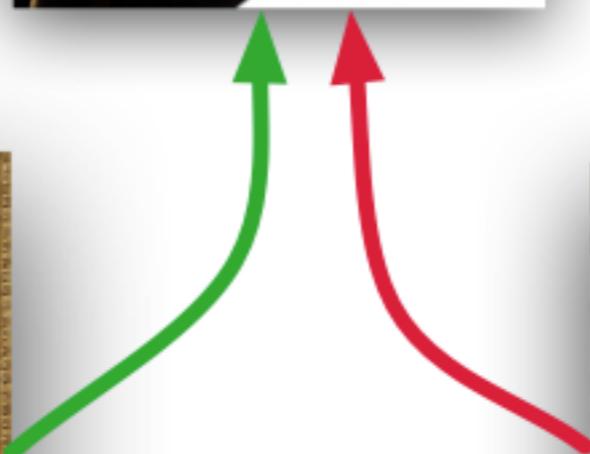
R: Real Data

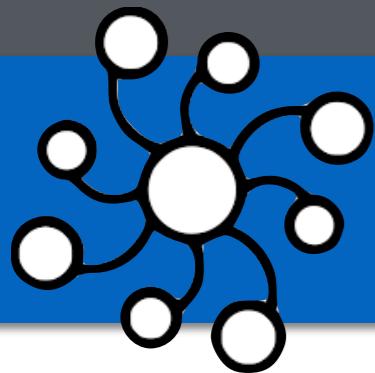


G: Generator (Forger)



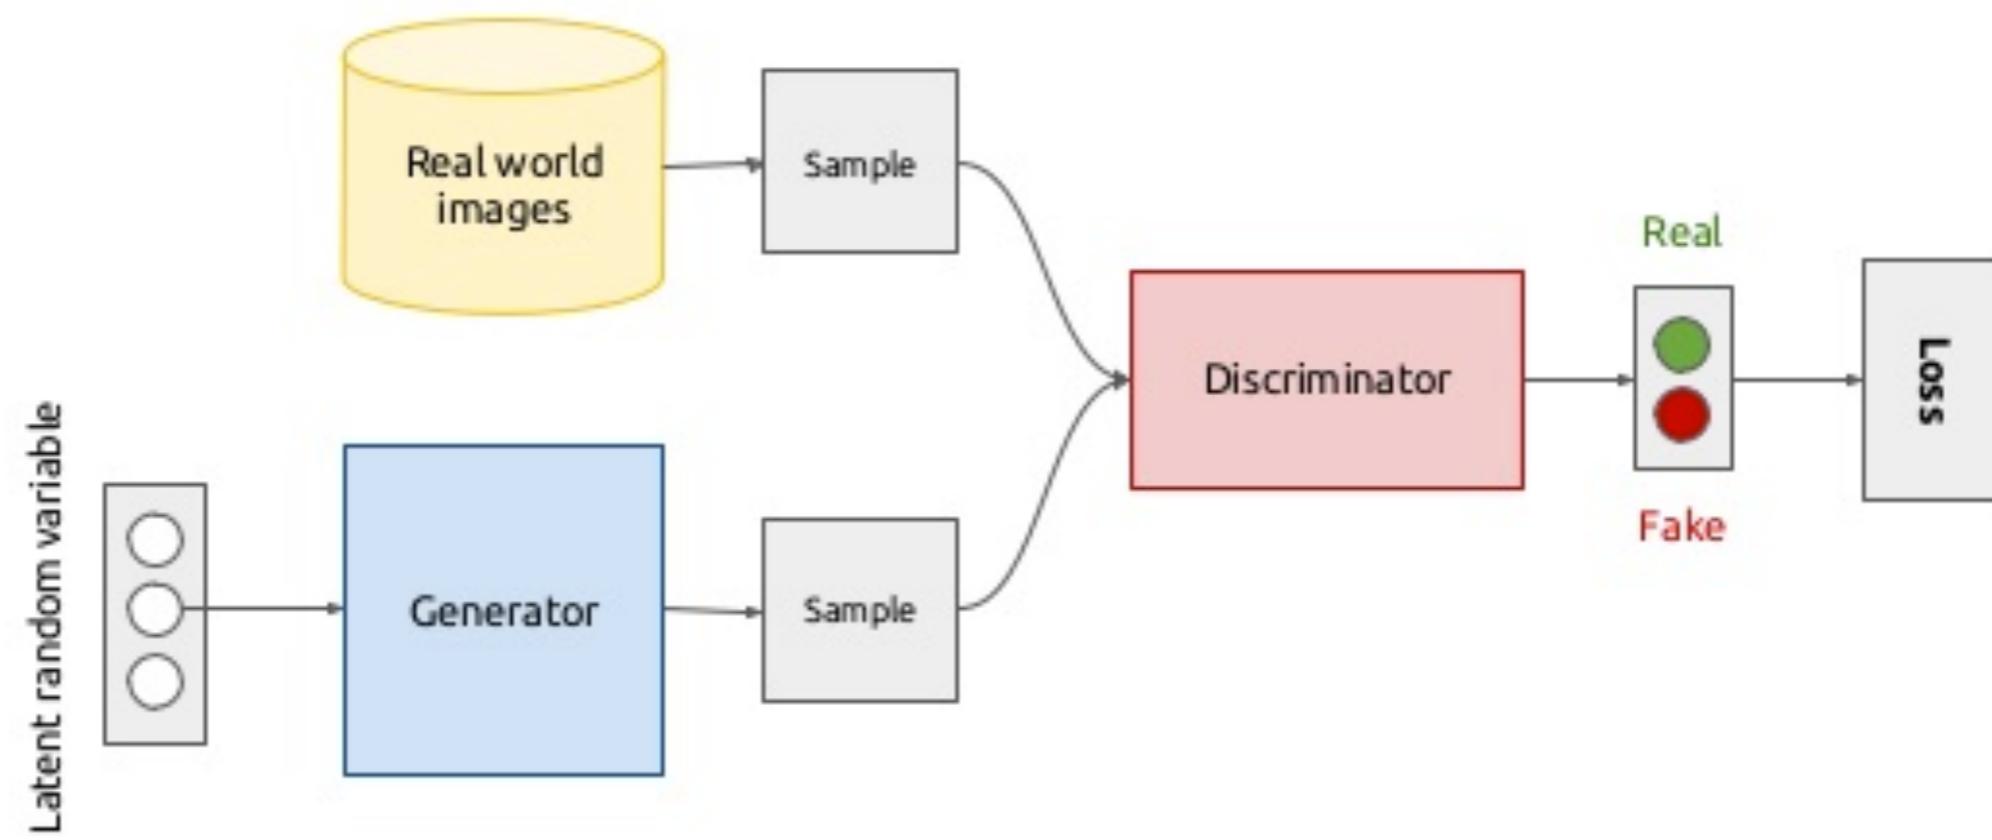
I: Input for Generator

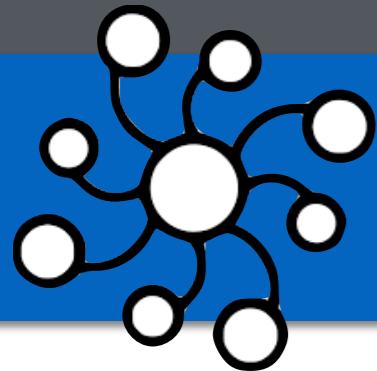




Generative Adversarial Network

Generative adversarial networks (conceptual)





Discriminator

Input



Picture of Money

Deep Convolutional
Neural Network

Output

True

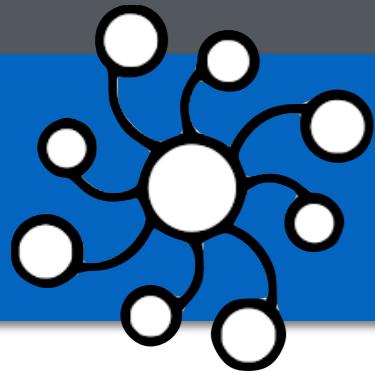
Whether or not the
picture contains money

Introdução

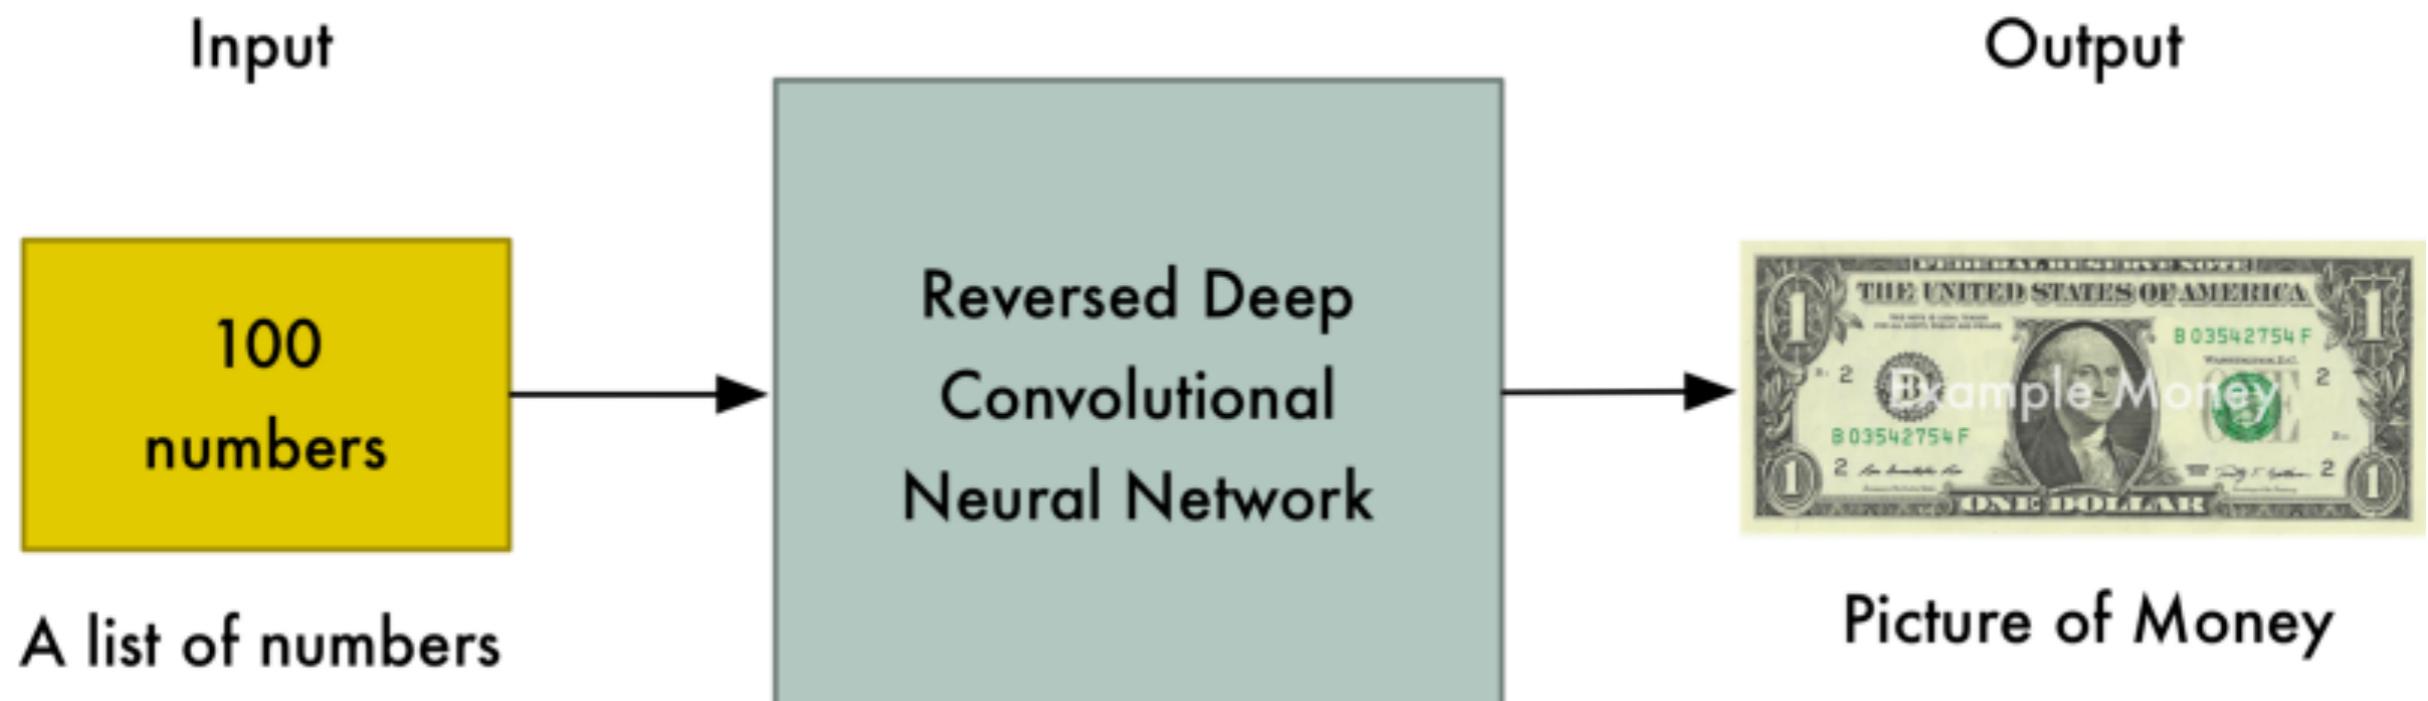
CNN

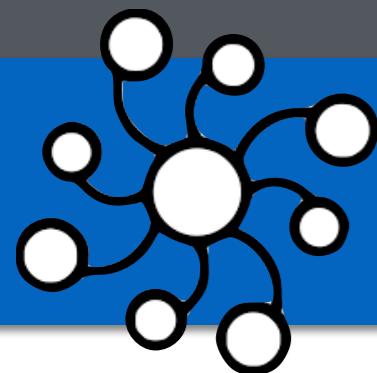
RNN

GAN



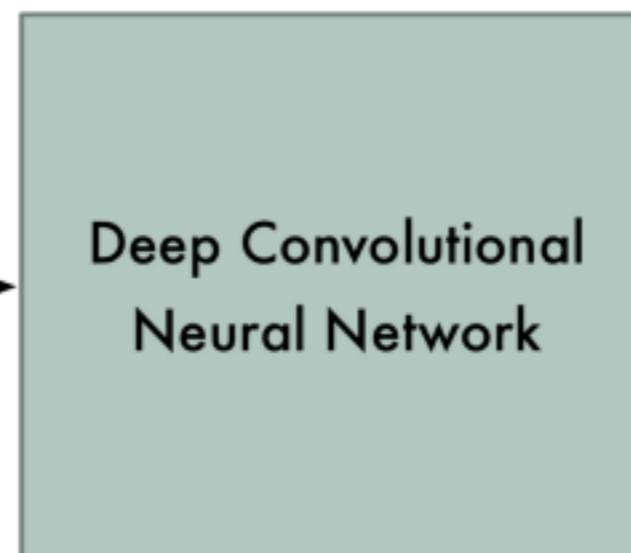
Generator





Discriminator Evolution

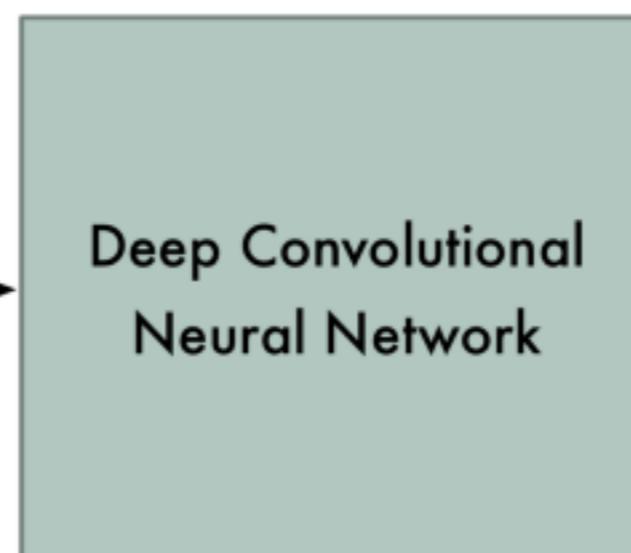
Input



Output

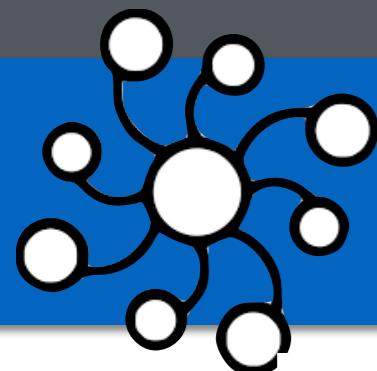
Whether or not the
picture contains money

Input

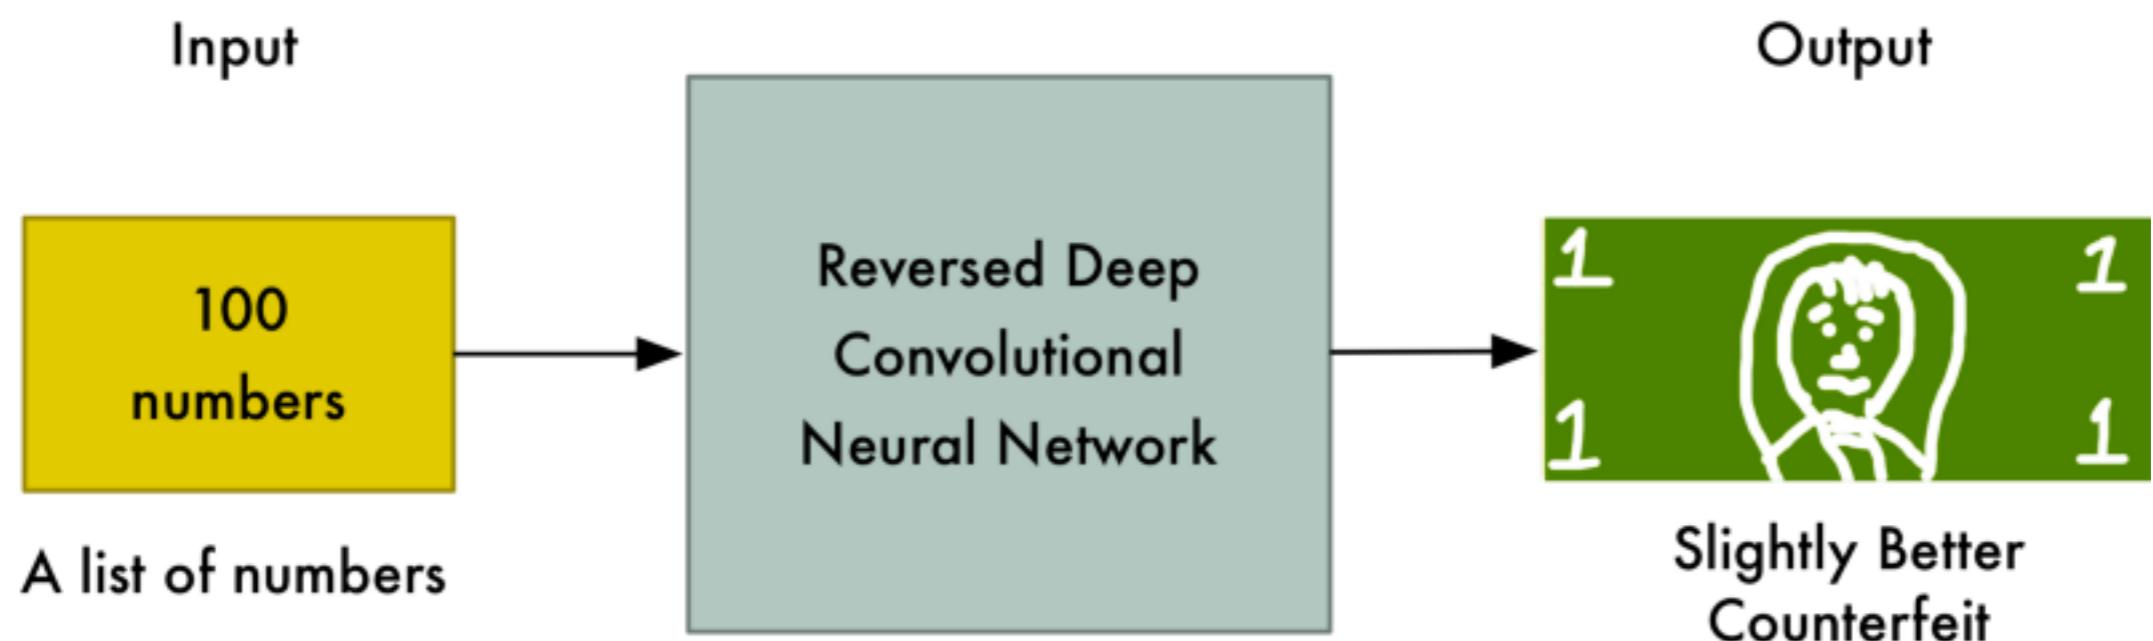
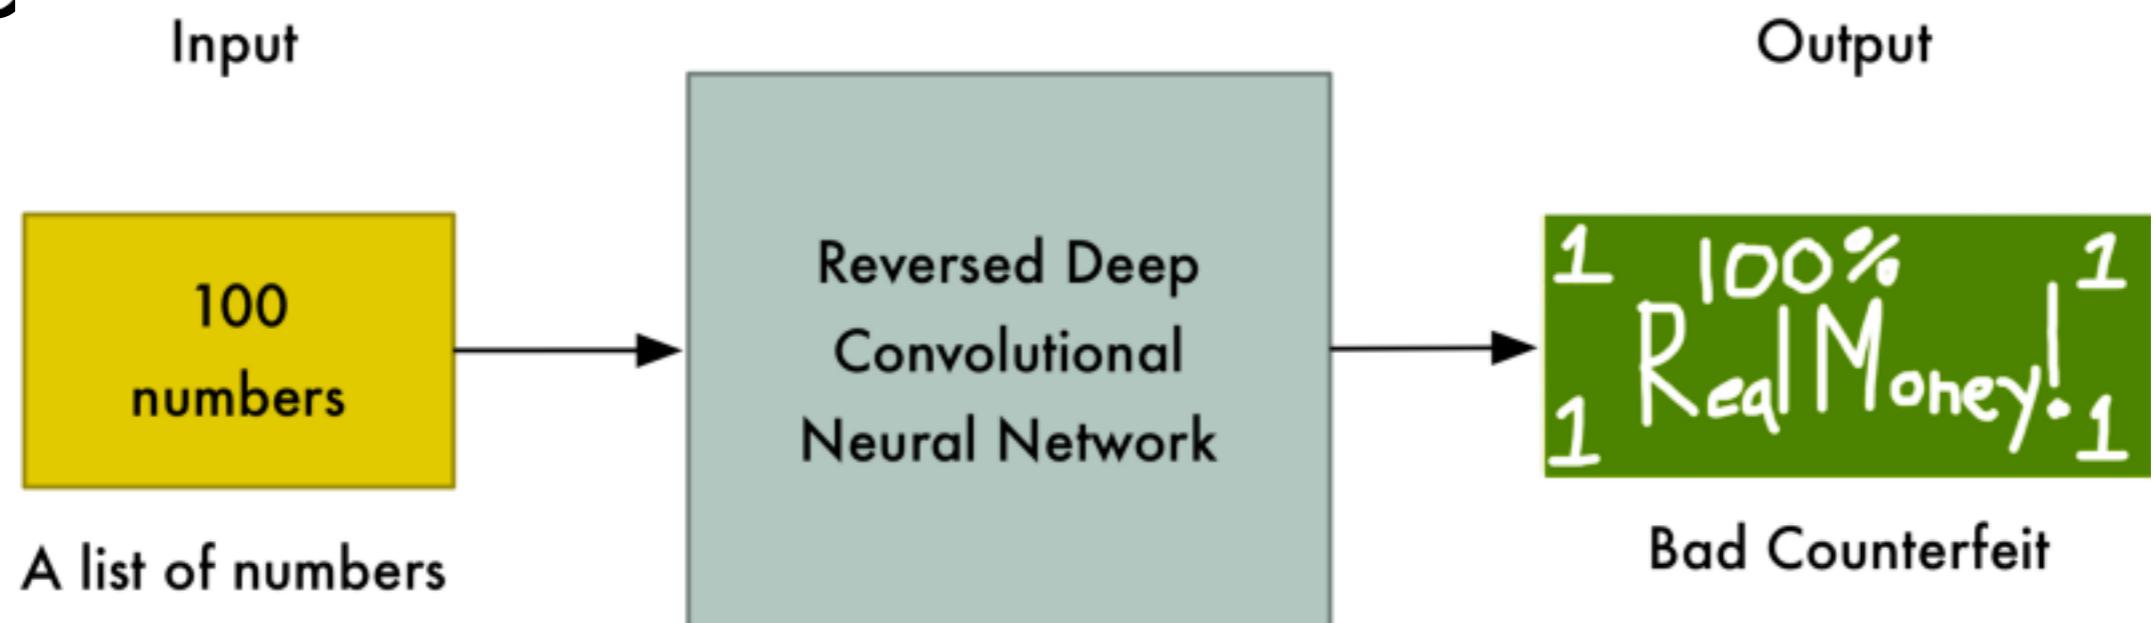


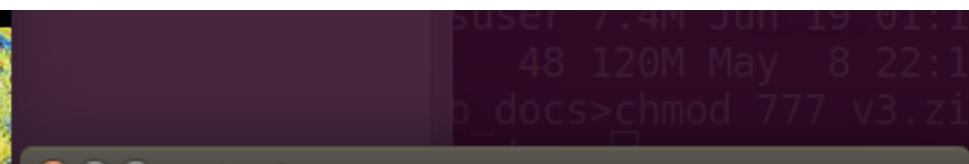
Output

Whether or not the
picture contains money



Discriminator Evolution





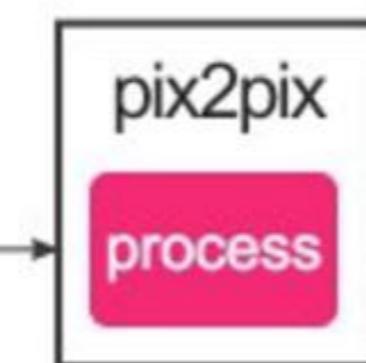
DOL

INPUT

OUTPUT

line 

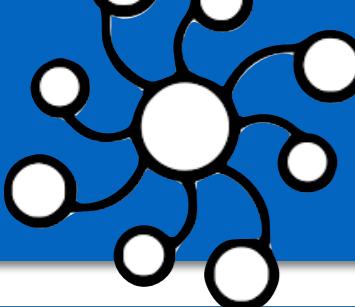
eraser 



undo

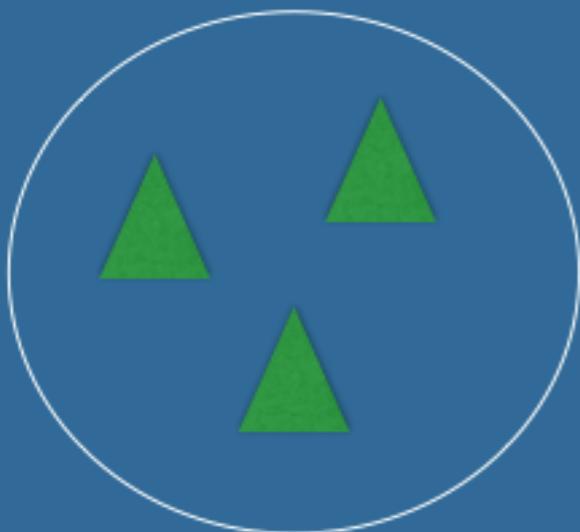
clear

save

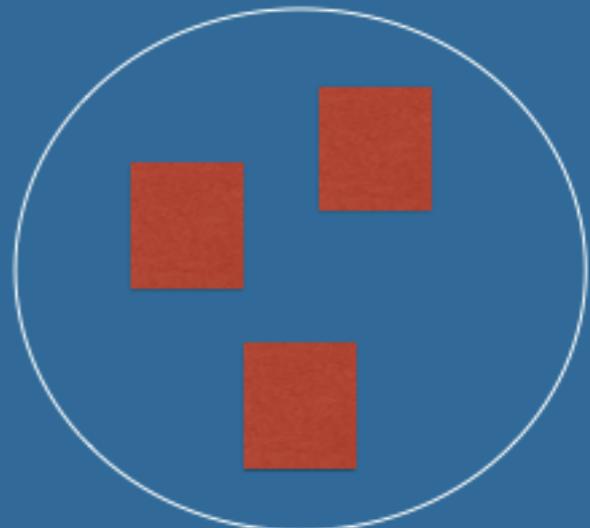


Transfer Learning

Traditional ML



Task / domain A

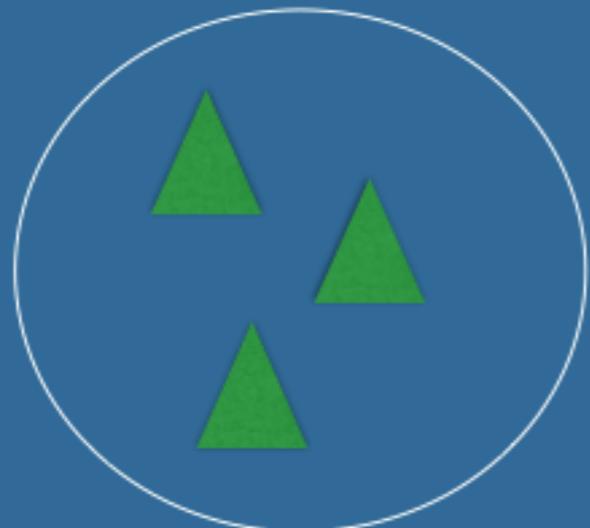


Task / domain B

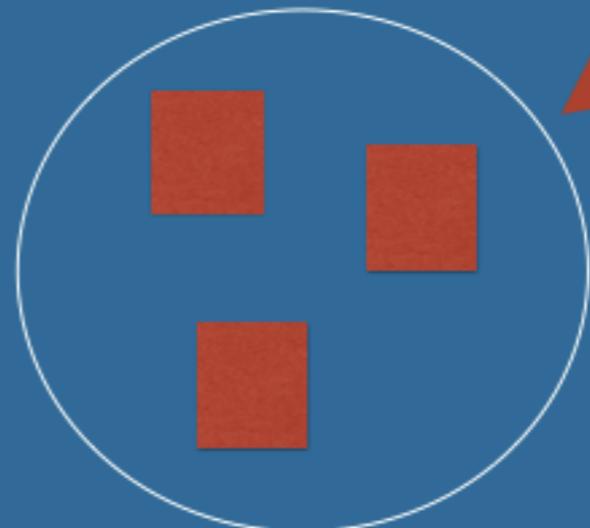
Training and evaluation on the same task or domain.

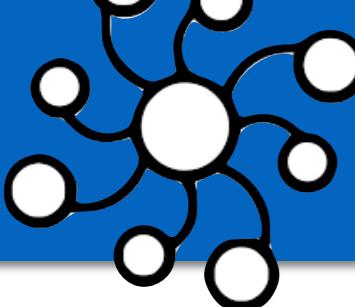


Model A



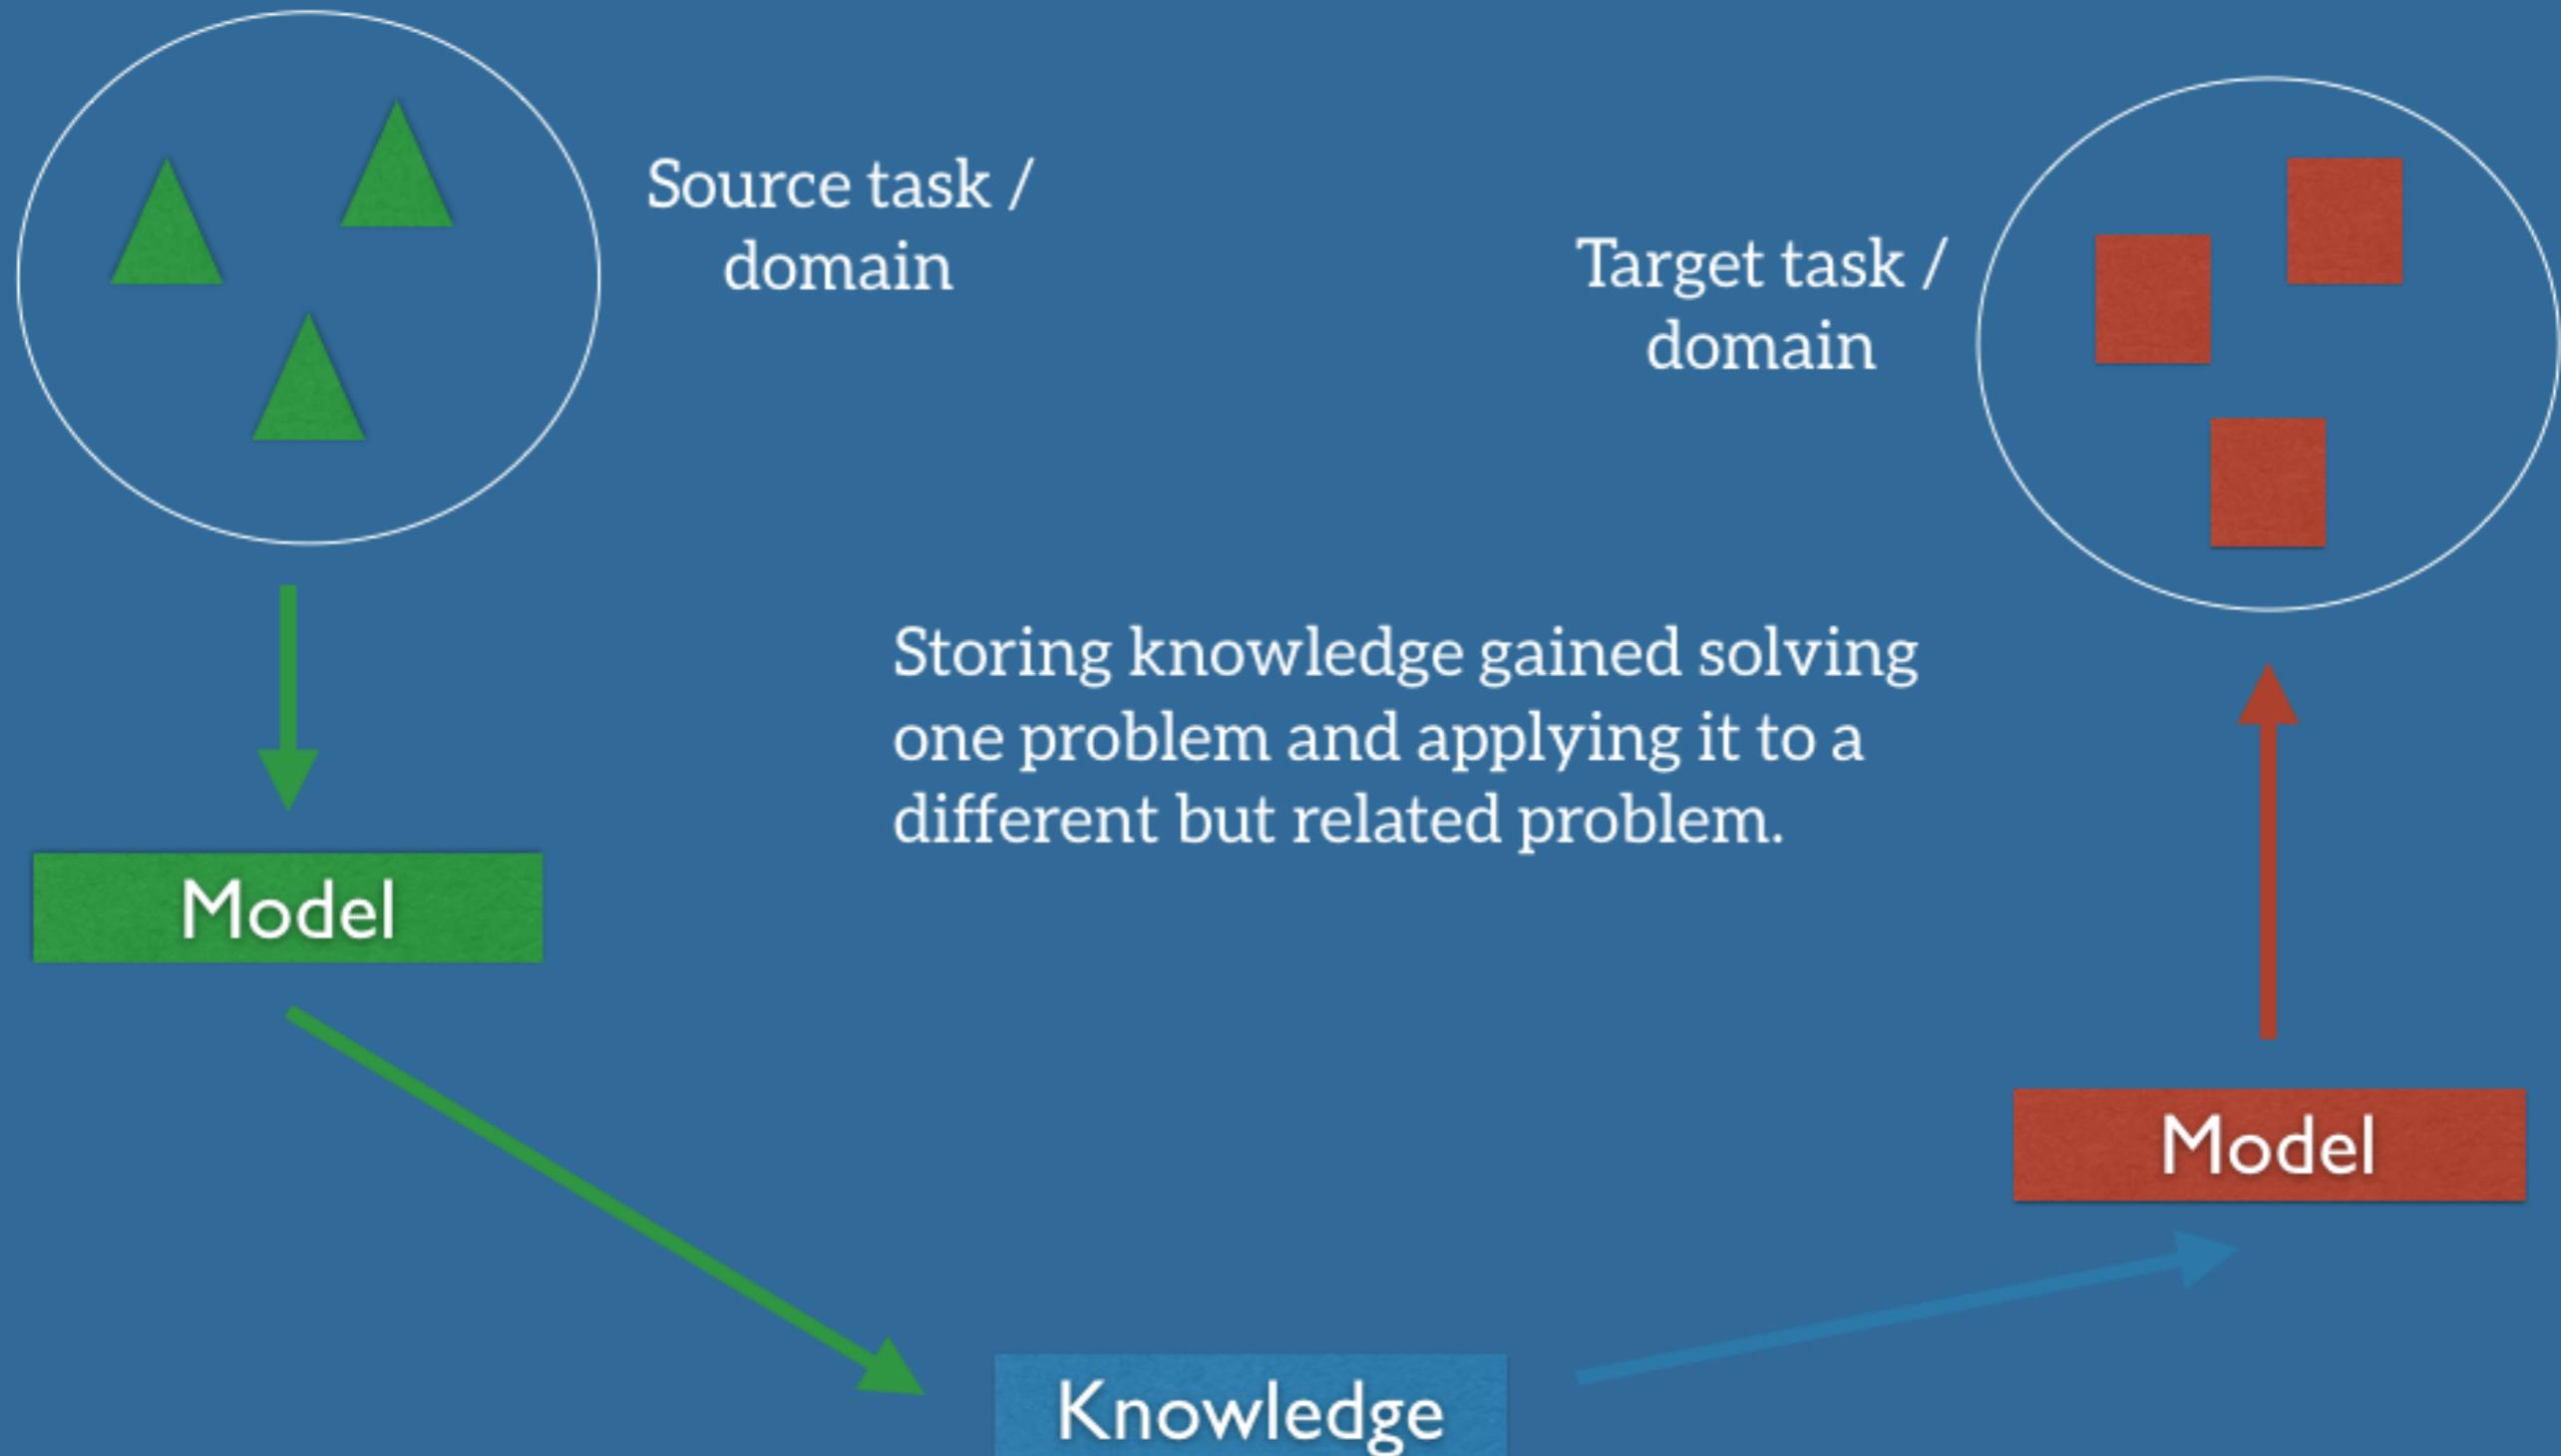
Model B

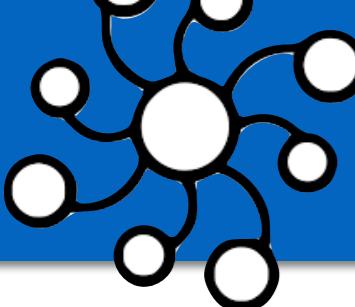




Transfer Learning

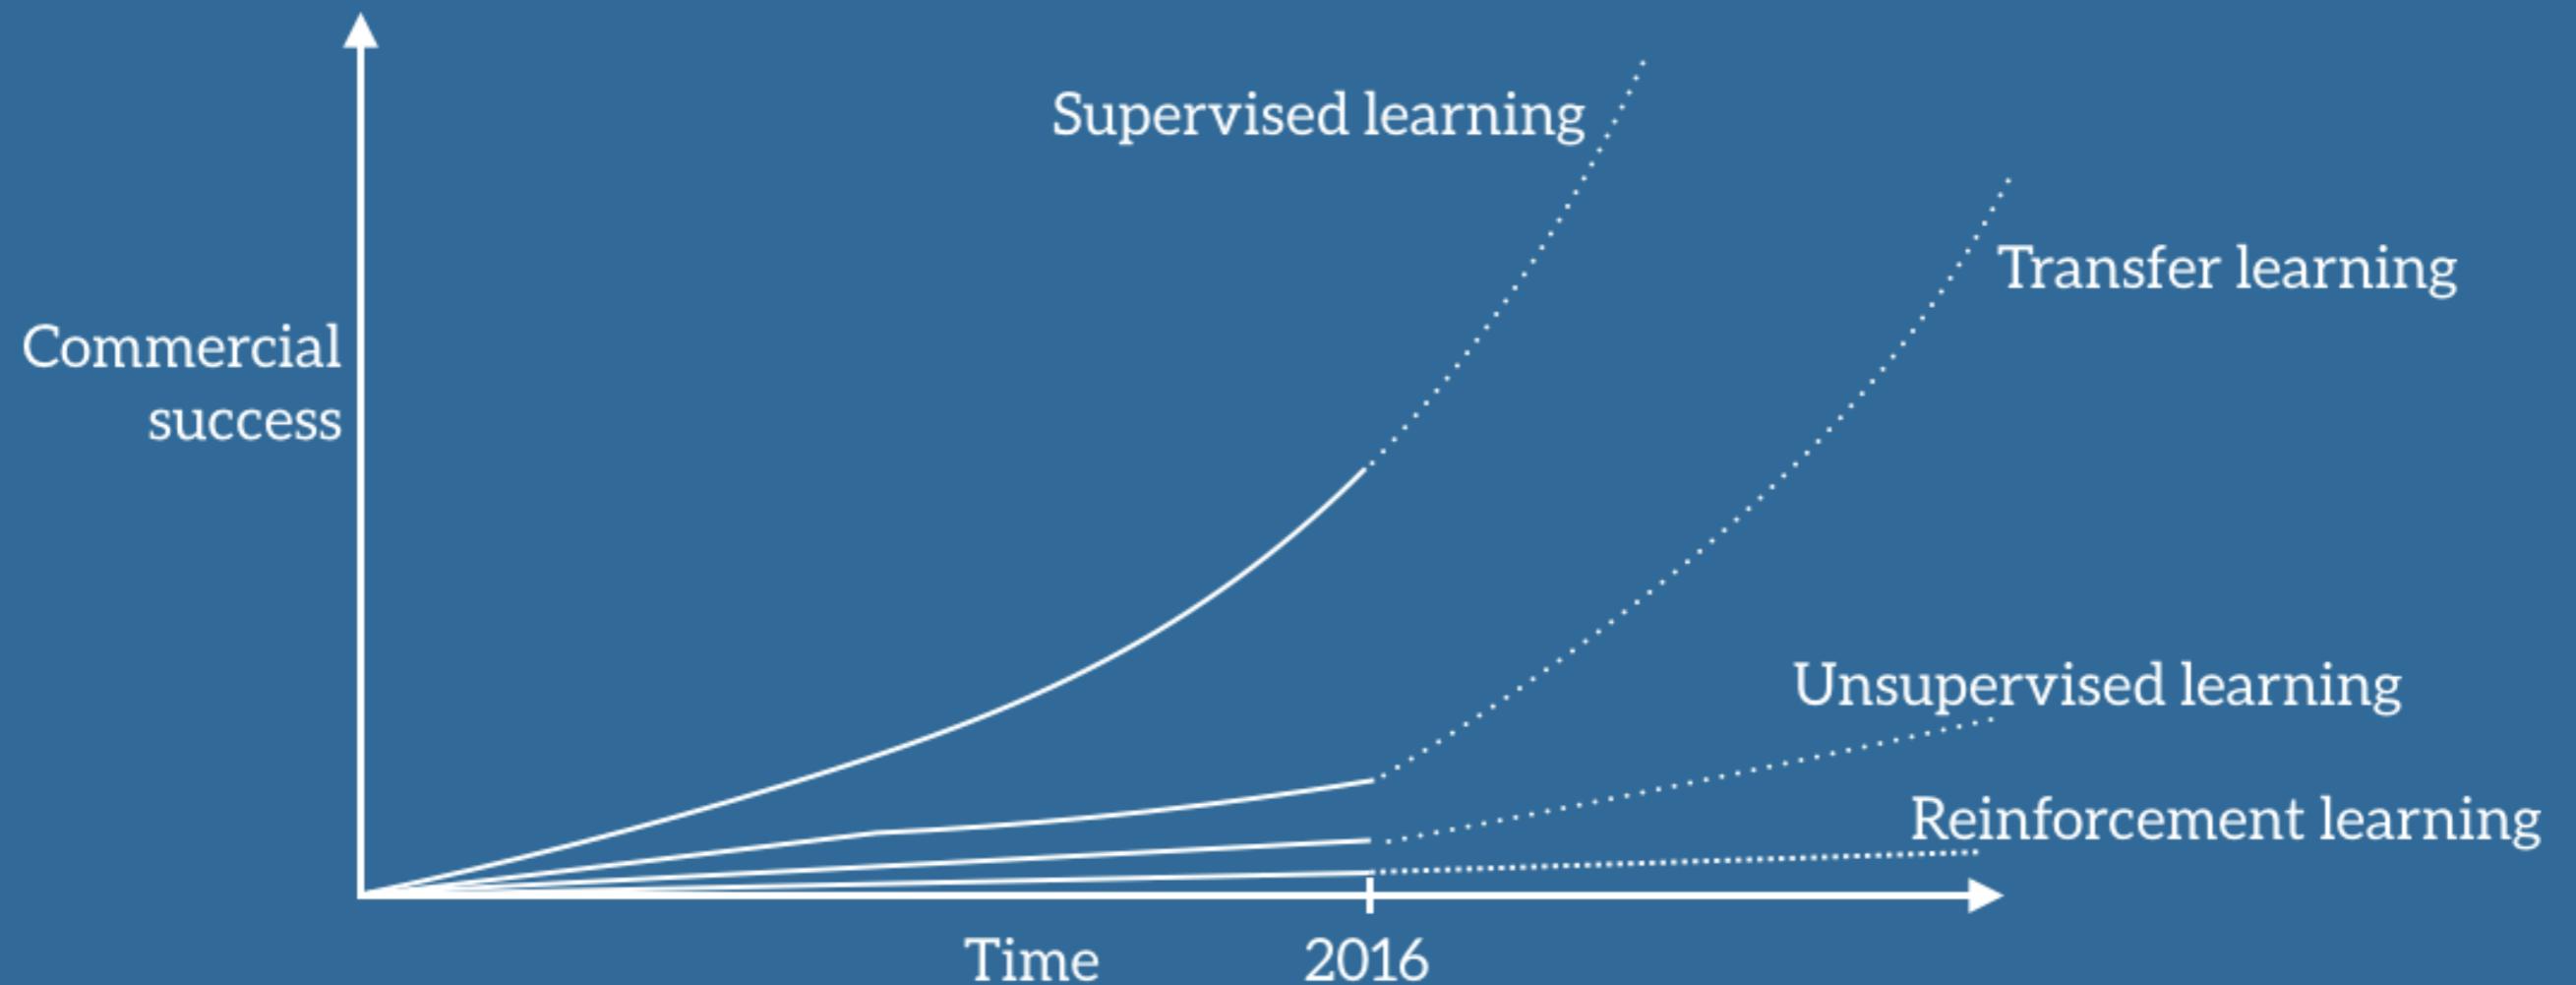
Transfer learning





Transfer Learning

Drivers of ML success in industry

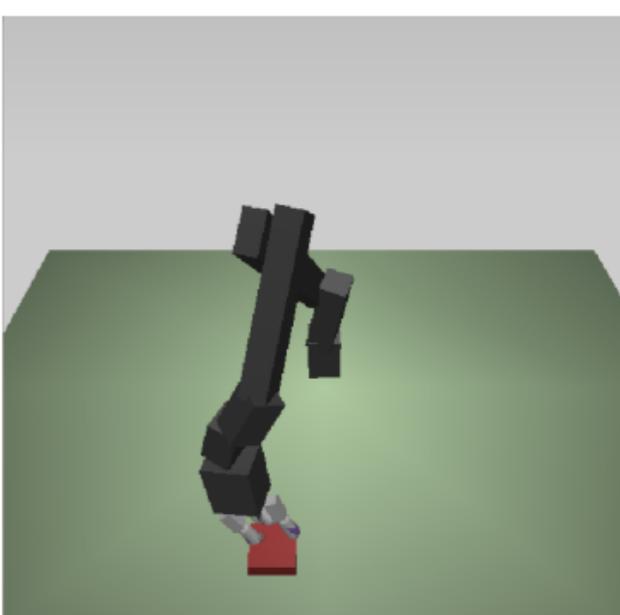
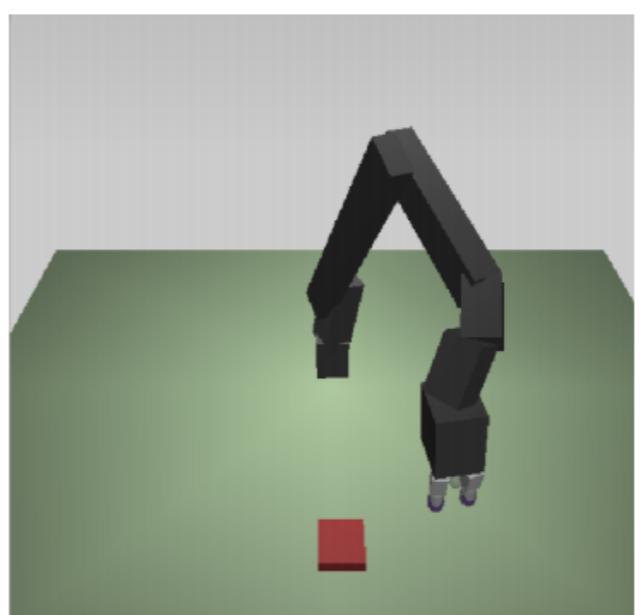
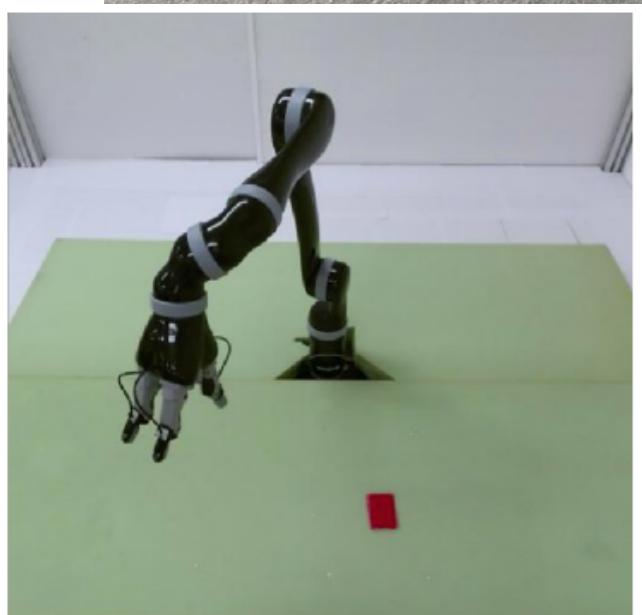
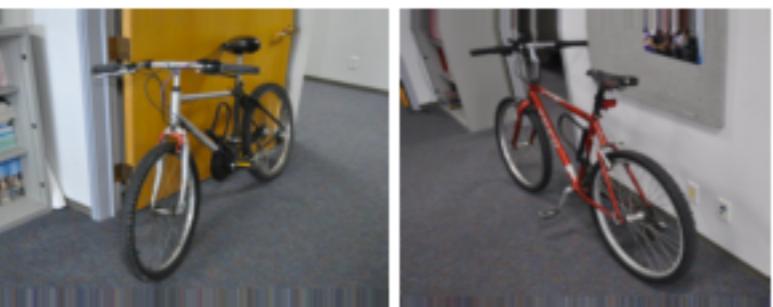


- Andrew Ng, NIPS 2016 tutorial

Domain 1



Domain 2



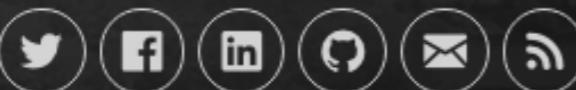
<https://alexisbcook.github.io/2017/using-transfer-learning-to-classify-images-with-keras/>



Alexis Cook

Deep Learning Professional

Blog



```
from keras.datasets import cifar10  
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

Extracting the InceptionV3 Bottleneck Features

Instead of building a CNN from scratch, I used **transfer learning** to leverage a pre-trained CNN that has demonstrated state-of-the-art performance in object classification tasks.

Keras makes it very easy to access several pre-trained **CNN architectures**. I decided to use the InceptionV3 architecture.



After importing the necessary Python class, it's only one line of code to get the model, along with the pre-trained weights.

```
from keras.applications.inception_v3 import InceptionV3  
base_model = InceptionV3(weights='imagenet', include_top=True)
```

The pre-trained InceptionV3 architecture is stored in the variable `base_model`.